

CS154

Lecture 6: Streaming Algorithms and Communication Complexity

Streaming Algorithms

Streaming Algorithms

Q



$L = \{x \mid x \text{ has more 1's than 0's}\}$



Initialize $C := 0$ and $B := 0$

Read the next bit x from the stream

If $(C = 0)$ then $B := x$, $C := 1$

If $(C \neq 0)$ and $(B = x)$ then $C := C + 1$

If $(C \neq 0)$ and $(B \neq x)$ then $C := C - 1$

When the stream stops,

accept if $B=1$ and $C > 0$, else ***reject***

B = the majority bit
 C = how many more times that B appears

On all strings of length n , the algorithm uses $(1 + \log_2 n)$ bits of space (to store B and C)

Streaming Algorithms

01011101



Can
recognize
non-regular
languages

Streaming algorithms differ from DFAs in several significant ways:

1. Streaming algorithms can output more than one bit
2. The “memory” or “space” of a streaming algorithm can (slowly) increase as it reads longer strings
3. Could also make multiple passes over the data, could be randomized

DFAs and Streaming

01011101



Theorem: Suppose a language L can be recognized by a **DFA** with $\leq 2^p$ states. Then L is computable by a **streaming algorithm A** using $\leq p$ **bits** of space.

Proof Idea: Algorithm A stores the DFA's current state in memory, beginning with the start state. Alg. A makes decisions based on DFA transitions. When the string ends, A outputs *accept* if the DFA state is accepting, *reject* otherwise.

DFAs and Streaming

01011101



For any $L \subseteq \Sigma^*$ define $L_n = L \cap \Sigma^n$

Theorem: Suppose L is computable by a streaming algorithm A using $f(n)$ bits of space, on all strings of length n . Then for all n , L_n is recognized by a DFA with $\leq 2^{f(n)}$ states.

Proof Idea: The new DFA will have a state for each of the $2^{f(n)}$ possible configurations of A 's memory. When A sees a symbol, its memory will update; the transition function of the DFA can simulate that.

$L = \{x \mid x \text{ has more 1's than 0's}\}$



Is there a streaming algorithm for L using much *less than* $(\log_2 n)$ space?

Theorem: Every streaming algorithm for L needs at least $(\log_2 n) - 1$ bits of space

We will use:

- Myhill-Nerode Theorem
- The connection between DFAs and streaming

$L = \{x \mid x \text{ has more 1's than 0's}\}$

Theorem: Every streaming algorithm for L requires at least $(\log_2 n) - 1$ bits of space

Proof Idea: Let n be even, and $L_n \subseteq \{0,1\}^n \cap L$

We will give a set S_n of $n/2 + 1$ strings such that each pair in S_n is *distinguishable* in L_n


Myhill-Nerode \Rightarrow Every DFA recognizing L_n needs at least $n/2 + 1$ states

\Rightarrow Every streaming algorithm for L requires at least $(\log n) - 1$ bits of memory

$L = \{x \mid x \text{ has more 1's than 0's}\}$

Theorem: Every streaming algorithm for L requires at least $(\log_2 n) - 1$ bits of space

Suppose we partition all strings into their equivalence classes under \equiv_{L_n}



But the number of states in every DFA recognizing L_n is *at least* the number of equivalence classes under \equiv_{L_n} 10

$L = \{x \mid x \text{ has more 1's than 0's}\}$

Theorem: Every streaming algorithm for L requires at least $(\log_2 n) - 1$ bits of space

Proof (Slide 1): Let $S_n = \{0^{n/2-i} 1^i \mid i = 0, \dots, n/2\}$

Let $x = 0^{n/2-k} 1^k$ and $y = 0^{n/2-j} 1^j$ be from S_n , $k > j$

Claim: $z = 0^{k-1} 1^{n/2-(k-1)}$ distinguishes x and y in L_n

xz has $n/2 - 1$ zeroes and $n/2 + 1$ ones $\Rightarrow xz \in L_n$

yz has $n/2 + (k - j - 1)$ zeroes and $n/2 - (k - j - 1)$ ones

But $k - j - 1 \geq 0$... so $yz \notin L_n$

So $x \not\equiv_{L_n} y$, because z distinguishes x and y

$L = \{x \mid x \text{ has more 1's than 0's}\}$

Theorem: Every streaming algorithm for L requires at least $(\log_2 n) - 1$ bits of space

Proof (Slide 2):

All pairs of strings in S_n are distinguishable in L_n

\Rightarrow There are at least $|S_n|$ equiv classes of \equiv_{L_n}

Then, from the Myhill-Nerode Theorem:

\Rightarrow All DFAs recognizing L_n need $\geq |S_n|$ states

\Rightarrow Every streaming algorithm for L requires at least $(\log_2 |S_n|)$ bits of space.

Recall $|S_n| = n/2 + 1$ and we're done!

Number of Distinct Elements

The DE problem

Input: $x \in \{0, 1, \dots, 2^k\}^*$, $2^k > |x|^2$

Output: The number of distinct elements
appearing in x

Note: There is a streaming algorithm for
DE using $O(k n)$ space

Theorem: Every streaming algorithm for
DE requires $\Omega(k n)$ space

Randomized Algorithms Help!

The DE problem

Input: $x \in \{0, 1, \dots, 2^k\}^*$, $2^k > |x|^2$

Output: The number of distinct elements
appearing in x

Theorem: There is a *randomized* streaming
algorithm that can approximate DE
to within 0.1% error, using $O(k + \log n)$ space!

See the lecture notes for more details.


Communication Complexity

Communication Complexity

A theoretical model of distributed computing

- **Function** $f: \{0,1\}^* \times \{0,1\}^* \rightarrow \{0,1\}$
 - Two inputs, $x \in \{0,1\}^*$ and $y \in \{0,1\}^*$
 - **We assume $|x|=|y|=n$. Think of n as HUGE**
 - **Two computers: Alice and Bob**
 - **Alice only** knows x , **Bob only** knows y
 - **Goal:** Compute $f(x, y)$ by communicating as few bits as possible between Alice and Bob
- We do not count computation cost.*** We only care about the number of bits communicated.


Alice and Bob Have a Conversation



In every step: Each bit sent is a function of the party's input and all the bits communicated so far in the conversation.

Communication cost = number of bits communicated
= 4 (in the example)

We assume Alice and Bob alternate in communicating, and the last bit sent is the value of $f(x,y)$



Def. A *protocol* for a function f is a pair of functions $A, B : \{0,1\}^* \times \{0,1\}^* \rightarrow \{0, 1, \text{STOP}\}$ with the semantics:

On input (x, y) , let $r := 0$, $b_0 = \epsilon$.


While $(b_r \neq \text{STOP})$,

$r++$

If r is odd, Alice sends $b_r = A(x, b_1 \cdots b_{r-1})$

else Bob sends $b_r = B(y, b_1 \cdots b_{r-1})$


Output b_{r-1} . Number of rounds = $r - 1$



Def. The **cost** of a protocol P for f on n -bit strings is

$$\max_{x, y \in \{0,1\}^n} [\text{number of rounds in } P \text{ to compute } f(x, y)]$$

The **communication complexity** of f on n -bit strings is the **minimum cost** over all protocols for f on n -bit strings
 = the minimum number of rounds used in any protocol for computing $f(x, y)$ over all n -bit x, y



Example. Let $f : \{0,1\}^* \times \{0,1\}^* \rightarrow \{0,1\}$ be arbitrary


There is always a “trivial” protocol:

Alice sends bits of x in odd rounds

Bob sends bits of y in even rounds

After $2n$ rounds, they both know each other's input!

The *communication complexity* of every f is at most $2n$



Example. $\text{PARITY}(x, y) = \sum_i x_i + \sum_i y_i \bmod 2$.

What's a good protocol for computing PARITY?

Alice sends $b_1 = (\sum_i x_i \bmod 2)$

Bob sends $b_2 = (b_1 + \sum_i y_i \bmod 2)$. **Alice** stops.

The communication complexity of PARITY is 2



$$f(x, y) = 0$$

x



$$f(x, y) = 0$$

y

Example. MAJORITY(x, y) = most frequent bit in xy

What's a good protocol for computing MAJORITY?

Alice sends **b** = number of 1s in **x**

Bob computes **c** = number of 1s in **y** ,

sends 1 iff **$b + c$** is greater than **$(|x| + |y|)/2 = n$**

Communication complexity of MAJORITY is $O(\log n)$



$$f(x,y)=0$$

x



$$f(x,y)=0$$

y


Example. $\text{EQUALS}(x, y) = 1 \Leftrightarrow x = y$

What's a good protocol for computing EQUALS?

?????

Communication complexity of EQUALS is at most $2n$

Connection to Streaming and DFAs



Let $L \subseteq \{0,1\}^*$

Def. $f_L: \{0,1\}^* \times \{0,1\}^* \rightarrow \{0,1\}$

for x, y with $|x|=|y|$ as:

$$f_L(x, y) = 1 \Leftrightarrow xy \in L$$

Examples:

$L = \{x \mid x \text{ has an odd number of 1s}\}$

$$\Rightarrow f_L(x, y) = \text{PARITY}(x, y) = \sum_i x_i + \sum_i y_i \bmod 2$$

$L = \{x \mid x \text{ has more 1s than 0s}\}$

$$\Rightarrow f_L(x, y) = \text{MAJORITY}(x, y)$$

$L = \{xx \mid x \in \{0,1\}^*\}$

$$\Rightarrow f_L(x, y) = \text{EQUALS}(x, y)$$

Connection to Streaming and DFAs



x

y

Let $L \subseteq \{0,1\}^*$

Def. $f_L: \{0,1\}^* \times \{0,1\}^* \rightarrow \{0,1\}$

for x, y with $|x|=|y|$ as:

$$f_L(x, y) = 1 \Leftrightarrow xy \in L$$

Theorem: If L has a streaming algorithm using $\leq s$ space, then the comm. complexity of f_L is at most $4s + 5$.

Proof: Alice runs streaming algorithm A on x .

Sends the *memory content* of A : this is s bits of space

Bob starts up A with that memory content, runs A on y .

Gets an output bit, sends to Alice.

(...why $4s+5$ rounds? Can you do better?)