

CS154

Lecture 7: Finish up Communication, Start up Turing Machines

Communication Complexity

A theoretical model of distributed computing

- **Function** $f: \{0,1\}^* \times \{0,1\}^* \rightarrow \{0,1\}$
 - Two inputs, $x \in \{0,1\}^*$ and $y \in \{0,1\}^*$
 - **We assume $|x|=|y|=n$. Think of n as HUGE**
- **Two computers: Alice and Bob**
 - **Alice only** knows x , **Bob only** knows y
- **Goal:** Compute $f(x, y)$ by communicating as few bits as possible between Alice and Bob

We do not count computation cost. We *only* care about the number of bits communicated.



$$f(x, y) = 0$$

$$A(x, \varepsilon) = 0$$

$$B(y, 0) = 1$$

$$A(x, 01) = 1$$

$$B(y, 011) = 0$$



$$f(x, y) = 0$$

x

y

Def. A *protocol* for a function f is a pair of functions $A, B : \{0,1\}^* \times \{0,1\}^* \rightarrow \{0, 1, \text{STOP}\}$ with the semantics:

On input (x, y) , let $r := 0$, $b_0 = \varepsilon$.


While $(b_r \neq \text{STOP})$,

$r++$

If r is odd, Alice sends $b_r = A(x, b_1 \cdots b_{r-1})$

else Bob sends $b_r = B(y, b_1 \cdots b_{r-1})$

Output b_{r-1} . Number of rounds = $r - 1$




Def. The **cost** of a protocol P for f on n -bit strings is

$$\max_{x, y \in \{0, 1\}^n} [\text{number of rounds in } P \text{ to compute } f(x, y)]$$

The **communication complexity** of f on n -bit strings is the **minimum cost** over all protocols for f on n -bit strings

= the minimum number of rounds used in any protocol for computing $f(x, y)$ over all n -bit x, y



Example. Let $f : \{0,1\}^* \times \{0,1\}^* \rightarrow \{0,1\}$ be arbitrary

There is always a “trivial” protocol:

Alice sends bits of x in odd rounds

Bob sends bits of y in even rounds

After $2n$ rounds, they both know each other's input!

The *communication complexity* of every f is at most $2n$



$$f(x,y)=0$$

x



$$f(x,y)=0$$

y

Example. $\text{EQUALS}(x, y) = 1 \Leftrightarrow x = y$

What's a good protocol for computing EQUALS?

?????

Communication complexity of EQUALS is at most $2n$

Connection to Streaming and DFAs



x

y

Let $L \subseteq \{0,1\}^*$

Def. $f_L: \{0,1\}^* \times \{0,1\}^* \rightarrow \{0,1\}$

for x, y with $|x|=|y|$ as:

$$f_L(x, y) = 1 \Leftrightarrow xy \in L$$

Examples:

$L = \{x \mid x \text{ has an odd number of 1s}\}$

$$\Rightarrow f_L(x, y) = \text{PARITY}(x, y) = \sum_i x_i + \sum_i y_i \bmod 2$$

$L = \{x \mid x \text{ has more 1s than 0s}\}$

$$\Rightarrow f_L(x, y) = \text{MAJORITY}(x, y)$$

$L = \{xx \mid x \in \{0,1\}^*\}$

$$\Rightarrow f_L(x, y) = \text{EQUALS}(x, y)$$

Connection to Streaming and DFAs



x

y

Let $L \subseteq \{0,1\}^*$

Def. $f_L: \{0,1\}^* \times \{0,1\}^* \rightarrow \{0,1\}$

for x, y with $|x|=|y|$ as:

$$f_L(\mathbf{x}, \mathbf{y}) = 1 \Leftrightarrow \mathbf{xy} \in L$$

Theorem: If L has a streaming algorithm using $\leq s$ space, then the comm. complexity of f_L is at most $4s + 5$.

Proof: Alice runs streaming algorithm **A** on **x**.

Sends the *memory content* of **A**: this is s bits of space

Bob starts up **A** with that memory content, runs **A** on **y**.

Gets an output bit, sends to Alice.

(...why $4s+5$ rounds? Can you do better?)

Connection to Streaming and DFAs

Let $L \subseteq \{0,1\}^*$ Def. $f_L(x, y) = 1 \Leftrightarrow xy \in L$

Theorem: If L has a streaming algorithm using $\leq s$ space, then the comm. complexity of f_L is at most $4s + 5$.

Corollary: For every regular language L , the comm. complexity of f_L is $O(1)$.

Example: Comm. Complexity of PARITY is $O(1)$

Corollary: Comm. Complexity of MAJORITY is $O(\log n)$

What about the Comm. Complexity of EQUALS?

Communication complexity of EQUALS

Theorem: The comm. complexity of EQUALS is $\Theta(n)$.
In particular, *every* protocol for EQUALS needs $\geq n$ bits of communication.

*No communication protocol can do much better than
“send your entire input”!*


Corollary: $L = \{ww \mid w \text{ in } \{0,1\}^*\}$ is not regular.
Moreover, every streaming algorithm for L needs $c n$ bits of memory, for some constant $c > 0$!

Communication complexity of EQUALS


Theorem: The comm. complexity of EQUALS is $\Theta(n)$.
In particular, *every* protocol for EQUALS needs $\geq n$ bits of communication.

Idea: Consider all possible ways they can communicate.

Definition: The *communication pattern* of a protocol on (x, y) is the sequence of bits that Alice & Bob send.



x



Pattern: 0110



y

Communication complexity of EQUALS

Theorem: The comm. complexity of EQUALS is $\Theta(n)$.
In particular, *every* protocol for EQUALS needs $\geq n$ bits of communication.

Proof: By contradiction. Suppose CC of EQUALS is $\leq n - 1$.
Then there are $\leq 2^n - 1$ possible communication *patterns* of that protocol, over all pairs of inputs (x, y) .

Claim: There are $x \neq y$ such that on (x, x) and on (y, y) , the protocol uses the *same* pattern P .

What is the communication pattern on (x, y) ?

This pattern is also P ! (WHY?)

So Alice & Bob *output the same bit* on (x, y) and (x, x) .

But $\text{EQUALS}(x, y) = 0$ and $\text{EQUALS}(x, x) = 1$. **Contradiction!**

Randomized Protocols Help!

EQUALS needs $c n$ bits of communication,
but...

Theorem: For all (x, y) of n bits each, there is a *randomized* protocol for **EQUALS** (x, y) using only $O(\log n)$ bits of communication, which works with probability **99.9%**!

Turing Machines



Turing Machine



Turing Machine (1936)

230

A. M. TURING


[Nov. 12,

ON COMPUTABLE NUMBERS, WITH AN APPLICATION TO THE ENTSCHEIDUNGSPROBLEM

By A. M. TURING.

[Received 28 May, 1936.—Read 12 November, 1936.]

The “computable” numbers may be described briefly as the real numbers whose expressions as a decimal are calculable by finite means. Although the subject of this paper is ostensibly the computable *numbers*, it is almost equally easy to define and investigate computable functions of an integral variable or a real or computable variable, computable predicates, and so forth. The fundamental problems involved are, however, the same in each case, and I have chosen the computable numbers for explicit treatment as involving the least cumbrous technique. I hope shortly to give an account of the relations of the computable numbers, functions and predicates to the Entscheidungsproblem.



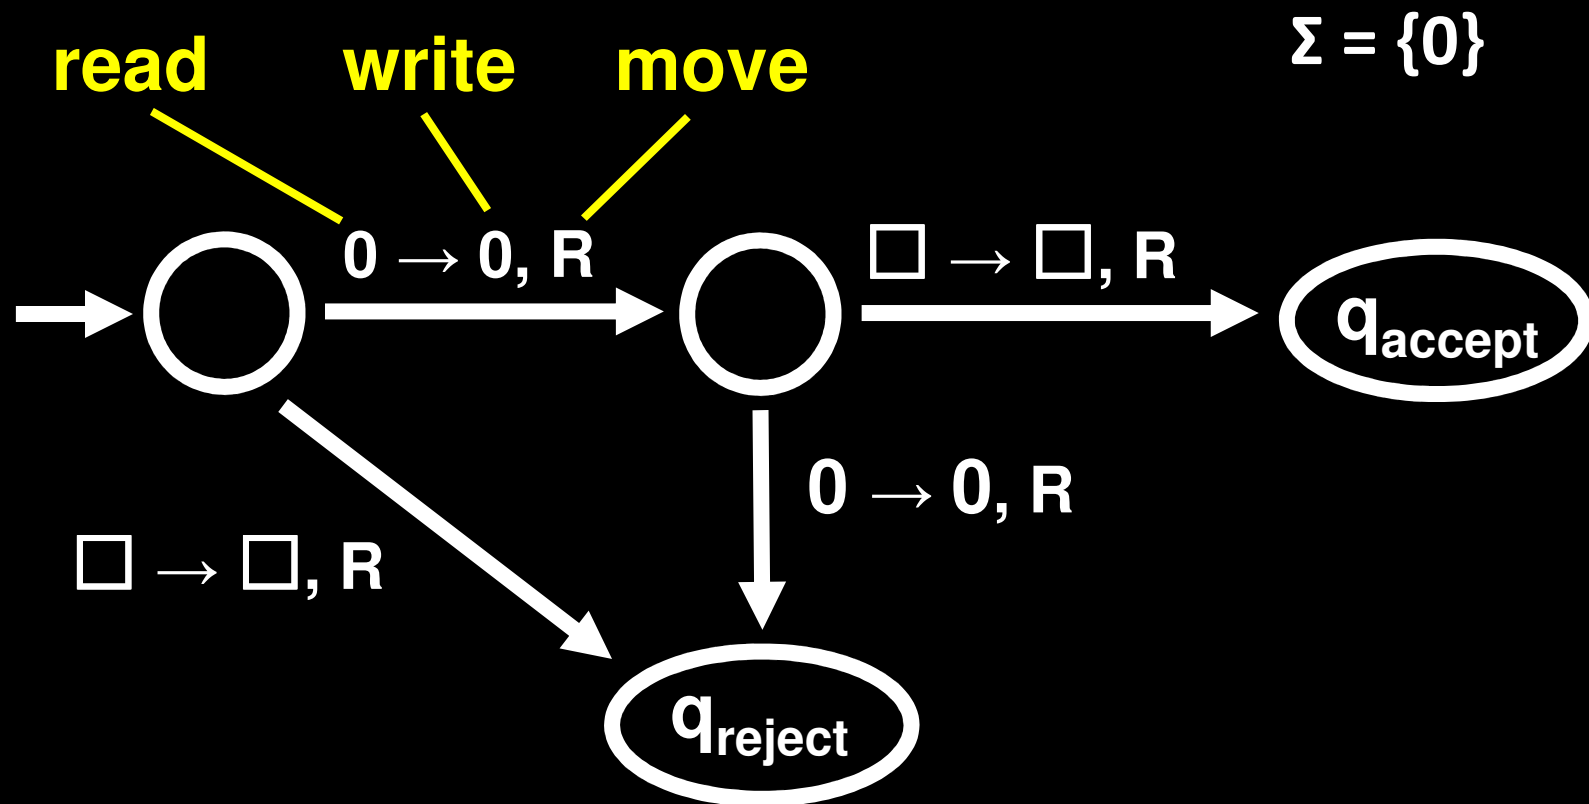
Turing Machines versus DFAs

TM can both *write to* and *read from* the tape

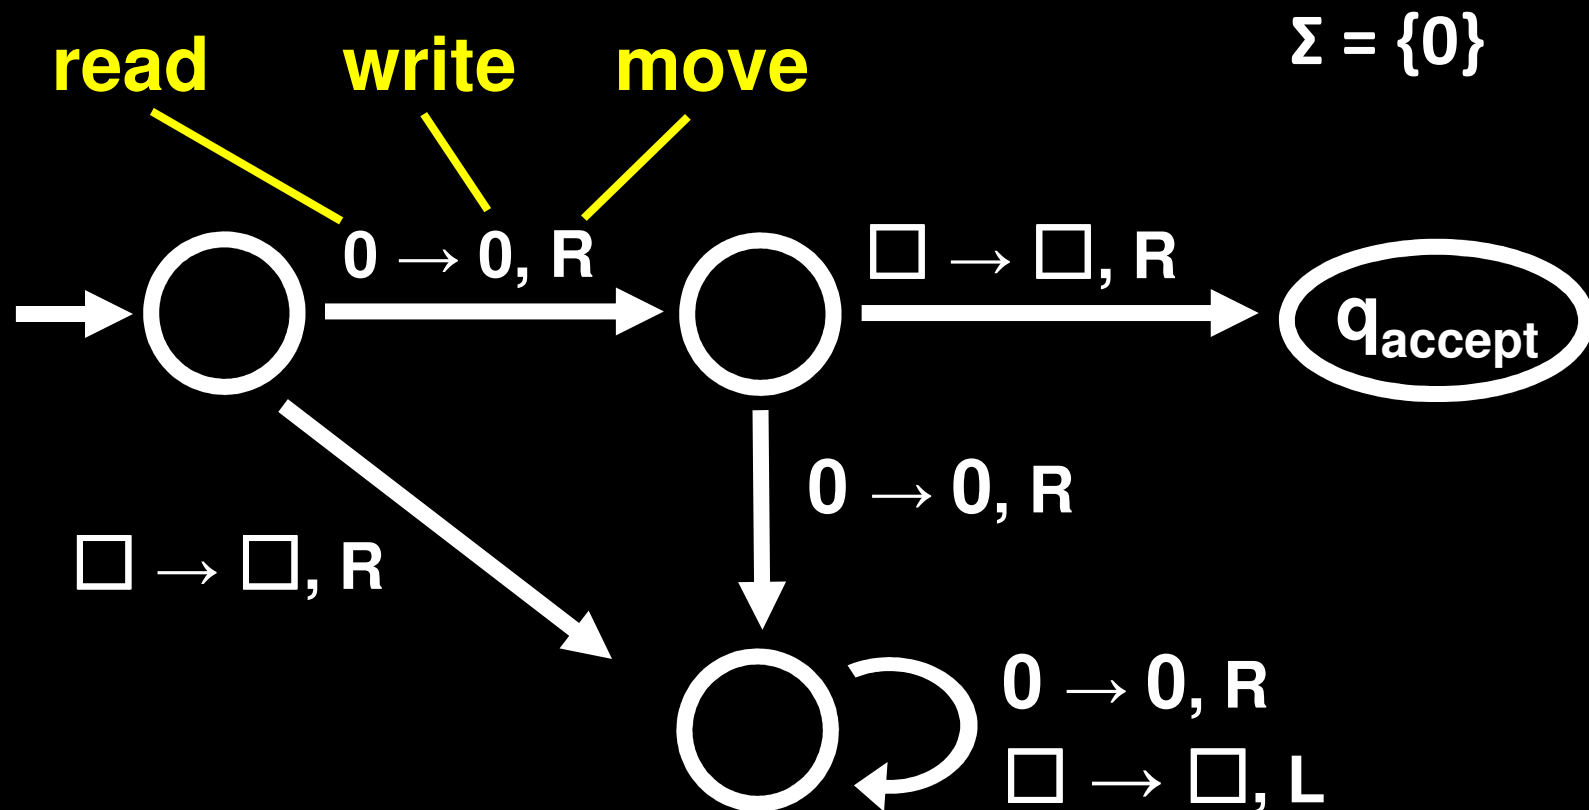
The head can move *left and right*

The input doesn't have to be read entirely,
and the computation can continue further
(even, *forever*) after all input has been read

Accept and **Reject** take immediate effect



This Turing machine *decides* the language $\{0\}$



This Turing machine *recognizes* the language $\{0\}$

Deciding the language $L = \{ w\#w \mid w \in \{0,1\}^* \}$



1. If there's no # on the tape (or more than one #), *reject*.
2. While there is a bit to the left of #,
 Replace the first bit with X, and check if the first bit b
 to the right of the # is identical. (If not, *reject*.)
 Replace that bit b with an X too.
3. If there's a bit to the right of #, then *reject* else *accept*

Definition: A Turing Machine is a 7-tuple

$T = (Q, \Sigma, \Gamma, \delta, q_0, q_{\text{accept}}, q_{\text{reject}})$, where:

Q is a finite set of states

Σ is the input alphabet, where $\square \notin \Sigma$

Γ is the tape alphabet, where $\square \in \Gamma$ and $\Sigma \subseteq \Gamma$


$\delta : Q \times \Gamma \rightarrow Q \times \Gamma \times \{L, R\}$

$q_0 \in Q$ is the start state

$q_{\text{accept}} \in Q$ is the accept state

$q_{\text{reject}} \in Q$ is the reject state, and $q_{\text{reject}} \neq q_{\text{accept}}$

Turing Machine Configurations



corresponds to the *configuration*:

$$11010q_700110 \in \{Q \cup \Gamma\}^*$$

Defining Acceptance and Rejection for TMs

Let C_1 and C_2 be configurations of a TM M

Definition. C_1 *yields* C_2 if M is in configuration C_2 after running M in configuration C_1 for one step

Suppose $\delta(q_1, b) = (q_2, c, L)$

Then aaq_1bb yields aq_2acb

Suppose $\delta(q_1, a) = (q_2, c, R)$

Then $cabq_1a$ yields $cabcq_2\Box$

Let $w \in \Sigma^*$ and M be a Turing machine

M *accepts* w if there are configs C_0, C_1, \dots, C_k , s.t.


- $C_0 = q_0w$
- C_i yields C_{i+1} for $i = 0, \dots, k-1$, and
- C_k contains the accept state q_{accept}

A TM M *recognizes* a language L
if M **accepts** exactly those strings in L


A language L is *recognizable*
(*a.k.a. recursively enumerable*)
if some TM **recognizes** L

A TM M *decides* a language L if M **accepts** all
strings in L and **rejects** all strings not in L

A language L is *decidable* (*a.k.a. recursive*)
if some TM **decides** L



L is **decidable**
(recursive)



L is **recognizable**
(recursively enumerable)

A Turing machine for deciding $\{ 0^{2^n} \mid n \geq 0 \}$

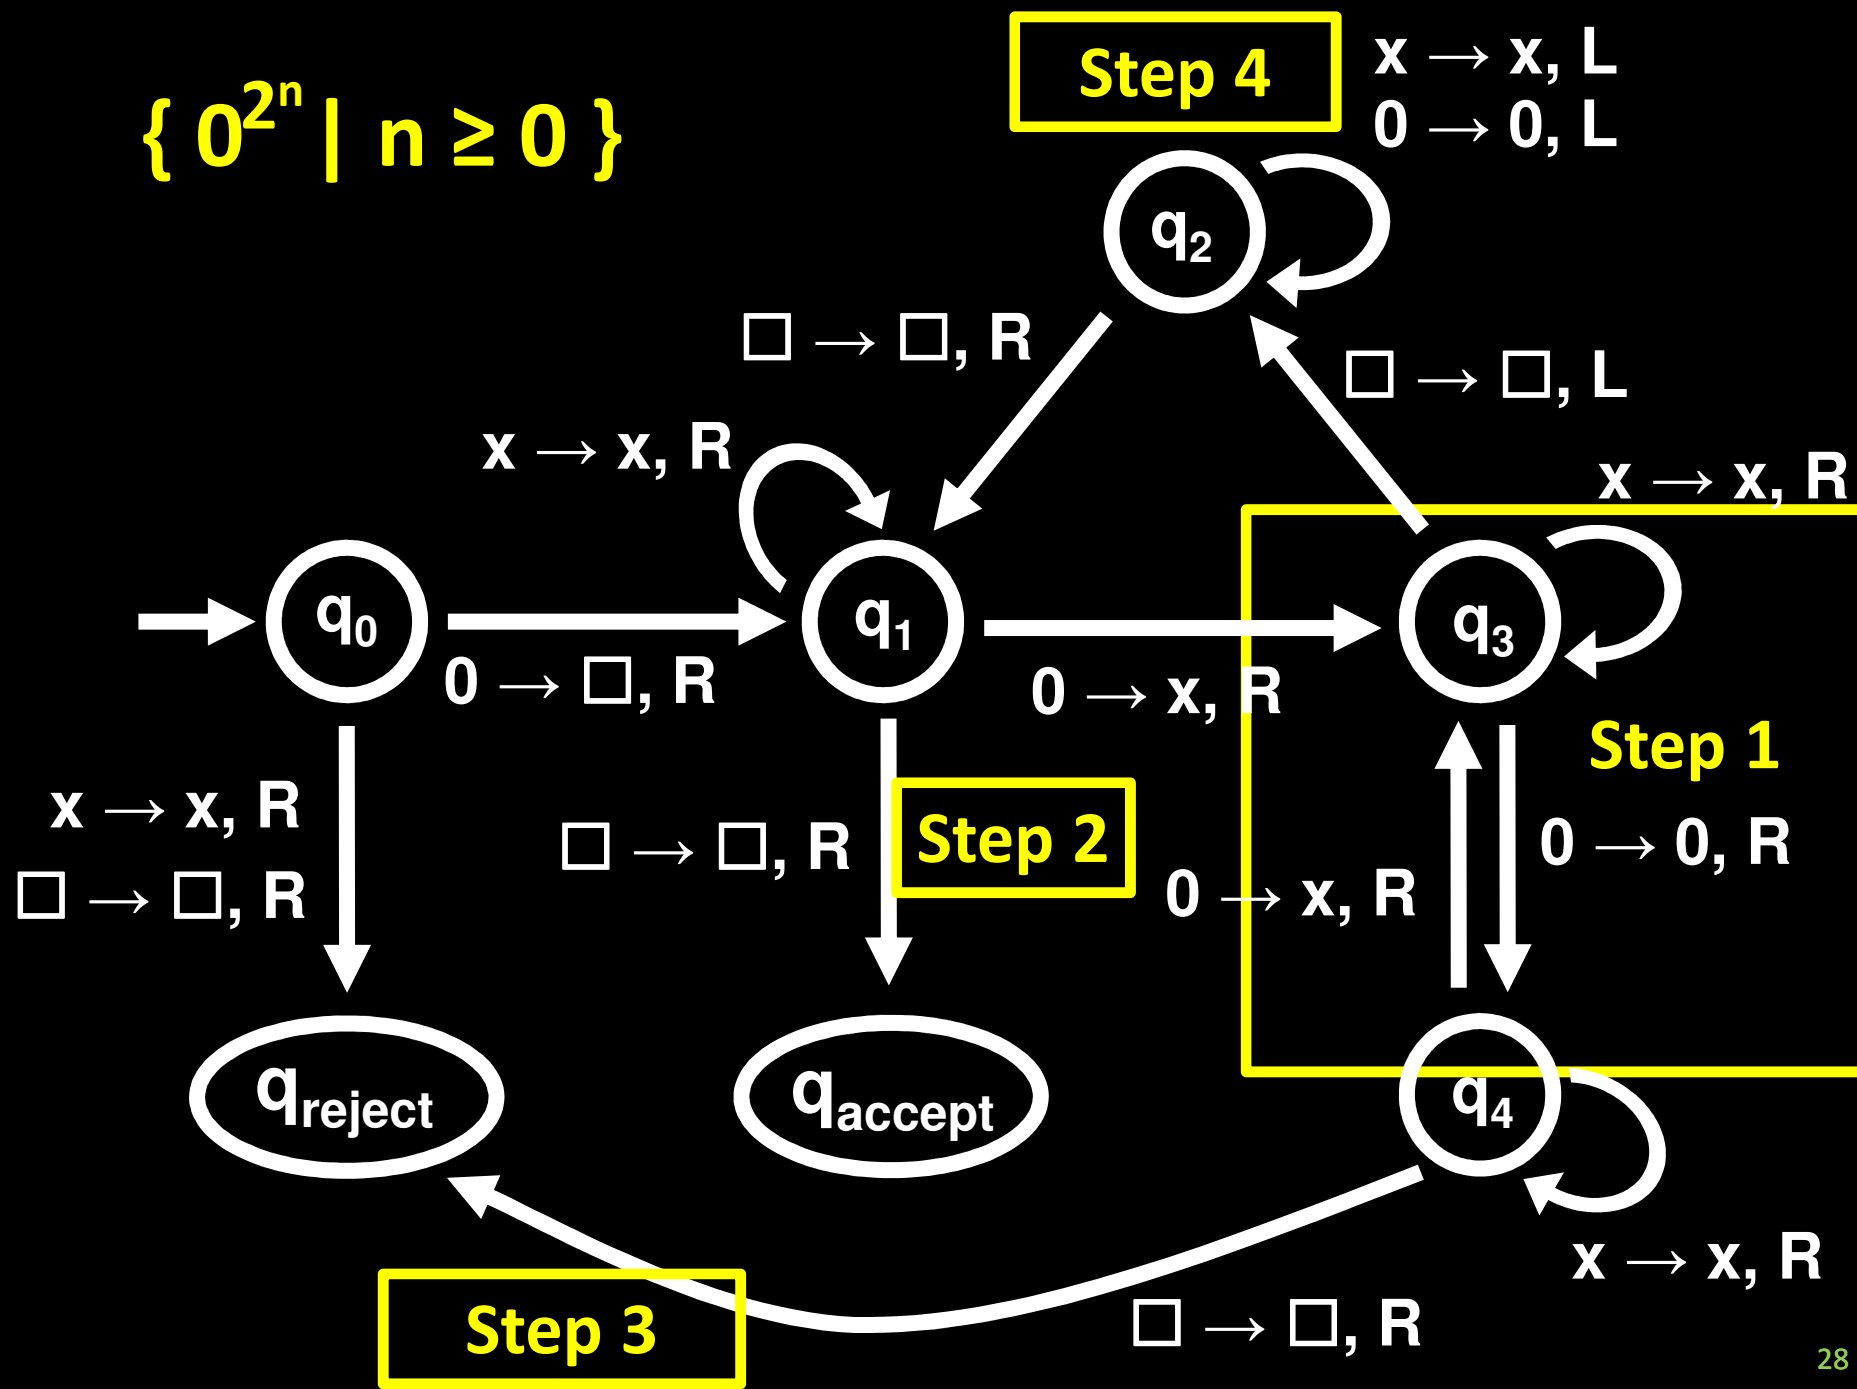
Turing Machine PSEUDOCODE:

1. Sweep from left to right, cross out every other 0
2. If in step 1, the tape had only one 0, *accept*
3. If in step 1, the tape had an odd number of 0's, *reject*
4. Move the head back to the first input symbol.
5. Go to step 1.


Why does this work?

Idea: Every time we return to stage 1, the number of 0's on the tape has been halved.

$\{0^{2^n} \mid n \geq 0\}$

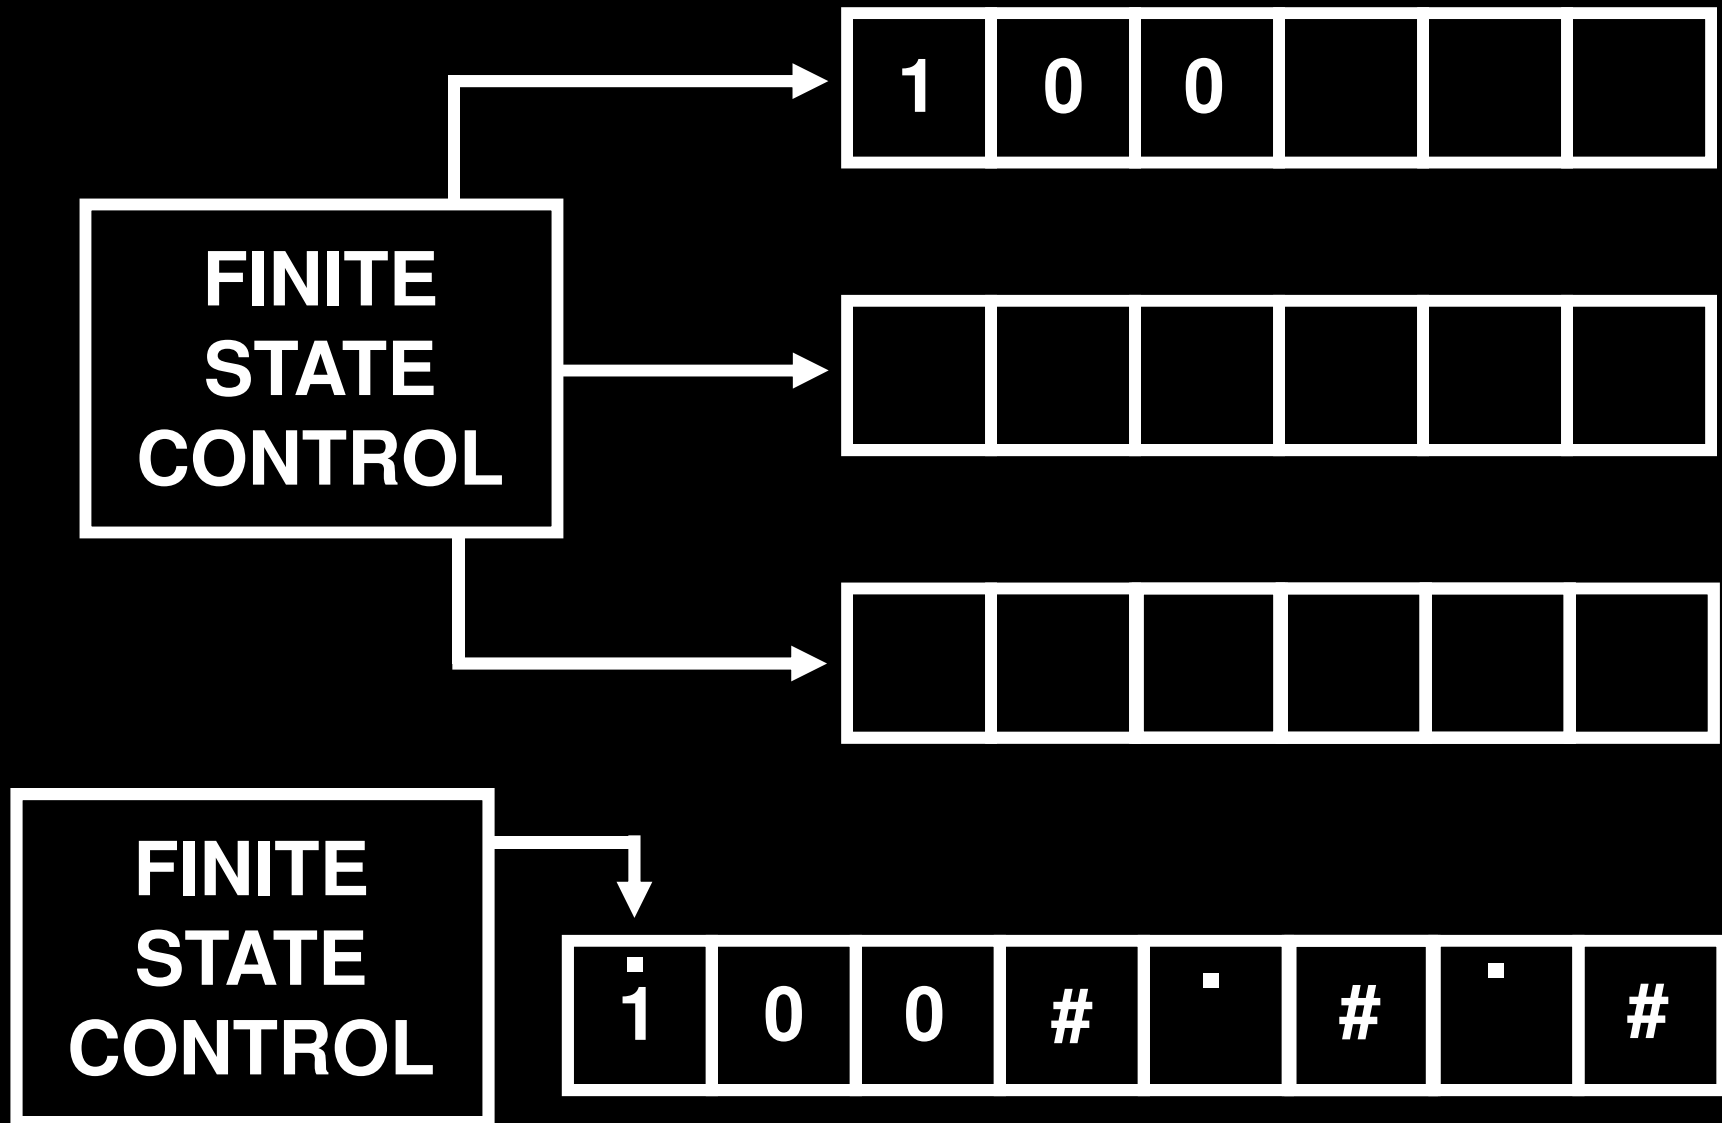


Multitape Turing Machines

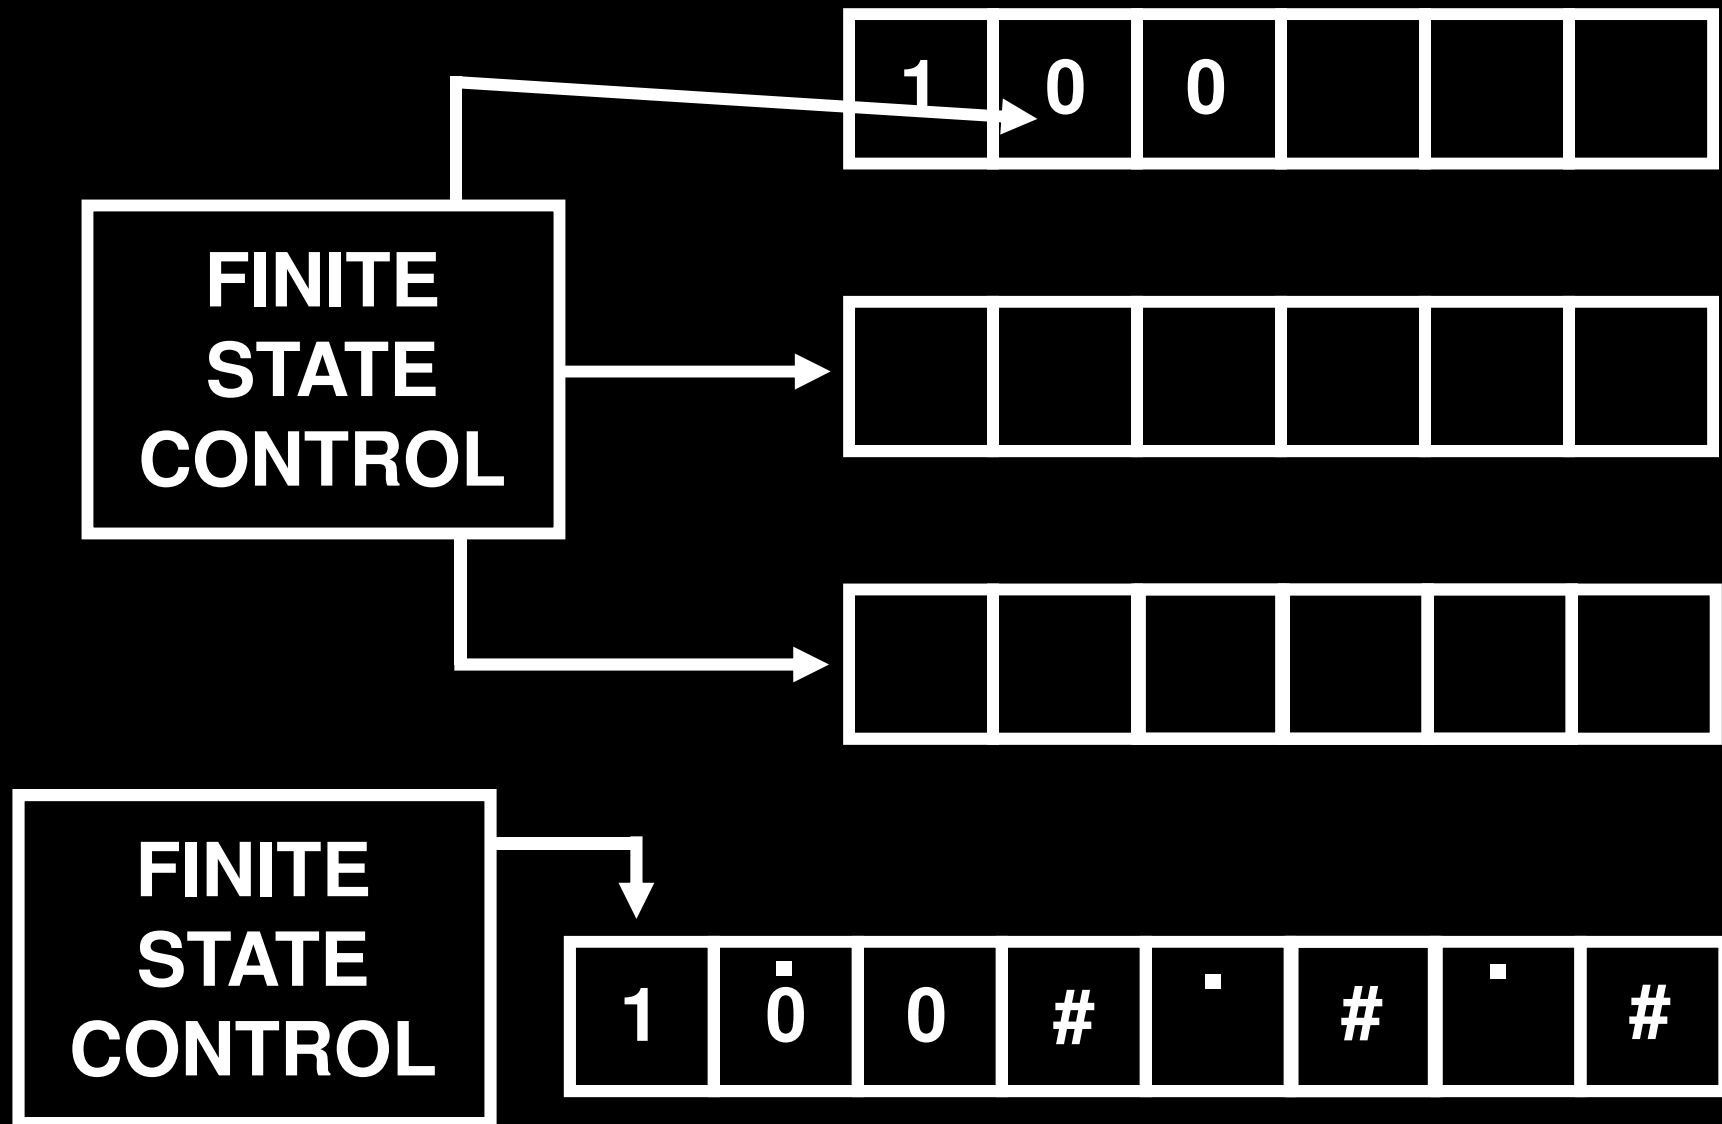


$$\delta : Q \times \Gamma^k \rightarrow Q \times \Gamma^k \times \{L, R\}^k$$

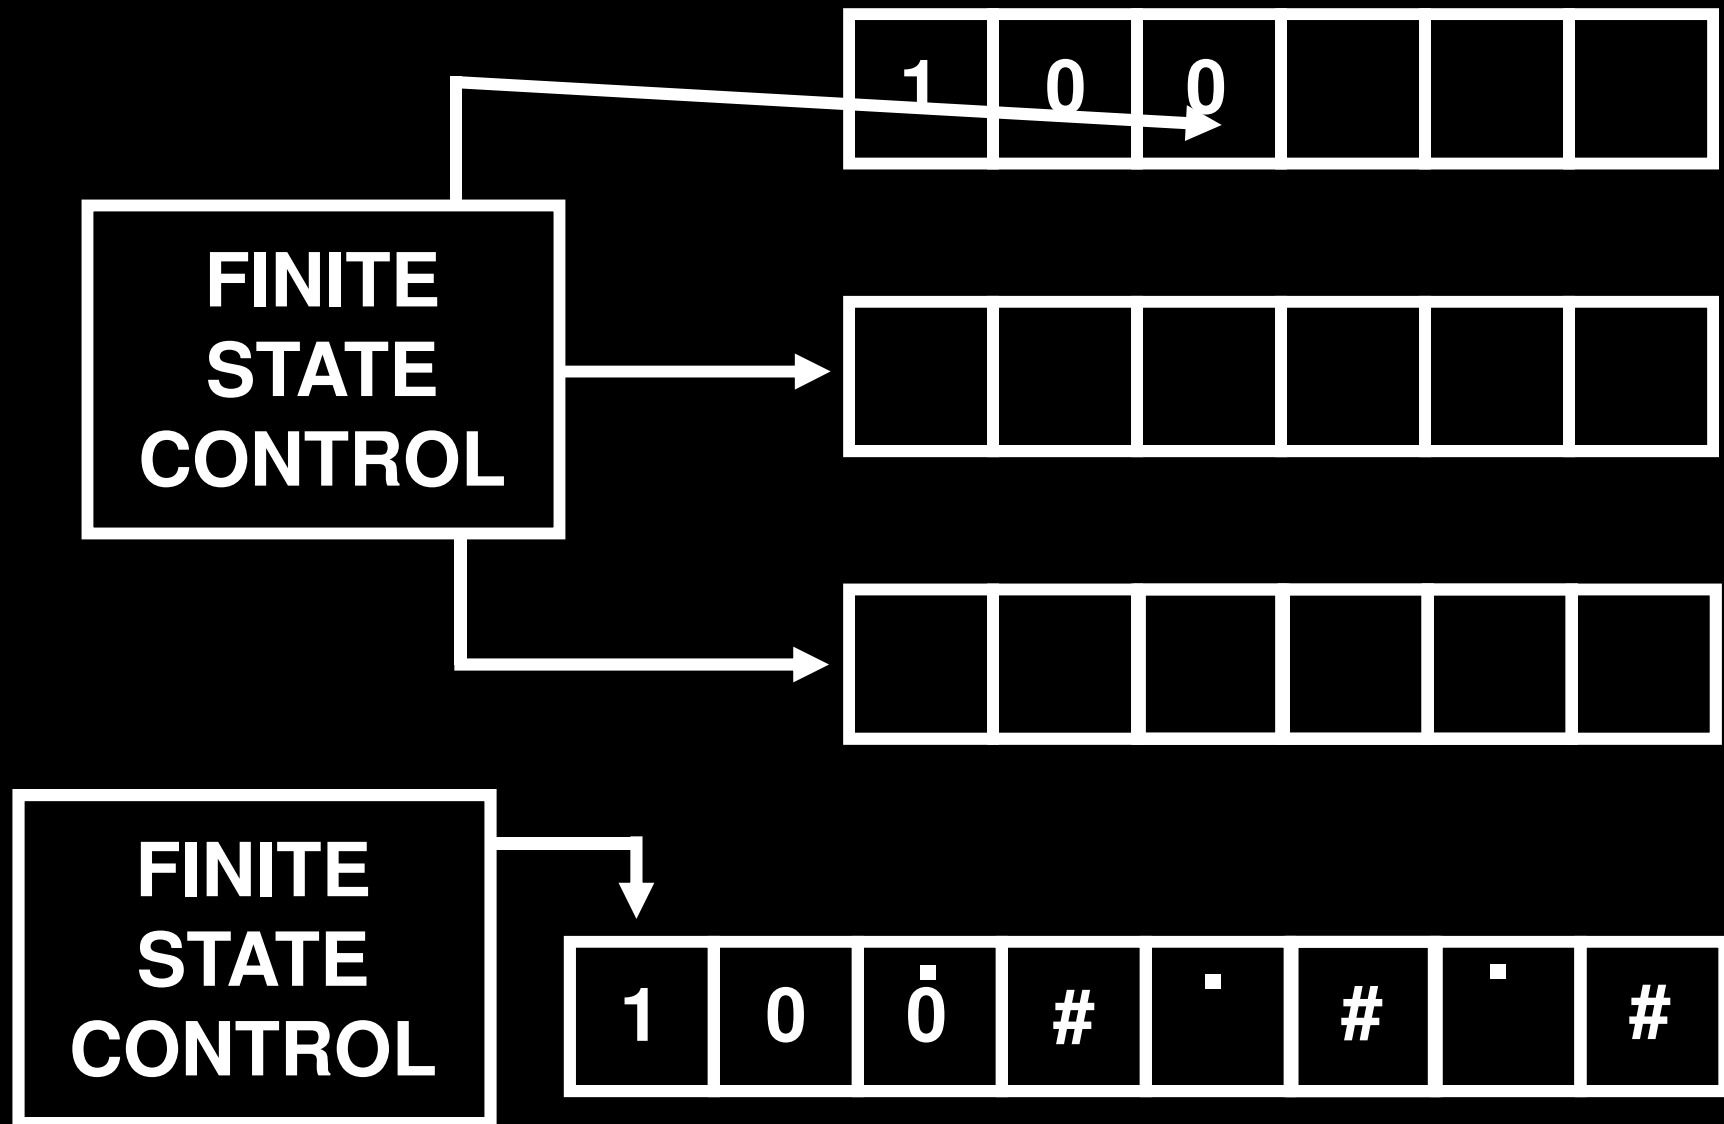
Theorem: Every Multitape Turing Machine can be transformed into a single tape Turing Machine



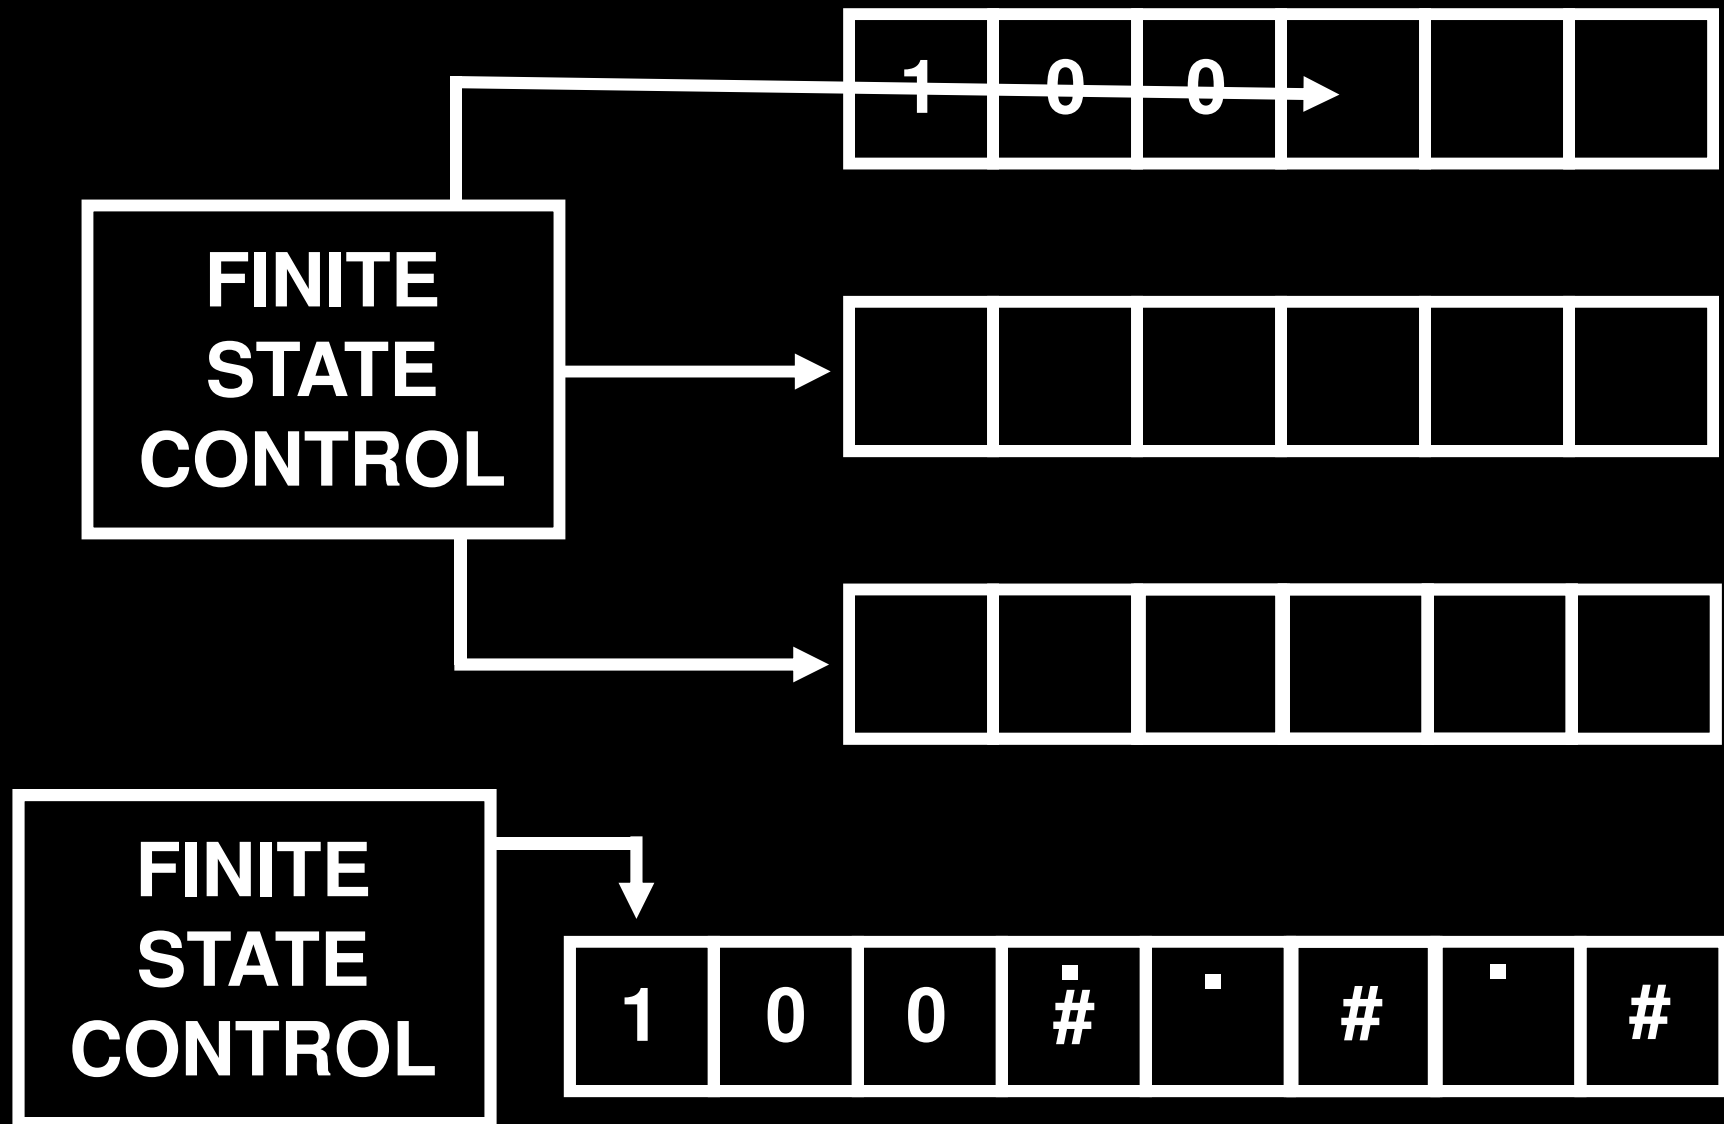
Theorem: Every Multitape Turing Machine can be transformed into a single tape Turing Machine



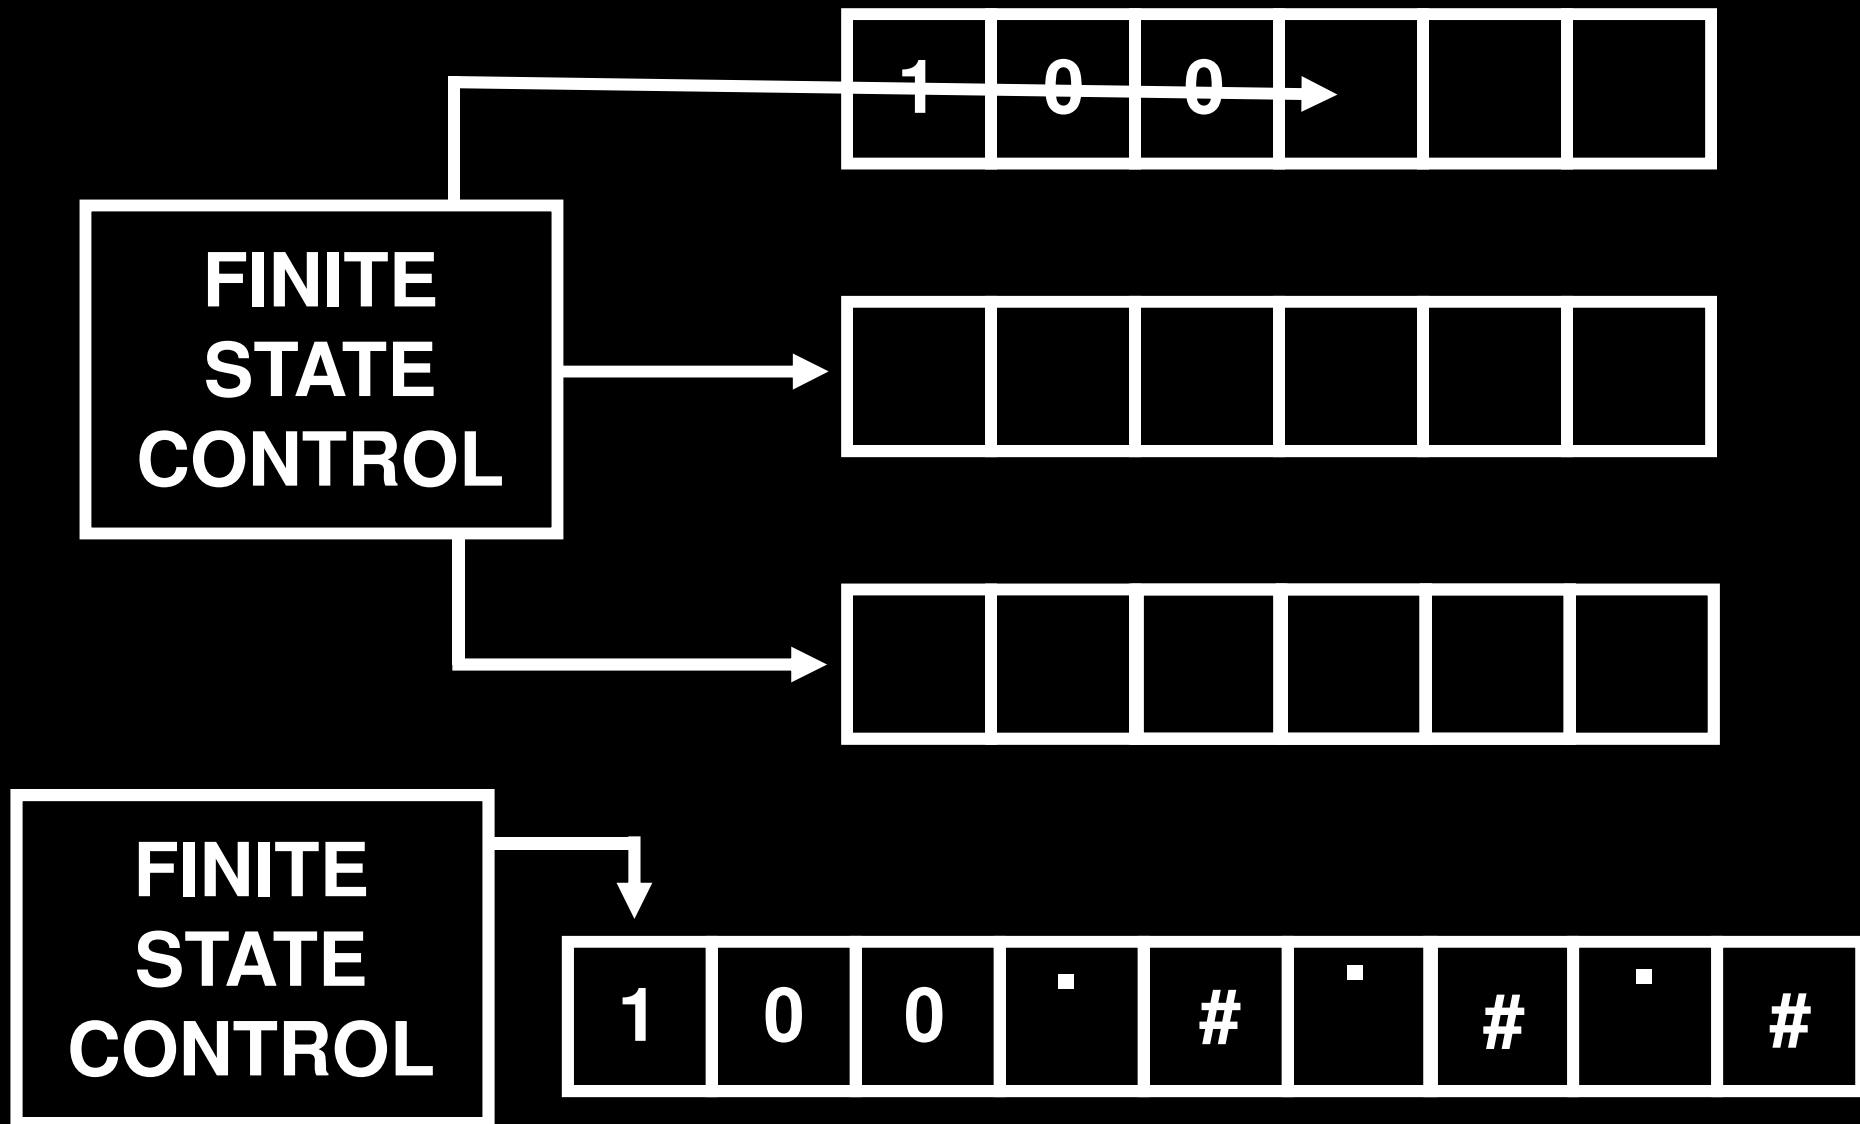
Theorem: Every Multitape Turing Machine can be transformed into a single tape Turing Machine




Theorem: Every Multitape Turing Machine can be transformed into a single tape Turing Machine




Theorem: Every Multitape Turing Machine can be transformed into a single tape Turing Machine





L is **decidable**
(recursive)



L is **recognizable**
(recursively enumerable)

Theorem: L is **decidable**
iff both L and $\neg L$ are **recognizable**