

CS 154

**Oracles, Self-Reference, and the
Foundations of Mathematics**

Next Tuesday (2/17)

Midterm: 12:50pm, Bishop Aud

We'll allow one single-sided page of notes

**Midterm will cover everything up to and including
Tuesday's lecture**

**If you are an SCPD student, contact SCPD for details
about how you will receive your exam**

Rice's Theorem

Suppose L is a language that satisfies two conditions:

1. **(Nontrivial)** There are TMs M_{YES} and M_{NO} ,
where $M_{\text{YES}} \in L$ and $M_{\text{NO}} \notin L$
2. **(Semantic)** For all TMs M_1 and M_2 such that
 $L(M_1) = L(M_2)$, $M_1 \in L$ if and only if $M_2 \in L$

Then, L is undecidable.

A Huge Hammer for Undecidability!



The Regularity Problem for Turing Machines

$\text{REGULAR}_{\text{TM}} = \{ M \mid M \text{ is a TM and } L(M) \text{ is regular} \}$

Given a program, is it equivalent to some DFA?

Theorem: $\text{REGULAR}_{\text{TM}}$ is *not* recognizable

Proof: Use Rice's Theorem!

$\text{REGULAR}_{\text{TM}}$ is nontrivial:

- there's an M_\emptyset which never halts: $M_\emptyset \in \text{REGULAR}_{\text{TM}}$
- there's M' deciding $\{0^n 1^n \mid n \geq 0\}$: $M' \notin \text{REGULAR}_{\text{TM}}$

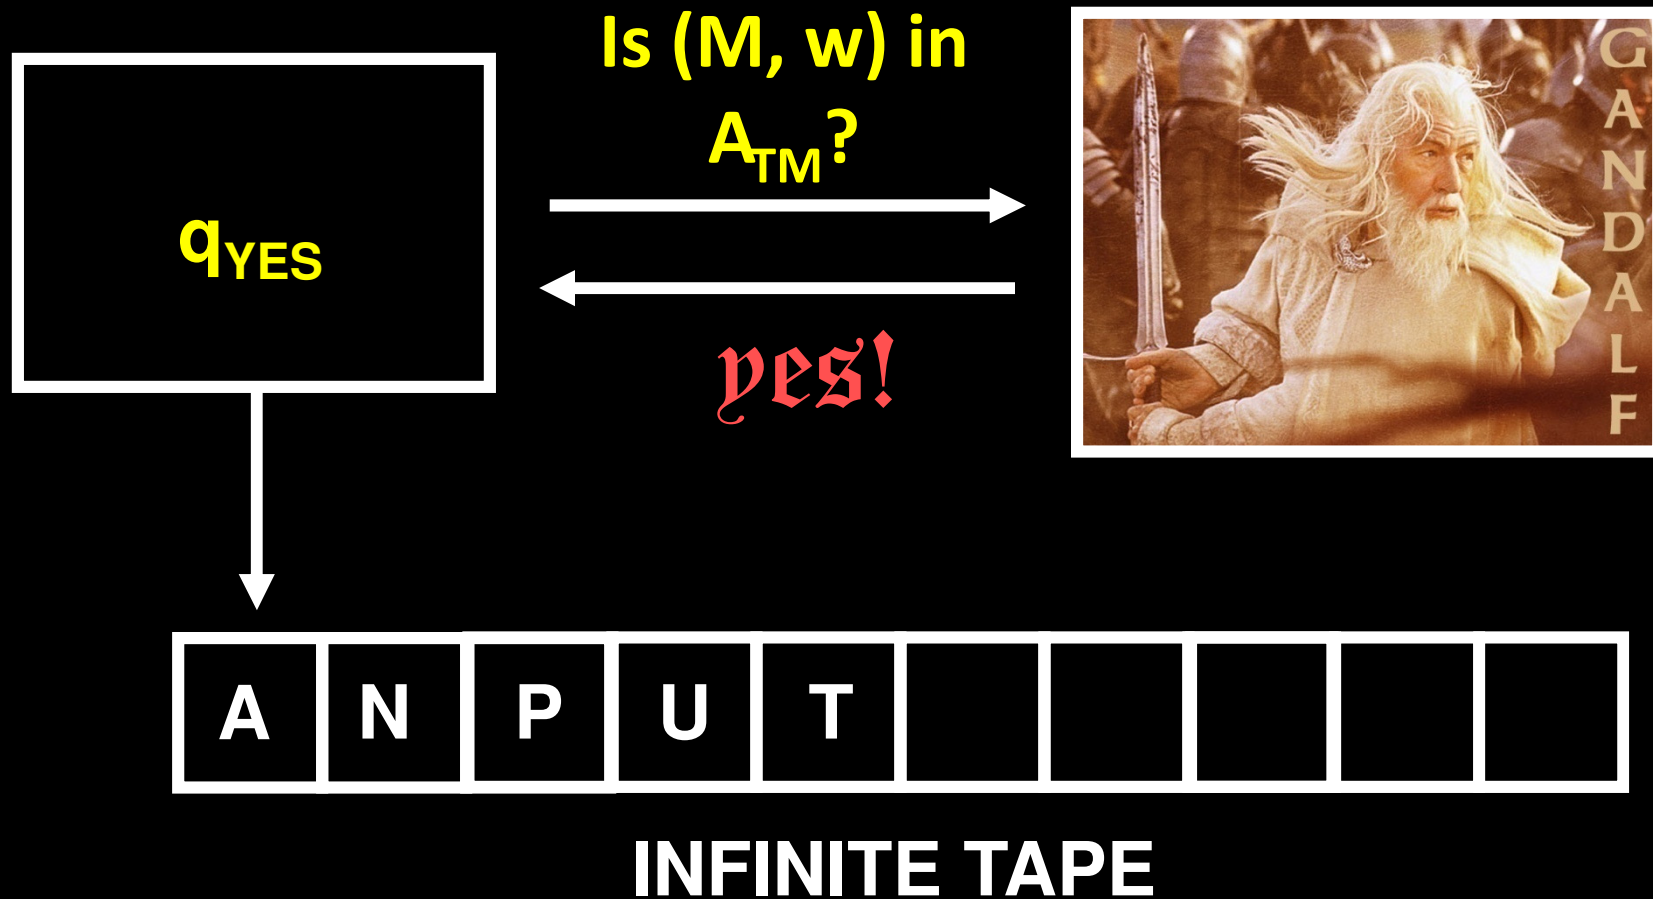
$\text{REGULAR}_{\text{TM}}$ is semantic:

If $L(M) = L(M')$ then $L(M)$ is regular iff $L(M')$ is regular,
therefore $M \in \text{REGULAR}_{\text{TM}}$ iff $M' \in \text{REGULAR}_{\text{TM}}$

By Rice, we have $\neg A_{\text{TM}} \leq_m \text{REGULAR}_{\text{TM}}$

Oracle Turing Machines and Hierarchies of Undecidable Problems

Oracle Turing Machines



Oracle Turing Machines

An **oracle Turing machine** M is equipped with a set $B \subseteq \Gamma^*$ to which a TM M may ask membership queries on a special “oracle tape”

[Formally, M enters a special state $q_?$]

and the TM receives a query answer in one step

[Formally, the transition function on $q_?$ is defined in terms of the *entire oracle tape*:

if the string y written on the oracle tape is in B ,
then state $q_?$ is changed to q_{YES} , otherwise q_{NO}]

This notion makes sense even if B is not decidable!

Definition: A is recognizable with B
if there is an *oracle TM M with oracle B*
that recognizes A

Definition: A is decidable with B
if there is an *oracle TM M with oracle B*
that decides A

Language **A** “Turing-Reduces” to **B**

$$A \leq_T B$$

A_{TM} is decidable with $HALT_{TM}$ ($A_{TM} \leq_T HALT_{TM}$)

On input (M, w) , decide if M accepts w as follows:

If (M, w) is in $HALT_{TM}$ then

run $M(w)$ and output its answer.

else **REJECT**.

HALT_{TM} is decidable with A_{TM} ($\text{HALT}_{\text{TM}} \leq_T A_{\text{TM}}$)

On input (M, w) , decide if M halts on w as follows:

1. If (M, w) is in A_{TM} then **ACCEPT**
2. Else, switch the accept and reject states of M to get a machine M' . If (M', w) is in A_{TM} then **ACCEPT**
3. **REJECT**

\leq_T versus \leq_m

Theorem: If $A \leq_m B$ then $A \leq_T B$

Proof (Sketch):

If $A \leq_m B$ then there is a computable function
 $f : \Sigma^* \rightarrow \Sigma^*$, where for every w ,

$$w \in A \Leftrightarrow f(w) \in B$$

We can simply use one “oracle call” to B to decide A

Theorem: $\neg\text{HALT}_{\text{TM}} \leq_T \text{HALT}_{\text{TM}}$

Theorem: $\neg\text{HALT}_{\text{TM}} \not\leq_m \text{HALT}_{\text{TM}}$

Why?

Limitations on Oracle TMs

The following problem cannot be decided by a TM with an oracle for the Halting Problem:

SUPERHALT = { (M,x) | M, with an oracle for the Halting Problem, halts on x }

Can still use the diagonalization argument here!

Suppose H decides SUPERHALT (with HALT oracle)

Define **D(X) := “if H(X,X) accepts (with HALT oracle) then LOOP, else ACCEPT.”**

Then D(D) halts \Leftrightarrow H(D,D) accepts \Leftrightarrow D(D) loops...

Limits on Oracle TMs

“Theorem” There is an infinite hierarchy
of unsolvable problems!


*Given **ANY** oracle O , there is always a harder problem
that can't be decided with that oracle O*

$\text{SUPERHALT}^0 = \text{HALT} = \{ (M, x) \mid M \text{ halts on } x \}.$

$\text{SUPERHALT}^1 = \{ (M, x) \mid M, \text{ with an oracle for } \text{HALT}_{\text{TM}}, \text{ halts on } x \}$


$\text{SUPERHALT}^n = \{ (M, x) \mid M, \text{ with an oracle for } \text{SUPERHALT}^{n-1}, \text{ halts on } x \}$

Self-Reference and the Recursion Theorem




Lemma: There is a computable function $q : \Sigma^* \rightarrow \Sigma^*$ such that for any string w , $q(w)$ is the *description* of a TM P_w that on every input, prints out w and then accepts

“Proof” Define a TM Q :




Theorem: There is a Self-Printing TM

Proof: First define a TM B:



Now consider the TM:



QED


Another Way of Looking At It

Suppose in general we want to design a program that prints its own description. How?

“Print this sentence.”

Print two copies of the following, the
second copy in quotes: $= B$

“Print two copies of the following, the
second copy in quotes:” $= P_B$



The Recursion Theorem


Theorem: For every TM T computing a function

$$t : \Sigma^* \times \Sigma^* \rightarrow \Sigma^*$$

there is a Turing machine R computing a function


$r : \Sigma^* \rightarrow \Sigma^*$, such that for every string w ,

$$r(w) = t(R, w)$$




Proof: $(a,b) \rightarrow$ T $\rightarrow t(a,b)$

Define M =




Define R:




Proof: $(a,b) \rightarrow$ **T** $\rightarrow t(a,b)$


Define M =



Define R:




Proof: $(a,b) \rightarrow$ T $\rightarrow t(a,b)$



Define R:

$S = Y = R$



For every computable t , there is a computable r such that $r(w) = t(R, w)$ where *R is a description of r*

Suppose we can design a TM T of the form:

*“On input (x, w) , do bla bla bla with x ,
do bla bla bla with w , etc. etc.”*

We can then find a TM R with the behavior:

*“On input w , do bla bla bla with the description R ,
do bla bla bla with w , etc. etc.”*

We can use the operation:

“Obtain your own description”
in Turing machine pseudocode!

Theorem: A_{TM} is undecidable

Proof (using the recursion theorem)

Assume H decides A_{TM}

Construct machine B such that on input w :

1. Obtains its own description B
2. Runs H on (B, w) and flips the output

Running B on input w always does the opposite of what H says it should!

Reminiscent of “free will” paradoxes!

The Fixed-Point Theorem

Theorem: Let $t : \Sigma^* \rightarrow \Sigma^*$ be computable. There is a TM F such that $t(F)$ outputs the description of a TM G such that $L(F)=L(G)$.

Proof: Here is pseudocode for the TM F :

On input w :

1. Obtain the description of F
2. Run $t(F)$ and get an output string G .
Interpret G as the description of a TM
3. Accept w if and only if G accepts w

Computability and the Foundations of Mathematics

The Foundations of Mathematics

A **formal system** describes a formal language for

- writing (finite) mathematical statements,
- has a definition of what statements are “true”
- has a definition of a proof of a statement

Example: Every TM M defines some formal system \mathcal{F}

- **{Mathematical statements in \mathcal{F} } = Σ^***

String w represents the statement “ **M accepts w** ”

- **{True statements in \mathcal{F} } = $L(M)$**
- A **proof** that “ **M accepts w** ” can be defined to be an accepting computation history for M on w

Consistency and Completeness

A formal system \mathcal{F} is **consistent** or **sound** if
no false statement has a valid proof in \mathcal{F}
(Proof in \mathcal{F} implies Truth in \mathcal{F})

A formal system \mathcal{F} is **complete** if
every true statement has a valid proof in \mathcal{F}
(Truth in \mathcal{F} implies Proof in \mathcal{F})

Interesting Formal Systems

Define a formal system \mathcal{F} to be *interesting* if:

1. Any mathematical statement about computation can be described as a statement of \mathcal{F} .

Given (M, w) , there is an $S_{M,w}$ in \mathcal{F} such that $S_{M,w}$ is true in \mathcal{F} if and only if M accepts w .

2. Proofs are “convincing” – a TM can check that a proof of a theorem is correct

This set is decidable: $\{(S, P) \mid P \text{ a proof of } S \text{ in } \mathcal{F}\}$

3. If S is in \mathcal{F} and there is a proof of S describable as a computation, then there’s a proof of S in \mathcal{F} .

If M accepts w , then there is a proof P in \mathcal{F} of $S_{M,w}$

Limitations on Mathematics

For every consistent and interesting \mathcal{F} ,

Theorem 1. (Gödel 1931) \mathcal{F} is *incomplete*:

There are mathematical statements in \mathcal{F} that are *true* but cannot be proved in \mathcal{F} .

Theorem 2. (Gödel 1931) The **consistency** of \mathcal{F} cannot be proved in \mathcal{F} .

Theorem 3. (Church-Turing 1936) The problem of checking whether a given statement in \mathcal{F} has a proof is undecidable.

Unprovable Truths in Mathematics

(Gödel) Every consistent interesting \mathcal{F} is *incomplete*: there are true statements that cannot be proved.

Let $S_{M,w}$ in \mathcal{F} be true if and only if M accepts w

Proof: Define Turing machine $G(x)$:

1. Obtain own description G [Recursion Theorem]
2. Construct statement $S' = \neg S_{G,\epsilon}$
3. Search for a proof of S' in \mathcal{F} over all finite length strings. *Accept* if a proof is found.

Claim: S' is *true in \mathcal{F}* , but has no proof in \mathcal{F}

S' basically says “There is no proof of S' in \mathcal{F} ”

(Gödel 1931) The **consistency of \mathcal{F}** cannot be proved within any interesting consistent \mathcal{F}

Proof: Suppose we can prove “ \mathcal{F} is consistent” in \mathcal{F}

We constructed $\neg S_{G,\varepsilon}$ = “G does not accept ε ”

which we showed is *true*, but *has no proof* in \mathcal{F}

G does not accept $\varepsilon \Leftrightarrow$ There is no proof of $\neg S_{G,\varepsilon}$ in \mathcal{F}

But if there’s a proof in \mathcal{F} of “ \mathcal{F} is consistent” then there’s a proof in \mathcal{F} that $\neg S_{G,\varepsilon}$ is true (here’s the proof):

“If $S_{G,\varepsilon}$ is true, then there is a proof in \mathcal{F} of $\neg S_{G,\varepsilon}$.

\mathcal{F} is consistent, therefore $\neg S_{G,\varepsilon}$ is true.

But $S_{G,\varepsilon}$ and $\neg S_{G,\varepsilon}$ cannot both be true.

Therefore, $\neg S_{G,\varepsilon}$ is true”

This is a contradiction.

Undecidability in Mathematics

$\text{PROVABLE}_{\mathcal{F}} = \{S \mid \text{there's a proof in } \mathcal{F} \text{ of } S, \text{ or} \\ \text{there's a proof in } \mathcal{F} \text{ of } \neg S\}$

(Church-Turing 1936) For every interesting consistent \mathcal{F} , $\text{PROVABLE}_{\mathcal{F}}$ is undecidable

Proof: Suppose $\text{PROVABLE}_{\mathcal{F}}$ is decidable with TM P .

Then we can decide A_{TM} using the following procedure:

On input (M, w) , run the TM P on input $S_{M,w}$

If P accepts, examine all possible proofs in \mathcal{F}

If a proof of $S_{M,w}$ is found then **accept**

If a proof of $\neg S_{M,w}$ is found then **reject**

If P rejects, then **reject**.

Why does this work?

Next Episode:

Your Midterm... Good Luck!