

Memory and C++

Eric Roberts
CS 106B
January 30, 2015

Data Types in C++

The data types that C++ inherits from C:

- Atomic types:
 - **short**, **int**, **long**, and their **unsigned** variants
 - **float**, **double**, and **long double**
 - **char**
 - **bool**
- Enumerated types defined using the **enum** keyword
- Structure types defined using the **struct** keyword
- Arrays of some base type
- Pointers to a target type

Simple Arrays in C++



- We haven't actually used arrays in their low-level form this quarter, because the **Vector** class is so much better.
- From the client perspective, an array is like a brain-damaged form of **Vector** with the following differences:
 - The only operation is selection using `[]`
 - Array selection does not check that the index is in range
 - The length of an array is fixed at the time it is created
 - Arrays don't store their length, so programs that use them must pass an extra integer value that represents the *effective size*
- Array variables are declared using the following syntax:

```
type name[n];
```

where *type* is the element type, *name* is the array name, and *n* is a constant integer expression indicating the length.

A Simple Array Example

```
const int N = 10;  
  
int main() {  
    int array[N];  
    for ( int i = 0 ; i < N ; i++ ) {  
        array[i] = randomInteger(100, 999);  
    }  
    sort(array, N);  
}
```



The Structure of Memory


- The fundamental unit of memory inside a computer is called a **bit**, which is a contraction of the words *binary digit*. A bit can be in either of two states, usually denoted as 0 and 1.
- The hardware structure of a computer combines individual bits into larger units. In most modern architectures, the smallest unit on which the hardware operates is a sequence of eight consecutive bits called a **byte**. The following diagram shows a byte containing a combination of 0s and 1s:

0	0	1	0	1	0	1	0
---	---	---	---	---	---	---	---

- Numbers are stored in still larger units that consist of multiple bytes. The unit that represents the most common integer size on a particular hardware is called a **word**. Because machines have different architectures, the number of bytes in a word may vary from machine to machine.

Binary Notation

- Bytes and words can be used to represent integers of different sizes by interpreting the bits as a number in ***binary notation***.
- Binary notation is similar to decimal notation but uses a different ***base***. Decimal numbers use 10 as their base, which means that each digit counts for ten times as much as the digit to its right. Binary notation uses base 2, which means that each position counts for twice as much, as follows:



Numbers and Bases

- The calculation at the end of the preceding slide makes it clear that the binary representation 00101010 is equivalent to the number 42. When it is important to distinguish the base, the text uses a small subscript, like this:

$$00101010_2 = 42_{10}$$

- Although it is useful to be able to convert a number from one base to another, it is important to remember that the number remains the same. What changes is how you write it down.
- The number 42 is what you get if you count how many stars are in the pattern at the right. The number is the same whether you write it in English as *forty-two*, in decimal as 42, or in binary as 00101010.
- Numbers do not have bases; representations do.



Octal and Hexadecimal Notation

- Because binary notation tends to get rather long, computer scientists often prefer **octal** (base 8) or **hexadecimal** (base 16) notation instead. Octal notation uses eight digits: 0 to 7. Hexadecimal notation uses sixteen digits: 0 to 9, followed by the letters A through F to indicate the values 10 to 15.
- The following diagrams show how the number forty-two appears in both octal and hexadecimal notation:

octal

5	2
---	---

$$\begin{array}{rcl} & \nearrow & 2 \times 1 = 2 \\ & \searrow & 5 \times 8 = 40 \\ & & \hline & & 42 \end{array}$$

hexadecimal

2	A
---	---

$$\begin{array}{rcl} & \nearrow & 10 \times 1 = 10 \\ & \searrow & 2 \times 16 = 32 \\ & & \hline & & 42 \end{array}$$

- The advantage of using either octal or hexadecimal notation is that doing so makes it easy to translate the number back to individual bits because you can convert each digit separately.

Exercises: Number Bases

- What is the decimal value for each of the following numbers?

10001_2
17

177_8
127

AD_{16}
173

- As part of a code to identify the file type, every Java class file begins with the following sixteen bits:

1	1	0	0	1	0	1	0	1	1	1	1	1	1	1	0
---	---	---	---	---	---	---	---	---	---	---	---	---	---	---	---

How would you express that number in hexadecimal


n

1	1	0	0	1	0	1	0	1	1	1	1	1	1	1	0
---	---	---	---	---	---	---	---	---	---	---	---	---	---	---	---

CAFE₁₆


Memory and Addresses

- Every byte inside the primary memory of a machine is identified by a numeric address. The addresses begin at 0 and extend up to the number of bytes in the machine, as shown in the diagram on the right.
- Memory diagrams that show individual bytes are not as useful as those that are organized into words. The revised diagram on the right now includes four bytes in each of the memory cells, which means that the address numbers increase by four each time.
- In these slides, addresses are four-digit hexadecimal numbers, which makes them easy to recognize.
- When you create memory diagrams, you don't know the actual memory addresses at which values are stored, but you do know that everything has an address. ***Just make something up.***



The Allocation of Memory to Variables

- When you declare a variable in a program, C++ allocates space for that variable from one of several memory regions.
- One region of memory is reserved for variables that persist throughout the lifetime of the program, such as constants. This information is called **static data**.
- Each time you call a method, C++ allocates a new block of memory called a **stack frame** to hold its local variables. These stack frames come from a region of memory called the **stack**.
- It is also possible to allocate memory dynamically, as described in Chapter 12. This space comes from a pool of memory called the **heap**.
- In classical architectures, the stack and heap grow toward each other to maximize the available space.



Sizes of the Fundamental Types

- The memory space required to represent a value depends on the type of value. Although the C++ standard actually allows compilers some flexibility, the following sizes are typical:

<u>1 byte</u> <u>(8 bits)</u>	<u>2 bytes</u> <u>(16 bits)</u>	<u>4 bytes</u> <u>(32 bits)</u>	<u>8 bytes</u> <u>(64 bits)</u>	<u>16 bytes</u> <u>(128 bits)</u>
char	short	int	long	long double
bool		float	double	

- Enumerated types are typically assigned the space of an **int**.
- Structure types have a size equal to the sum of their fields.
- Arrays take up the element size times the number of elements.
- Pointers take up the space needed to hold an address, which is 4 bytes on a 32-bit machine and 8 bytes on a 64-bit machine.
- The expression **sizeof(*t*)** returns the size of the type *t*.

Pointers

- In C++, every value is stored somewhere in memory and can therefore be identified with that address. Such addresses are called *pointers*.
- Because C++ is designed to allow programmers to control data at the lowest level, pointers can be manipulated just like any other kind of data. In particular, you can assign one pointer value to another, which means that the two pointers end up indicating the same data value.
- Diagrams that involve pointers are typically represented in two different ways. Using memory addresses emphasizes the fact that pointers are just like integers. Conceptually, it often makes more sense to represent a pointer as an arrow. The head of the arrow is positioned at the address in memory at which the object lives. The tail of the arrow is positioned inside the variable that holds the pointer value.

Declaring a Pointer Variable

- Pointer variables have a declaration syntax that may at first seem confusing. To declare a variable as a pointer to a particular type as opposed to a variable of that type, all you need to do is add a `*` in front of the variable name, like this:

```
type *var;
```

- For example, if you wanted to declare a variable **px** to be a pointer to a **double** value, you could do so as follows:

```
double *px;
```

Similarly, to declare a variable **pptr** as a pointer to a **Point** structure, you would write:

```
Point *pptr;
```

Pointer Operators

- C++ includes two built-in operators for working with pointers:
 - The address-of operator (&) is written before a variable name (or any expression to which you could assign a value) and returns the address of that variable. Thus, the expression **&total** gives the address of **total** in memory.
 - The deference operator (*) is written before a pointer expression and returns the actual value to which the pointer points.
- Suppose, for example, that you have declared and initialized the following variables:

```
double x = 2.5;  
double *px = &x;
```

- At this point, the variable **px** points to the variable **x**, and the expression ***px** is synonymous with the variable **x**.

Pointers and Arrays

- In C++, all arrays are represented internally as a pointer to their first element.
- For example, if you declare an array



```
int list[100];
```

the C++ compiler treats the name **list** as a shorthand for the expression **&list[0]**.

- You can freely intermix array and pointer notation in your code. If you declare something as an array of a particular type, you can use it as a pointer to that type, and vice versa.
- When you pass an array to a function, only the *address* of the array is copied into the parameter. This strategy has the effect of *sharing* the elements of the array between the function and its caller.

Arrays Are Passed as Pointers

```
void sort(int array[], int n) {  
    for ( int lh = 0 ; lh < n ; lh++ ) {  
        int rh = findSmallest(array, lh, n - 1);  
        swap(array[lh], array[rh]);  
    }  
}
```



Pointer Arithmetic

- Like C before it, C++ defines the $+$ and $-$ operators so that they work with pointers.
- Suppose, for example, that you have made the following declarations:


```
double array[5];  
double *dp = array;
```

How do those variables appear in memory?

- C++ defines pointer addition so that the following identity always holds:

$dp + i \equiv \&array[i]$

Thus, **$dp + 2$** points to **$array[2]$** .




C Strings are Pointers to Characters

- As you know from Chapter 3, C++ supports the old C style of strings, which is simply a pointer to a character, which is the first element of a character array terminated by the ***null character*** (`'\0'`).
- Given this definition, what does the declaration

```
char *msg = "hello, world";
```


generate in memory?
- You can still select characters in `msg` by their index because of the equivalence of arrays and pointers.



Exercises: C String Functions

- Implement the C library method **strlen(cstr)** that returns the length of the C string **cstr**.
- Implement the C library function **strcpy(dst, src)**, which copies the characters from the string **src** into the character array indicated by **dst**. For example, the code on the left should generate the memory state on the right:

```
char *msg = "hello, world";  
char buffer[16];  
strcpy(buffer, msg);
```



The Hot-Shot Solution

```
void strcpy(char *dst, char *src) {  
    while (*dst++ = *src++);  
}
```



- The pointer expression

***p++**

is one of the most common idioms in C and C++, and it is essential to understand what it means. It is, however, equally important that you avoid using it in your own code.

- When you work with C++, understanding the ***p++** idiom is important primarily because the same syntax comes up in STL iterators, which are used everywhere in professional code.

The Internet Worm

"All the News
That's Fit to Print"

The New York Times

Late Edition
New York Today, partly sunny, milder.
High 55-54. Tonight, mostly cloudy.
Low 48-54. Tomorrow, cloudy, windy,
rain developing. High 53-62. Yesterday:
High 55, low 41. Details, page D16.

VOL. CXXXVIII... No. 47,679 Copyright © 1989 The New York Times NEW YORK, FRIDAY, NOVEMBER 4, 1989 34 cents beyond 15 miles from New York City, except on Long Island. 35 CENTS

Gov. Michael S. Dukakis having his picture taken by a 10-year-old fan at a town meeting in Fairlee, Vt., Pa., during a tour of the Northeast in which he emphasized the drug problem. Page A10. Vice President Bush addressed supporters at a rally in Columbus, Ohio. Later than a week after Mr. Dukakis acknowledged being a liberal, Mr. Bush said yesterday that "this election is not about labels." Page A18.

**Registration Off
Since 1984 Vote**

There has been a pronounced decline in the percentage of eligible Americans who are registered to vote, a research group reports.

Nationally, the percentage of eligible Americans who are registered is estimated to be 78.3 percent, down 1.2 points from the 1984 level.

The group's study concluded that in many of the 50 states where final figures are available the decline was among

'Virus' in Military Computers Disrupts Systems Nationwide

By JOHN MARKOFF

In an address that raised questions about the vulnerability of the nation's computers, a Department of Defense network has been disrupted twice Wednesday by a rapidly spreading "virus" program apparently introduced by a computer science student.

The program reproduced itself through the computer network, making hundreds of copies in each machine it reached, effectively clogging systems linking thousands of military, corporate and university computers around the nation and preventing them from doing additional work. The virus is thought not to have destroyed any files.

By late yesterday afternoon computer experts were calling the virus the largest assault ever on the nation's computers.

Military officials, researchers and corporations.

While some sensitive military data are involved, the computers handling the nation's most sensitive secret information, that on the control of nuclear weapons, are thought not to have been touched by the virus.

Parallel to Biological Virus

Computer viruses are so named because they parallel in the computer world the behavior of biological viruses. A virus is a program, or a set of instructions in a computer, that is either placed on a floppy disk meant to be used with the computer or introduced when the computer is communicating over telephone lines or data networks with other computers.

The programs can copy themselves into the computer's memory software, or operating system, usually without calling any attention to themselves. From there, the program can be passed to additional computers.

Depending upon the intent of the software's creator, the program might cause a provocative but otherwise harmless message to appear on the computer's screen. Or it could systematically destroy data in the computer's memory. In this case, the virus program did nothing more than reproduce itself rapidly.

The program was apparently a result of an experiment, which

Continued on Page A21, Column 2

PENTAGON REPORTS IMPROPER CHARGES FOR CONSULTANTS

CONTRACTORS CRITICIZED

Inquiry Shows Routine Billing of Government by Industry on Fees, Some Dubious

By JOHN H. CUSHMAN Jr.
Special to the New York Times

WASHINGTON, Nov. 3 — A Pentagon investigation has found that the nation's largest military contractors routinely charge the Defense Department for hundreds of millions of dollars paid to consultants, often without justification.


The report of the investigation said that neither the military's current rules nor the contractors' own policies are adequate to assure that the Government does not improperly pay for privately arranged consulting work. Senior Defense Department officials said the Pentagon was proposing changes to correct the flaws.

While it is not improper for military contractors to use consultants in performing work for the Pentagon, the work must directly benefit the military if it is to be paid for by the Defense Department. Often, Pentagon investigators discovered, this cost is not met.

Reader Look at Consultants

The Justice Department's continuing criminal investigation has focused attention on consultants and their role in the designing and selling of weapons, and the Defense Department has been criticized for using consultants too freely. Now the Pentagon's own inves-


Robert Morris Jr.



How the Morris Worm Worked

If the user, however, enters a name string that overflows the buffer, the bytes in that name will overwrite the data on the stack.

Now when the function returns, it will jump into the code written as part of the name, thereby executing the worm's instructions.



The End