

CS110: API Programming Against the UNIX Filesystem

▪ Learning filesystem API calls

- Next lecture, we'll be talking about how components of a UNIX filesystem are actually implemented. That alone should be motivation for learning how to code against a collection of APIs that manipulate files, directories, and so forth.
- You'll benefit by understanding how clients interact with the file system through **system calls**, which are a collection of kernel-resident functions that user programs must go through in order to manipulate the hardware. Requests to open a file, read from a file, to extend the heap, etc, all eventually go through system calls, which are the only functions that can be trusted to touch the system.
- Today's lecture examples reside in **/usr/class/cs110/lecture-examples/spring-2015/filesystems**.
- The **/usr/class/cs110/lecture-examples/spring-2015** directory is a mercurial repository that will be updated with additional examples as the quarter progresses.
 - To get starter, type **hg clone /usr/class/cs110/lecture-examples/spring-2015 cs110-lecture-examples** at the command prompt to create a local copy of the master.
 - Each time I mention there are new examples, navigate into your local copy and type **hg pull; hg update**.
- More importantly, read [Sections 1 through 5](#) of the Saltzer & Kaashoek online textbook, paying special attention to the details in Section 5, which will help you with your first assignment (which goes out on Friday).

Filesystem API: Using `stat` and `lstat`

■ `stat` and `lstat`

- `stat` is a function that populates a `struct stat` with information about some named file (regular files, directories, links).
- `stat` and `lstat` operate exactly the same way, except when the named file is a link, `stat` returns information about the file the link references, and `lstat` returns information about the link itself.
- Manual (`man`) pages exist for both of these functions (e.g. `man 2 stat`, `man 2 lstat`, etc.)
- The `struct stat` contain the following fields ([source](#))

<code>dev_t</code>	<code>st_dev</code>	ID of device containing file
<code>ino_t</code>	<code>st_ino</code>	file serial number
<code>mode_t</code>	<code>st_mode</code>	mode of file
<code>nlink_t</code>	<code>st_nlink</code>	number of links to the file
<code>uid_t</code>	<code>st_uid</code>	user ID of file
<code>gid_t</code>	<code>st_gid</code>	group ID of file
<code>dev_t</code>	<code>st_rdev</code>	device ID (if file is character or block special)
<code>off_t</code>	<code>st_size</code>	file size in bytes (if file is a regular file)
<code>time_t</code>	<code>st_atime</code>	time of last access
<code>time_t</code>	<code>st_mtime</code>	time of last data modification
<code>time_t</code>	<code>st_ctime</code>	time of last status change
<code>blksize_t</code>	<code>st_blksize</code>	a filesystem-specific preferred I/O block size for this object. In some filesystem types, this may vary from file to file
<code>blkcnt_t</code>	<code>st_blocks</code>	number of blocks allocated for this object

- The `st_mode` field isn't so much a single value as it is a collection of bits encoding multiple pieces of information about file type and permissions.
- A collection of bit masks and macros can be used to extract information from the `st_mode` field.
- The next two examples illustrate how the `stat` and `lstat` functions can be used to navigate and otherwise manipulate files within the file system.

Filesystem API: search

■ Implementation of **search**

- **search** is our own simplified implementation of the **find** built-in.
- The following **main** relies on **listMatches**, which we'll implement a little later. (The full program of interest is online [right here](#).)

```
static void exitUnless(bool test, FILE *stream, int code, const char *control, ...) {
    if (test) return;
    va_list arglist;
    va_start(arglist, control);
    vfprintf(stream, control, arglist);
    va_end(arglist);
    exit(code);
}

int main(int argc, char *argv[]) {
    exitUnless(argc == 3, stderr, kWrongArgumentCount,
               "Usage: %s <directory> <pattern>\n", argv[0]);

    struct stat st;
    const char *directory = argv[1];
    stat(directory, &st);
    exitUnless(S_ISDIR(st.st_mode), stderr, kDirectoryNeeded,
               "<directory> must be an actual directory, %s is not", directory);

    size_t length = strlen(directory);
    if (length > kMaxPath) return 0;

    const char *pattern = argv[2];
    char path[kMaxPath + 1];
    strcpy(path, directory); // no buffer overflow because of above check
    listMatches(path, length, pattern);
    return 0;
}
```

Filesystem API: search (continued)

■ Implementation details

- Don't worry about the **va_list** trickery unless you're curious. I just wanted to unify error checking to a single helper function. My **exitUnless** is basically a sexier version of the **assert** macro.
- The first thing that's new to us is the call to **stat**, which lifts a bunch of information about the named file off of the file system and populates **st** with it.
- You'll also note the use of the **S_ISDIR** macro, which examines the upper four bits of the **st_mode** field to determine whether the named file is a directory (or a link to one).
- **S_ISDIR** has a few cousins: **S_ISREG** decides whether a file is a regular file, and **S_ISLNK** decided whether the file is a link. (We'll use all of these in our next example).
- Most of what's interesting is managed by the **listMatches** function, which does a depth-first traversal of the file system to see what files just happen to contain a named **pattern** as a substring.

Filesystem API: search (continued)

■ Implementation of `listMatches`

```
static void listMatches(char path[], size_t length, const char *pattern) {
    DIR *dir = opendir(path);
    if (dir == NULL) return; // path isn't a directory
    strcpy(path + length, "/");
    while (true) {
        struct dirent *de = readdir(dir);
        if (de == NULL) break;
        if (strcmp(de->d_name, ".") == 0 || strcmp(de->d_name, "..") == 0) continue;
        if (length + strlen(de->d_name) + 1 > kMaxPath) continue;
        strcpy(path + length + 1, de->d_name);
        struct stat st;
        lstat(path, &st);
        if (S_ISREG(st.st_mode)) {
            if (strstr(de->d_name, pattern) != NULL) {
                printf("%s\n", path);
            }
        } else if (S_ISDIR(st.st_mode)) {
            listMatches(path, length + 1 + strlen(de->d_name), pattern);
        }
    }
    closedir(dir);
}
```

Implementation of `listMatches` (continued)

■ Implementation details

- My implementation relies on `opendir`, which accepts what is presumably a directory. It returns a pointer to an opaque iterable that surfaces a sequence of `struct dirent`s via a series of `readdir` calls.
- If `opendir`'s parameter is something other than a directory, it'll return `NULL`.
- When the `DIR` has surfaced all of its entries, `readdir` returns `NULL`. A return value of `NULL` says it's all over.
- The `struct dirent` is only *guaranteed* to contain a `d_name` field, which is a C string expression of the directory entry's name. `.` and `..` are among the sequence of named entries, but I ignore them so I don't cycle through any single directory more than once.
- I use `lstat` instead of `stat` so I know whether an entry is really a link.
- If the status clearly identifies an entry as a regular file, then I print the entire path if and only if it contains the `pattern` of interest.
- If the status identifies an entry to be a directory, then I recursively descend into it to see if any of its named entries match the pattern we're looking for.
- `opendir` returns access to a record that eventually must be released via a call to `closedir`. That's why my implementation ends with it.

Filesystem API: **list**

▪ **list** implementation details

- I also present the implementation of **list**, which emulates the functionality of **ls** (in particular, **ls -lUa**).
- The implementation of **list** and **search** have many things in common, but the implementation of **list** is much longer.
- Full implementation of entire **list** executable is [right here](#).
- Sample output (notice this is my own **list**, not **ls**!):

```
myth22:/usr/class/cs110/staff/lectures/filesystem> list /usr/
lrwxrwxrwx 1 root root 26 Sep 19 2012 newsw -> /afs/ir/systems/@sys/newsw
drwxr-xr-x 2 root root 4096 Oct 10 2012 games
lrwxrwxrwx 1 root root 26 Sep 19 2012 pubsw -> /afs/ir/systems/@sys/pubsw
lrwxrwxrwx 1 root root 26 Sep 19 2012 sweet -> /afs/ir/systems/@sys/sweet
drwxr-xr-x 5 root root 4096 Sep 21 2012 lib32
drwxr-xr-x 2 root root 118784 Apr 01 16:28 bin
drwxr-xr-x 2 root root 12288 Mar 21 16:56 sbin
drwxr-xr-x >9 root root 4096 Oct 17 2012 ..
drwxr-xr-x >9 root root 4096 Sep 21 2012 local
drwxr-xr-x >9 root root 73728 Mar 21 16:56 lib
lrwxrwxrwx 1 root root 26 Sep 19 2012 class -> /afs/ir.stanford.edu/class
drwxr-xr-x >9 root root 12288 Mar 13 22:28 include
drwxr-xr-x >9 root root 16384 Oct 10 2012 share
drwxr-xr-x >9 root root 4096 Sep 19 2012 .
drwxr-xr-x 8 root root 4096 Sep 21 2012 src
```

- I don't present the entire implementation. I just show one key function: the one that knows how to print out the permissions information for an arbitrary entry.

Filesystem API: `list`

■ Implementation of `list`'s `listPermissions`:

```
static inline void updatePermissionsBit(bool flag, char permissions[],
                                       size_t column, char ch) {
    if (!flag) return;
    permissions[column] = ch;
}

static const size_t kNumPermissionColumns = 10;
static const char kPermissionChars[] = {'r', 'w', 'x'};
static const size_t kNumPermissionChars = sizeof(kPermissionChars);
static const mode_t kPermissionFlags[] = {
    S_IRUSR, S_IWUSR, S_IXUSR, // user flags
    S_IRGRP, S_IWGRP, S_IXGRP, // group flags
    S_IROTH, S_IWOTH, S_IXOTH, // everyone (other) flags
};
static const size_t kNumPermissionFlags = sizeof(kPermissionFlags)/sizeof(kPermissionFlags[0]);

static void listPermissions(mode_t mode) {
    char permissions[kNumPermissionColumns + 1];
    memset(permissions, '-', sizeof(permissions));
    permissions[kNumPermissionColumns] = '\0';
    updatePermissionsBit(S_ISDIR(mode), permissions, 0, 'd');
    updatePermissionsBit(S_ISLNK(mode), permissions, 0, 'l');
    for (size_t i = 0; i < kNumPermissionFlags; i++) {
        updatePermissionsBit(mode & kPermissionFlags[i], permissions, i + 1,
                             kPermissionChars[i % kNumPermissionChars]);
    }
    printf("%s ", permissions);
}
```

- Full implementation of `list` is in `list.c`.

Filesystem API: **copy**

■ Implementation of **copy**

- The implementation of **copy** (designed to mimic the behavior of **cp**) illustrates how to use **open**, **read**, **write**, **close**, **stat**. It also introduces the notion of a file descriptor.
- **man** pages exist for all of these functions (e.g. **man 2 open**, **man 2 read**, etc.)
- Full implementation of our own **copy** executable is [right here](#).

■ Pros and cons of file descriptors over **FILE** pointers and C++ **iostreams**

- The file descriptor abstraction provides direct, low level access to a stream of data without the fuss of data structures or objects. It certainly can't be slower, and depending on what you're doing, it may even be faster.
- **FILE** pointers and C++ **iostreams** work well when you know you're layering over standard output, standard input, and local files. They are less useful when the stream of bytes is associated with a network connection. (**FILE** pointers and C++ **iostreams** assume they can rewind and move the file pointer back and forth freely, but that's not the case with file descriptors associated with network connections).
- File descriptors, however, work with **read** and **write** and nothing else. C **FILE** pointers and C++ streams provide automatic buffering and more elaborate formatting options.

Filesystem API (continued)

■ Implementation of `copy`

```
int main(int argc, char *argv[]) {
    if (argc != 3) {
        fprintf(stderr, "%s <source-file> <destination-file>.\n", argv[0]);
        return kWrongArgumentCount;
    }

    int fdin = open(argv[1], O_RDONLY);
    if (fdin == -1) {
        fprintf(stderr, "%s: source file could not be opened.\n", argv[1]);
        return kSourceFileNonExistent;
    }

    struct stat st;
    stat(argv[1], &st);
    int fdout = open(argv[2],
        /* flags = */ O_WRONLY | O_CREAT | O_EXCL,
        /* mode = */ st.st_mode & S_IRWXU);
    if (fdout == -1) {
        switch (errno) {
            case EEXIST:
                fprintf(stderr, "%s: destination file already exists.\n", argv[2]);
                break;
            default:
                fprintf(stderr, "%s: destination file could not be created.\n", argv[2]);
                break;
        }
    }

    return kDestinationFileOpenFailure;
}
```

Filesystem API (continued)

■ Implementation of **copy**, continued

```
char buffer[1024];
while (true) {
    ssize_t bytesRead = read(fdin, buffer, sizeof(buffer));
    if (bytesRead == 0) break;
    if (bytesRead == -1) {
        fprintf(stderr, "%s: lost access to file while reading.\n", argv[1]);
        return kReadFailure;
    }

    size_t bytesWritten = 0;
    while (bytesWritten < bytesRead) {
        ssize_t count = write(fdout, buffer + bytesWritten, bytesRead - bytesWritten);
        if (count == -1) {
            fprintf(stderr, "%s: lost access to file while writing.\n", argv[2]);
            return kWriteFailure;
        }
        bytesWritten += count;
    }
}

if (close(fdin) == -1) fprintf(stderr, "%s: had trouble closing file.\n", argv[1]);
if (close(fdout) == -1) fprintf(stderr, "%s: had trouble closing file.\n", argv[2]);
return 0;
}
```

- Full implementation is in **copy.c** file.

Filesystem API (continued)

- We're equipped to understand another system call that knows how to duplicate a file descriptor: **dup2**.
- Here's the prototype:

```
int dup2(int oldfd, int newfd);
```

- **dup2** attaches **newfd** to the same file **oldfd** is attached to, closing **newfd** first if necessary (but not closing **oldfd**).
- **dup2** is often used to support pipes and redirections, and it's often used to simplify program logic when characters can be either drawn from standard input or from a file.
- Here's an example of a program—I call it **filedump**—that emulates the **cat** executable that comes with most UNIX/Linux distributions.
 - Our version copies all data from either standard input (if the program is invoked without any arguments) or the contents of a regular file (if the program is invoked with a single extra argument, assumed to be the name of a readable, regular file) to standard output until the relevant **read** call returns 0.
 - I contrive a requirement that all bytes must be drawn from file descriptor 0 (aka **STDIN_FILENO**), thereby requiring the conditional use of **dup2**.
 - Full implementation is online [right here](#).

```
int main(int argc, char *argv[]) {
    // error checking omitted
    if (argc == 2) {
        int fd = open(argv[1], O_RDONLY);
        dup2(fd, STDIN_FILENO);
        close(fd);
    }

    char buffer[1024];
    while (true) {
        size_t numBytesRead = read(STDIN_FILENO, buffer, sizeof(buffer));
        if (numBytesRead == 0) break;
        size_t numBytesWritten = 0;
        while (numBytesWritten < numBytesRead) {
            numBytesWritten +=
                write(STDOUT_FILENO, buffer + numBytesWritten,
                    numBytesRead - numBytesWritten);
        }
    }

    return 0;
}
```