

Announcements

- **Assignment 1 due Monday, Assignment 2 out on Monday as well.**
 - Due just before midnight. If you're new to mercurial and/or the CS110 submissions process (which you are unless you took CS107 and followed the same script), go through the submission process now to confirm everything works sans drama.
 - Assignment 2 will be an opportunity to revisit your Assignment 1 implementation and optimize it to run as quickly as possible.
- **Topics for today:**
 - Continue playing with **fork**, introduce **waitpid**, and maybe even get to the awesomeness that is **execvp**.
 - Lots of examples from previous slide decks to work through, and I have a few more here as well.
 - You can download Truman's lecture slides from today right [here](#)

Enter the **execvp** command

- **execvp** effectively reboots a process to run a different program from scratch.
 - **execvp** has many variants (**execle**, **execlp**, and so forth. Type **man execvp** to see all of them).
 - Here is the prototype:

```
int execvp(const char *path, char *argv[]);
```

- Here's what the arguments and the return type mean:
 - **path** identifies the name of the executable that should be invoked.
 - **argv** is the argument vector that should be funneled through to the new executable's **main** function.
 - Generally, at least for the purposes of CS110, **path** and **argv[0]** end up being the same exact string.
 - If **execvp** fails to consume the process and install a new process image within it, **execvp** will return -1.

Using execvp

▪ Core implementation of **mysystem** (to emulate the **system** builtin)

- Here we present our own implementation of the **mysystem** function, which executes the provided **command** (guaranteed to be a `'\0'`-terminated C string) by calling `"/bin/sh -c command"` and ultimately returning once the surrogate **command** has finished.
- If the execution of **command** exits normally (either via an **exit** system call, or via a normal return statement from **main**), then our **mysystem** implementation should return that exact same exit status.
- If the execution exits abnormally (e.g. it segfaults), then we'll assume it aborted because some signal was ignored, and we'll return that signal number (e.g. 11 for **SIGSEGV**).
- Here's the implementation (online right [here](#))

```
static int mysystem(const char *command) {
    pid_t pid = fork();
    if (pid == 0) {
        char *arguments[] = {"/bin/sh", "-c", (char *) command, NULL};
        execvp("/bin/sh", arguments);
        exitIf(true, kExecFailed, stderr, "execvp failed to invoke this: %s.\n", command);
    }

    int status;
    waitpid(pid, &status, 0);
    if (WIFEXITED(status))
        return WEXITSTATUS(status);
    else
        return WTERMSIG(status);
}
```

- Here's a trivial unit test that I'll run in lecture to prove this thing really works:

```
static const size_t kMaxLine = 2048;
int main(int argc, char *argv[]) {
    char buf[kMaxLine];
    while (true) {
        printf("> ");
        fgets(buf, kMaxLine, stdin);
        if (feof(stdin)) break;
        buf[strlen(buf) - 1] = '\0'; // overwrite '\n'
        printf("retcode = %d\n", mysystem(buf));
    }

    printf("\n");
    return 0;
}
```

Core implementation of **simplesh**, version 1.0

- This is the best example of **execvp** imaginable: a minimum version of a shell not unlike those you've been using since the day you learned UNIX.
 - Relies on **fork**, **waitpid**, and **execvp**.
 - This first version operates as a read-eval-print loop (keyword: repl), responding to many things we type in by forking off child processes.
 - Each child process is initially a deep clone of the shell.
 - Each child proceeds to replace its own process image with the new one we specify (e.g. **ls**, **cp**, our own CS110 **search** (which we wrote the second day of class), or even **emacs**.
 - A trailing ampersand (&) (e.g. **emacs &**) is an instruction to execute in the background, without blocking.
 - Implementation of **simplesh** is presented over the next three slides. Where helper functions don't rely on CS110 concepts, I omit their implementations (but will describe them in lecture).

Core of `simplersh` implementation

■ Implementation of `main`:

- Here's the first half (full implementation right [here](#))

```
int main(int argc, char *argv[]) {
    while (true) {
        char command[kMaxCommandLength + 1];
        readCommand(command, sizeof(command) - 1);
        if (feof(stdin)) break;
        char *arguments[kMaxArgumentCount + 1];
        int count = parseCommandLine(command, arguments, sizeof(arguments)/sizeof(arguments[0]));
        if (count == 0) continue;
        bool builtin = handleBuiltin(arguments);
        if (builtin) continue; // it's been handled, move on
        bool isBackgroundProcess = strcmp(arguments[count - 1], "&") == 0;
        if (isBackgroundProcess) arguments[--count] = NULL; // overwrite "&"
        pid_t pid = forkProcess();
    }
}
```

Core of **simplesh** implementation (continued)

- Implementation of **main**:

- Here's the second half

```
    if (pid == 0) {
        if (execvp(arguments[0], arguments) < 0) {
            printf("%s: Command not found\n", arguments[0]);
            exit(0);
        }
    }

    if (!isBackgroundProcess) {
        waitForChildProcess(pid);
    } else {
        printf("%d %s\n", pid, command);
    }
}

printf("\n");
return 0;
}
```

Core of **simplesh** implementation (continued)

▪ Helper routines

- Here are a few helper routines that do rely on CS110 material:

```
static bool handleBuiltin(char *arguments[]) {
    if (strcasecmp(arguments[0], "quit") == 0) exit(0);
    return strcmp(arguments[0], "&") == 0;
}

static pid_t forkProcess() {
    pid_t pid = fork();
    exitIf(pid == -1, kForkFailed, stderr, "fork function failed.\n");
    return pid;
}

static void waitForChildProcess(pid_t pid) {
    exitUnless(waitpid(pid, NULL, 0) == pid, kWaitFailed,
               stderr, "Error waiting in foreground for process %d to exit", pid);
}
```

Implementing subprocess

▪ Introducing pipe:

- The **pipe** system call takes an uninitialized array of two integers (let's call this array **fds**) and populates it with two file descriptors such that everything *written* to **fds[1]** can be *read* from **fds[0]**.

```
int pipe(int fds[]); // fds array should be of length 2, return -1 on error, 0 otherwise
```

- **pipe** is particularly useful for allowing parent processes to communicate with forked child processes. (Recall that **fork** clones the caller's virtual address space **and** duplicates all open file descriptors as well).
- Using **pipe**, **fork**, **dup2**, **execvp**, **close**, and **waitpid**, we can implement the **subprocess** function, which relies on the following record definition and is implemented to the following prototype:

```
typedef struct {  
    pid_t pid;  
    int infd;  
} subprocess_t;  
subprocess_t subprocess(const char *command);
```

- The subprocess created by **subprocess** executes the provided command (guaranteed to be a '\0'-terminated C string) by calling **"/bin/sh -c command"**. Rather than waiting for **command** to finish, the implementation returns a **subprocess_t** with the **command** process's pid and a single file descriptor. In particular, arbitrary data can be published to the return value's **infd** with the understanding that it'll be read by **command**'s standard input.

Implementing subprocess (continued)

▪ Sample client application, with output

- The following client program and test run illustrate precisely how **subprocess** should work:

```
int main(int argc, char *argv[]) {
    subprocess_t sp = subprocess("/usr/bin/sort");
    const char *words[] = {
        "felicity", "umbrage", "susurraration", "halcyon",
        "pulchritude", "ablution", "somnolent", "indefatigable"
    };
    for (size_t i = 0; i < sizeof(words)/sizeof(words[0]); i++) {
        dprintf(sp.infd, "%s\n", words[i]);
    }
    close(sp.infd); // effectively sends ctrl-D to child process
    int status;
    pid_t pid = waitpid(sp.pid, &status, 0);
    return pid == sp.pid && WIFEXITED(status) ? WEXITSTATUS(status) : -1;
}
```

- The output of the above program, given a properly implemented **subprocess** routine, should look like so:

```
myth22> ./subprocess-test
ablution
felicity
halcyon
indefatigable
pulchritude
somnolent
susurraration
umbrage
```

Implementing subprocess (continued)

■ Implementation is right here (dense):

- Here is the implementation of **subprocess** (where I omit error checking so that the meat of the implementation is clear):

```
subprocess_t subprocess(const char *command) {
    int fds[2];
    pipe(fds);
    subprocess_t process = { fork(), fds[1] };
    if (process.pid == 0) {
        dup2(fds[0], STDIN_FILENO);
        close(fds[0]);
        close(fds[1]);
        char *argv[] = { "/bin/sh", "-c", (char *) command, NULL };
        execvp(argv[0], argv);
    }
    close(fds[0]);
    return process;
}
```

- Note that the write end of the pipe is embedded into the **subprocess_t**. That way, the parent knows where to publish text so that it flows to the other end of the pipe, across the parent process/child process boundary.
- Further note that the child process reassociates the read end of the pipe with its own standard input (using **dup2**).
- Once the **dup2** reassociation has been made, the child process can close both ends of its copy of the pipe.
- The parent doesn't need the read end of the pipe, so it too can close its own copy of **fds[0]**.