

# CS 124/LINGUIST 180

## From Languages to Information


# Dan Jurafsky

# Stanford University

# Recommender Systems & Collaborative Filtering

Slides adapted from Jure Leskovec

# Recommender Systems



## Customer X

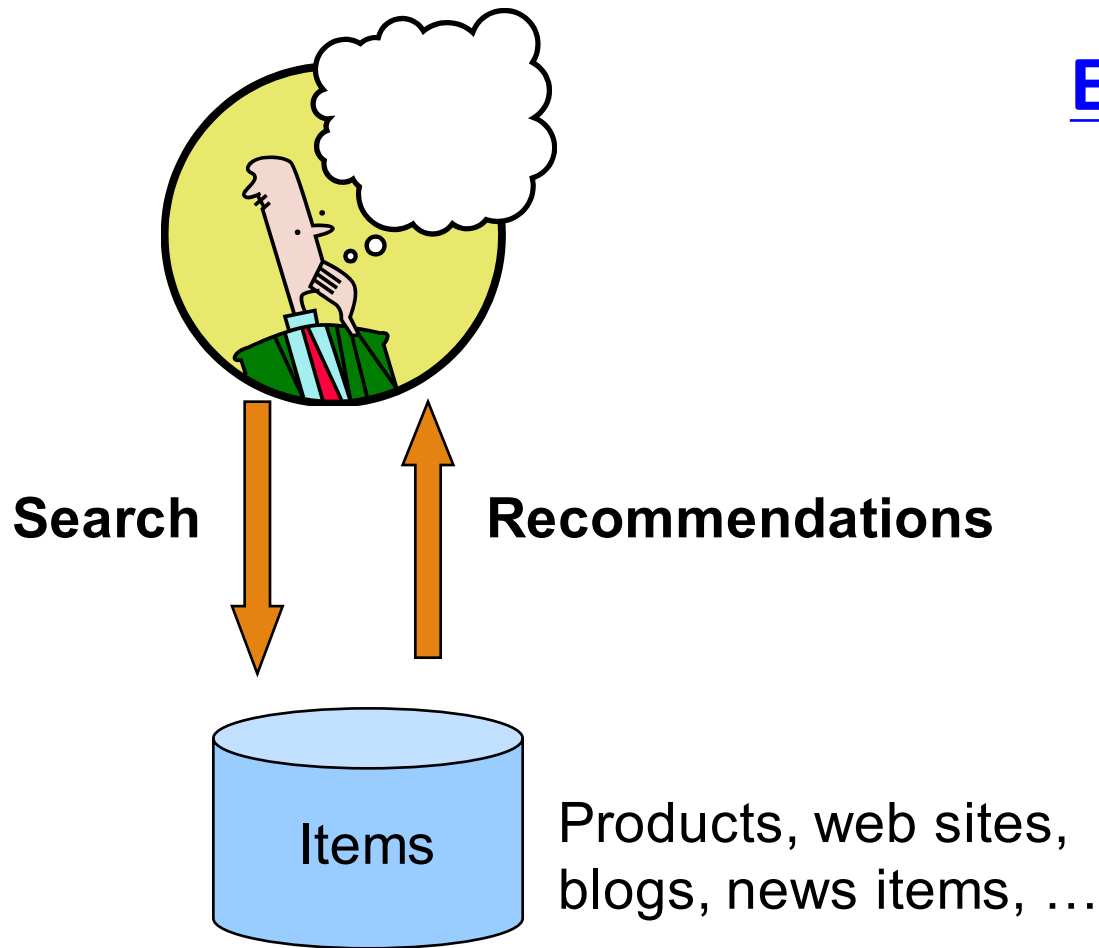
- Buys Metallica CD
- Buys Megadeth CD



## Customer Y


- Does search on Metallica
- Recommender system suggests Megadeth from data collected about customer X

# Recommendations



## Examples:

amazon.com.



**movie lens**  
helping you find the *right* movies

last.fm™  
the social music revolution

Google™  
News

YouTube

XBOX  
LIVE

# From Scarcity to Abundance

## **Shelf space is a scarce commodity for traditional retailers**

- Also: TV networks, movie theaters,...

## **Web enables near-zero-cost dissemination of information about products**

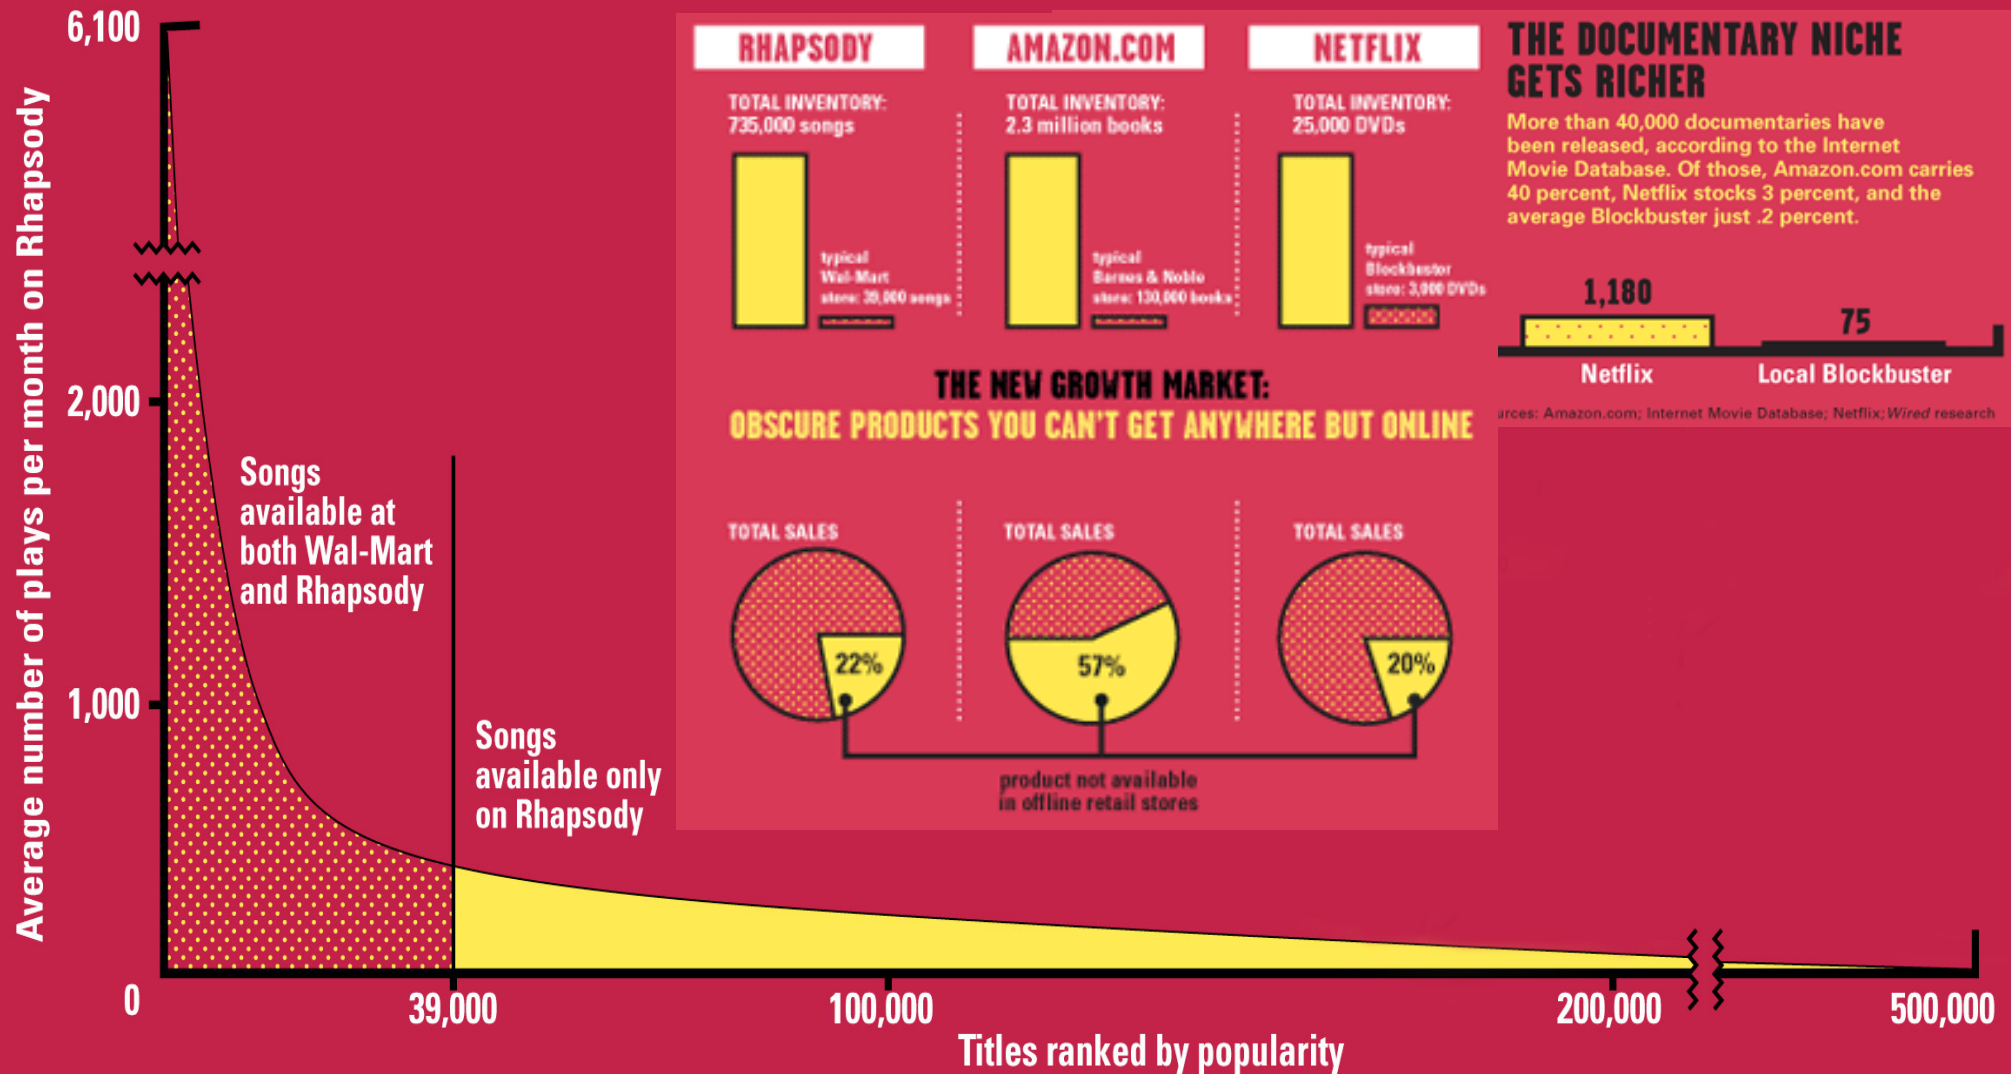
- From scarcity to abundance

## **More choice necessitates better filters**

- Recommendation engines
- How **Into Thin Air** made **Touching the Void** a bestseller:

<http://www.wired.com/wired/archive/12.10/tail.html>

# Sidenote: The Long Tail



Sources: Erik Brynjolfsson and Jeffrey Hu, MIT, and Michael Smith, Carnegie Mellon; Barnes & Noble; Netflix; RealNetworks

Source: Chris Anderson (2004)

SLIDES ADAPTED FROM JURE LESKOVEC




# Physical vs. Online

## THE BIT PLAYER ADVANTAGE

Beyond bricks and mortar there are two main retail models – one that gets halfway down the Long Tail and another that goes all the way. The first is the familiar hybrid model of Amazon and Netflix, companies that sell physical goods online. Digital catalogs allow them to offer unlimited selection along with search, reviews, and recommendations, while the cost savings of massive warehouses and no walk-in customers greatly expands the number of products they can sell profitably.

Pushing this even further are pure digital services, such as iTunes, which offer the additional savings of delivering their digital goods online at virtually no marginal cost. Since an extra database entry and a few megabytes of storage on a server cost effectively nothing, these retailers have no economic reason not to carry *everything* available.



Read <http://www.wired.com/wired/archive/12.10/tail.html> to learn more!

# Types of Recommendations

## Editorial and hand curated

- List of favorites
- Lists of “essential” items

## Simple aggregates

- Top 10, Most Popular, Recent Uploads

## Tailored to individual users

- Amazon, Netflix, ...



# Formal Model

$X$  = set of **Customers**

$S$  = set of **Items**

**Utility function**  $u: X \times S \rightarrow R$

- $R$  = set of ratings
- $R$  is a totally ordered set
- e.g., **0-5** stars, real number in **[0,1]**



# Utility Matrix

	Avatar	LOTR	Matrix	Pirates
Alice	1		0.2	
Bob		0.5		0.3
Carol	0.2		1	
David				0.4

# Key Problems

## **(1) Gathering “known” ratings for matrix**

- How to collect the data in the utility matrix

## **(2) Extrapolate unknown ratings from known ones**

- Mainly interested in high unknown ratings
  - We are not interested in knowing what you don't like but what you like

## **(3) Evaluating extrapolation methods**

- How to measure success/performance of recommendation methods

# (1) Gathering Ratings

## Explicit

- Ask people to rate items
- Doesn't work well in practice – people can't be bothered
- Crowdsourcing: Pay people to label items

## Implicit

- Learn ratings from user actions
  - E.g., purchase implies high rating
- What about low ratings?

## (2) Extrapolating Utilities

**Key problem:** Utility matrix  $U$  is **sparse**

- Most people have not rated most items
- **Cold start:**
  - New items have no ratings
  - New users have no history

**Three approaches to recommender systems:**

1. Content-based
  2. Collaborative
  3. Latent factor based
- } **Today!**

A decorative vertical bar on the left side of the slide, composed of two parallel lines in shades of orange and brown.

# Content-based Recommender Systems

# Content-based Recommendations

**Main idea:** Recommend items to customer  $x$  similar to previous items rated highly by  $x$

*Example:*

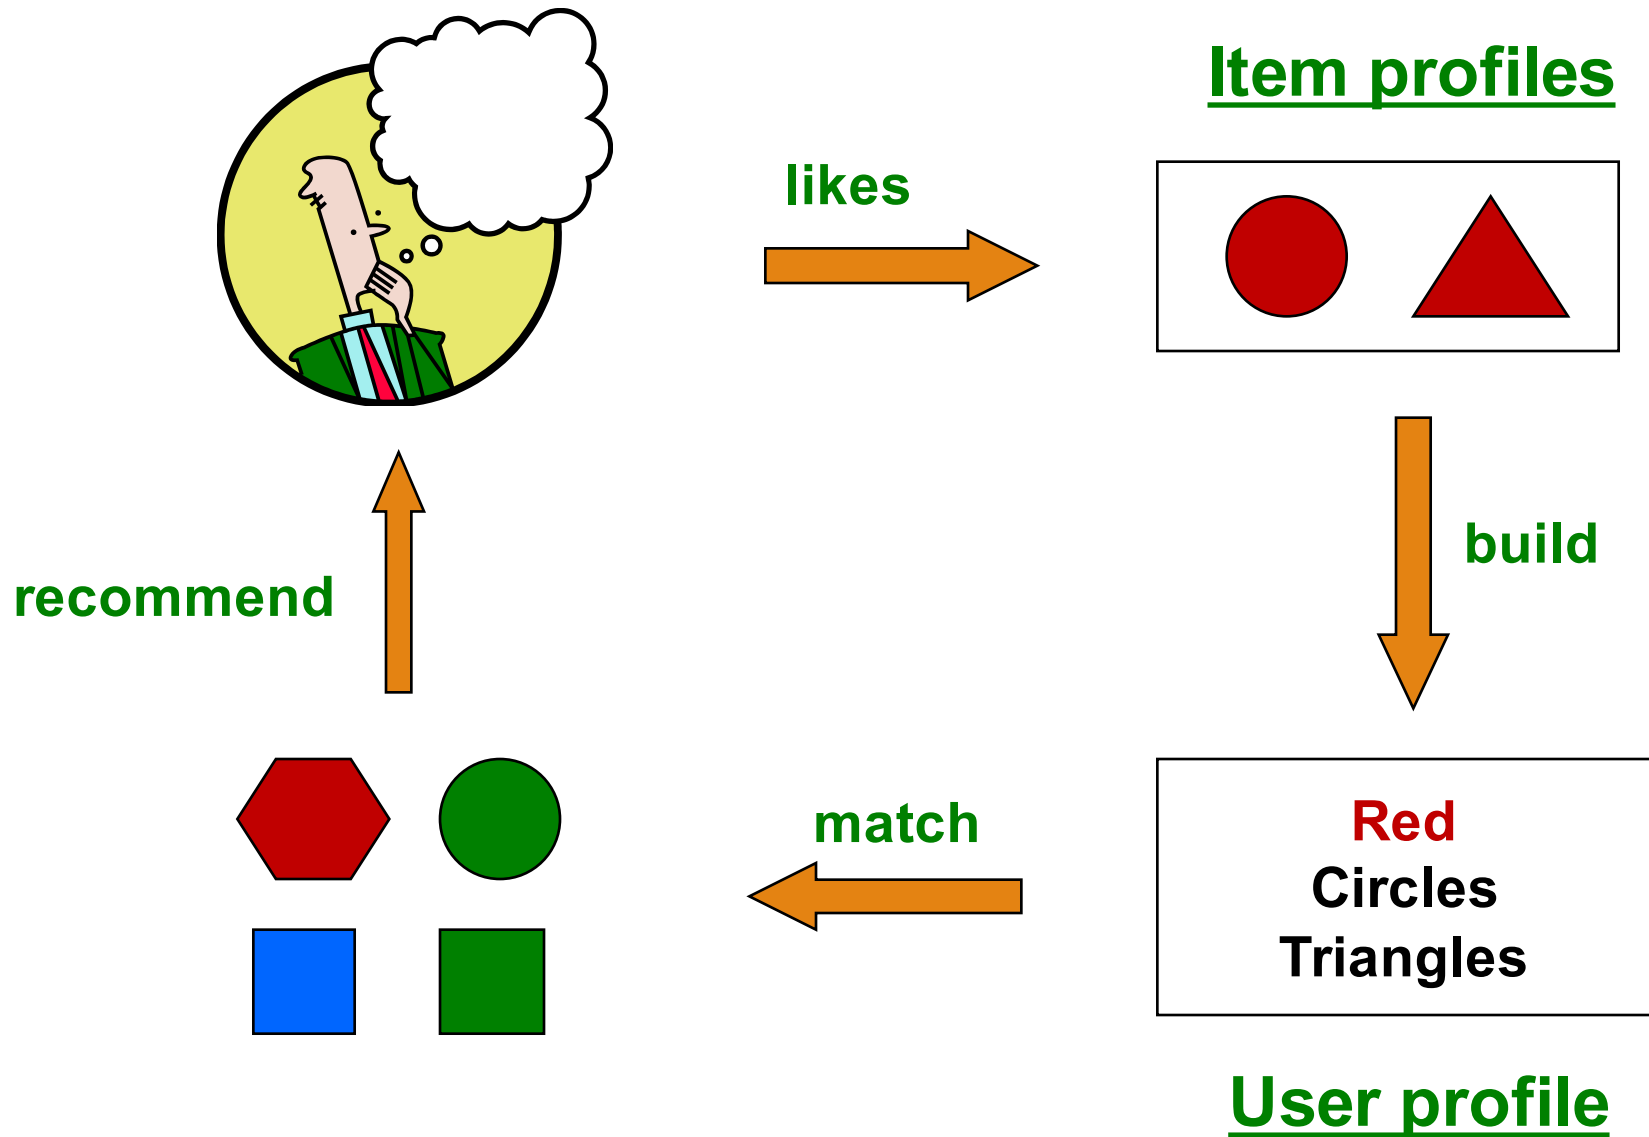
## Movie recommendations

- Recommend movies with same actor(s), director, genre, ...

## Websites, blogs, news

- Recommend other sites with “similar” content

# Plan of Action





# Item Profiles

For each item, create an **item profile**

**Profile is a set (vector) of features**

- **Movies:** author, genre, director, actors, year...
- **Text:** Set of “important” words in document

**How to pick important features?**

- **TF-IDF** (Term frequency \* Inverse Doc Frequency)
  - **Term ... Feature**
  - **Document ... Item**

# Content-based Item Profiles

	Melissa McCarthy	Actor A	Actor B	...	Johnny Depp	Comic Genre	Spy Genre	Pirate Genre	Avg Rating
Movie X	0	1	1	0	1	1	0	1	$3\alpha$
Movie Y	1	1	0	1	0	1	1	0	$4\alpha$

- Maybe there is a scaling factor  $\alpha$  between binary and numeric features
- Or maybe  $\alpha=1$

$$\text{Cosine}(\text{Movie X}, \text{Movie Y}) = \frac{2 + 12\alpha^2}{\sqrt{5 + 9\alpha^2} \sqrt{5 + 16\alpha^2}}$$

# User Profiles

**Want a vector with the same components/dimensions as items**

- Could be 1s representing user purchases
- Or arbitrary numbers from a rating

**User profile is aggregate of items:**

- Average(weighted?) of rated item profiles

# Sample user profile

- Items are movies
- Utility matrix has 1 if user has seen movie
- 20% of the movies user U has seen have Melissa McCarthy
- $U[\text{"Melissa McCarthy"}] = 0.2$

	Melissa McCarthy	Actor A	Actor B	...	
User U	0.2	.0005	0	0	...

# Prediction

- User and item vectors have the same components/dimensions!
- So just recommend the items whose vectors are most similar to the user vector!
- Given user profile  $\mathbf{x}$  and item profile  $\mathbf{i}$ ,
- estimate  $u(\mathbf{x}, \mathbf{i}) = \cos(\mathbf{x}, \mathbf{i}) = \frac{\mathbf{x} \cdot \mathbf{i}}{||\mathbf{x}|| \cdot ||\mathbf{i}||}$

# Pros: Content-based Approach

## **+: No need for data on other users**

- No cold-start or sparsity problems

## **+: Able to recommend to users with unique tastes**

## **+: Able to recommend new & unpopular items**

- No first-rater problem

## **+: Able to provide explanations**

- Can provide explanations of recommended items by listing content-features that caused an item to be recommended

# Cons: Content-based Approach

## –: Finding the appropriate features is hard

- E.g., images, movies, music

## –: Recommendations for new users

- How to build a user profile?

## –: Overspecialization

- Never recommends items outside user's content profile
- People might have multiple interests
- Unable to exploit quality judgments of other users



A decorative vertical bar on the left side of the slide, composed of two parallel lines in shades of orange and brown.

# Collaborative Filtering

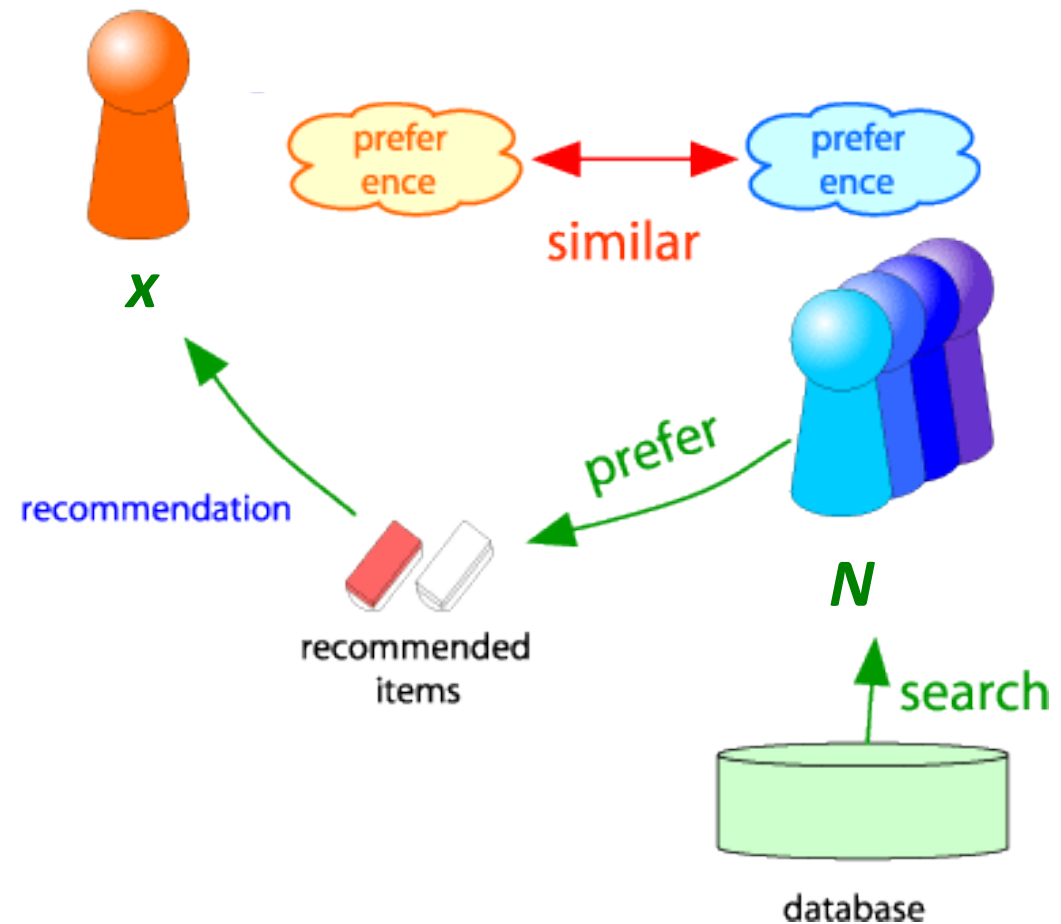
**Harnessing quality judgments of other users**

# Collaborative Filtering

Consider user  $x$

Find set  $N$  of other users whose ratings are “**similar**” to  $x$ ’s ratings

Estimate  $x$ ’s ratings based on ratings of users in  $N$



# Finding Similar Users

$$\begin{aligned} r_x &= [*, \_, \_, *, ***] \\ r_y &= [*, \_, **, **, \_] \end{aligned}$$

Let  $r_x$  be the vector of user  $x$ 's ratings

## Jaccard similarity measure

- **Problem:** Ignores the value of the rating

*$r_x, r_y$  as sets:*

$$r_x = \{1, 4, 5\}$$

$$r_y = \{1, 3, 4\}$$

## Cosine similarity measure

- $\text{sim}(x, y) = \cos(r_x, r_y) = \frac{r_x \cdot r_y}{\|r_x\| \cdot \|r_y\|}$

*$r_x, r_y$  as points:*

$$r_x = \{1, 0, 0, 1, 3\}$$

$$r_y = \{1, 0, 2, 2, 0\}$$

- **Problem:** Treats missing ratings as “negative”

— —

# Utility Matrix

	HP1	HP2	HP3	TW	SW1	SW2	SW3
<i>A</i>	4			5	1		
<i>B</i>	5	5	4				
<i>C</i>				2	4	5	
<i>D</i>		3					3

**Intuitively we want:**  $\text{sim}(A, B) > \text{sim}(A, C)$

**Jaccard similarity:**  $1/5 < 2/4$

**Cosine similarity:**  $0.386 > 0.322$

Considers missing ratings as “negative”

# Utility Matrix

	HP1	HP2	HP3	TW	SW1	SW2	SW3
<i>A</i>	4			5	1		
<i>B</i>	5	5	4				
<i>C</i>				2	4	5	
<i>D</i>		3					3

- Problem with cosine: 0 acts like a negative review
  - C really loves SW
  - A hates SW
  - B just hasn't seen it
- Another problem: we'd like to normalize for raters
  - D rated everything the same; not very useful

# Modified Utility Matrix:

subtract the means of each row

	HP1	HP2	HP3	TW	SW1	SW2	SW3
<i>A</i>	4			5	1		
<i>B</i>	5	5	4				
<i>C</i>				2	4	5	
<i>D</i>		3					3

	HP1	HP2	HP3	TW	SW1	SW2	SW3
<i>A</i>	$2/3$			$5/3$	$-7/3$		
<i>B</i>	$1/3$	$1/3$	$-2/3$				
<i>C</i>				$-5/3$	$1/3$	$4/3$	
<i>D</i>		0					0

- Now a 0 means no information
- And negative ratings means viewers with opposite ratings will have vectors in opposite directions!

# Modified Utility Matrix:

subtract the means of each row

	HP1	HP2	HP3	TW	SW1	SW2	SW3
<i>A</i>	2/3			5/3	-7/3		
<i>B</i>	1/3	1/3	-2/3				
<i>C</i>				-5/3	1/3	4/3	
<i>D</i>		0					0

$$\text{Cos}(A,B) = \frac{(2/3) \times (1/3)}{\sqrt{(2/3)^2 + (5/3)^2 + (-7/3)^2} \sqrt{(1/3)^2 + (1/3)^2 + (-2/3)^2}} = 0.092$$

$$\text{Cos}(A,C) = \frac{(5/3) \times (-5/3) + (-7/3) \times (1/3)}{\sqrt{(2/3)^2 + (5/3)^2 + (-7/3)^2} \sqrt{(-5/3)^2 + (1/3)^2 + (4/3)^2}} = -0.559$$

Now A and C are (correctly) way further apart than A,B



# Cosine after subtracting mean

Turns out to be the same as Pearson correlation coefficient!!!

Cosine similarity is correlation when the data is centered at 0

- Terminological Note: subtracting the mean is **zero-centering**, not **normalizing** (normalizing is dividing by a norm to turn something into a probability), but the textbook (and in common use) we sometimes overload the term “normalize”

# Finding Similar Users

$$\begin{aligned} r_x &= [*, \_, \_, *, ***] \\ r_y &= [*, \_, **, **, \_] \end{aligned}$$

Let  $r_x$  be the vector of user  $x$ 's ratings

## Cosine similarity measure

$r_x, r_y$  as points:

$$\begin{aligned} r_x &= \{1, 0, 0, 1, 3\} \\ r_y &= \{1, 0, 2, 2, 0\} \end{aligned}$$

- $\text{sim}(x, y) = \cos(r_x, r_y) = \frac{r_x \cdot r_y}{\|r_x\| \cdot \|r_y\|}$
- **Problem:** Treats missing ratings as “negative”

## Pearson correlation coefficient

- $S_{xy}$  = items rated by both users  $x$  and  $y$

$$\text{sim}(x, y) = \frac{\sum_{s \in S_{xy}} (r_{xs} - \bar{r}_x)(r_{ys} - \bar{r}_y)}{\sqrt{\sum_{s \in S_{xy}} (r_{xs} - \bar{r}_x)^2} \sqrt{\sum_{s \in S_{xy}} (r_{ys} - \bar{r}_y)^2}}$$

$\bar{r}_x, \bar{r}_y \dots$  avg.  
rating of  $x, y$

# Rating Predictions

## From similarity metric to recommendations:

Let  $\mathbf{r}_x$  be the vector of user  $x$ 's ratings

Let  $N$  be the set of  $k$  users most similar to  $x$  who have rated item  $i$

## Prediction for item $i$ of user $x$ :

- $r_{xi} = \frac{1}{k} \sum_{y \in N} r_{yi}$
- $r_{xi} = \frac{\sum_{y \in N} s_{xy} \cdot r_{yi}}{\sum_{y \in N} s_{xy}}$

**Shorthand:**

$$s_{xy} = \textit{sim}(x, y)$$

- **Many other tricks possible...**

# Item-Item Collaborative Filtering

So far: **User-user collaborative filtering**

**Another view: Item-item**

- For item  $i$ , find other similar items
- Estimate rating for item  $i$  based on ratings for similar items
- Can use same similarity metrics and prediction functions as in user-user model

$$r_{xi} = \frac{\sum_{j \in N(i;x)} s_{ij} \cdot r_{xj}}{\sum_{j \in N(i;x)} s_{ij}}$$


$s_{ij}$ ... similarity of items  $i$  and  $j$   
 $r_{xj}$ ...rating of user  $x$  on item  $j$   
 $N(i;x)$ ...set of items rated by  $x$  similar to  $i$

# Item-Item CF ( $|N|=2$ )

users

movies

	1	2	3	4	5	6	7	8	9	10	11	12
1	1		3			5			5		4	
2			5	4			4			2	1	3
3	2	4		1	2		3		4	3	5	
4		2	4		5			4			2	
5			4	3	4	2					2	5
6	1		3		3			2			4	



- unknown rating



- rating between 1 to 5

# Item-Item CF ( $|N|=2$ )

		users											
		1	2	3	4	5	6	7	8	9	10	11	12
movies	1	1		3		?	5			5		4	
	2			5	4			4			2	1	3
	3	2	4		1	2		3		4	3	5	
	4		2	4		5			4			2	
	5			4	3	4	2					2	5
	6	1		3		3			2			4	



- estimate rating of movie 1 by user 5

# Item-Item CF ( $|N|=2$ )

	users												
	1	2	3	4	5	6	7	8	9	10	11	12	$\text{sim}(1,m)$
1	1		3		?	5			5		4		1.00
2			5	4			4			2	1	3	-0.18
<u>3</u>	2	4		1	2		3		4	3	5		<u>0.41</u>
4		2	4		5			4			2		-0.10
5			4	3	4	2					2	5	-0.31
<u>6</u>	1		3		3			2			4		<u>0.59</u>

## Neighbor selection:

Identify movies similar to  
movie 1, **rated by user 5**

Here we use Pearson correlation as similarity:

1) Subtract mean rating  $m_i$  from each movie  $i$

$$m_1 = (1+3+5+5+4)/5 = 3.6$$

row 1:  $[-2.6, 0, -0.6, 0, 0, 1.4, 0, 0, 1.4, 0, 0.4, 0]$

2) Compute cosine similarities between rows



# Item-Item CF ( $|N|=2$ )

		users												sim(1,m)
		1	2	3	4	5	6	7	8	9	10	11	12	
movies	1	1		3		?	5			5		4		1.00
	2			5	4			4			2	1	3	-0.18
	<u>3</u>	2	4		1	2		3		4	3	5		<u>0.41</u>
	4		2	4		5			4			2		-0.10
	5			4	3	4	2					2	5	-0.31
	<u>6</u>	1		3		3			2			4		<u>0.59</u>

Compute similarity weights:

$$s_{1,3}=0.41, s_{1,6}=0.59$$

# Item-Item CF ( $|N|=2$ )

		users											
		1	2	3	4	5	6	7	8	9	10	11	12
movies	1	1		3		2.6	5			5		4	
	2			5	4			4			2	1	3
	<u>3</u>	2	4		1	2		3		4	3	5	
	4		2	4		5			4			2	
	5			4	3	4	2					2	5
	<u>6</u>	1		3		3			2			4	

Predict by taking weighted average:

$$r_{1,5} = (0.41 \cdot 2 + 0.59 \cdot 3) / (0.41 + 0.59) = 2.6$$

$$r_{ix} = \frac{\sum_{j \in N(i;x)} s_{ij} \cdot r_{jx}}{\sum s_{ij}}$$

# Item-Item vs. User-User

	Avatar	LOTR	Matrix	Pirates
Alice	1		0.8	
Bob		0.5		0.3
Carol	0.9		1	0.8
David			1	0.4

- In practice, it has been observed that item-item often works better than user-user
- **Why?** Items are simpler, users have multiple tastes

# Pros/Cons of Collaborative Filtering

## + Works for any kind of item

- No feature selection needed

## - Cold Start:

- Need enough users in the system to find a match

## - Sparsity:

- The user/ratings matrix is sparse
- Hard to find users that have rated the same items

## - First rater:

- Cannot recommend an item that has not been previously rated
- New items, Esoteric items

## - Popularity bias:

- Cannot recommend items to someone with unique taste
- Tends to recommend popular items

# Hybrid Methods

**Implement two or more different recommenders and combine predictions**

- Perhaps using a linear model

**Add content-based methods to collaborative filtering**

- Item profiles for new item problem
- Demographics to deal with new user problem

# Evaluation

**movies**

**users**

1	3	4			
	3	5			5
		4	5		5
		3			
		3			
2			2		2
				5	
	2	1			1
	3			3	
1					

# Evaluation

**movies**

**users**

1	3	4			
	3	5			5
		4	5		5
		3			
		3			
2			?		?
				?	
	2	1			?
	3			?	
1					

**Test Data Set**

# Evaluating Predictions

## Compare predictions with known ratings

- **Root-mean-square error (RMSE)**

- $\sqrt{\sum_{xi} (r_{xi} - r_{xi}^*)^2}$  where  $r_{xi}$  is predicted,  $r_{xi}^*$  is the true rating of  $x$  on  $i$

- **Rank Correlation:**

- Spearman's *correlation* between system's and user's complete rankings



# Problems with Error Measures

**Narrow focus on accuracy sometimes misses the point**

- Prediction Diversity
- Prediction Context
- Order of predictions

**In practice, we care only to predict high ratings:**

- RMSE might penalize a method that does well for high ratings and badly for others

# There's No Data like Mo' Data

## Leverage all the data

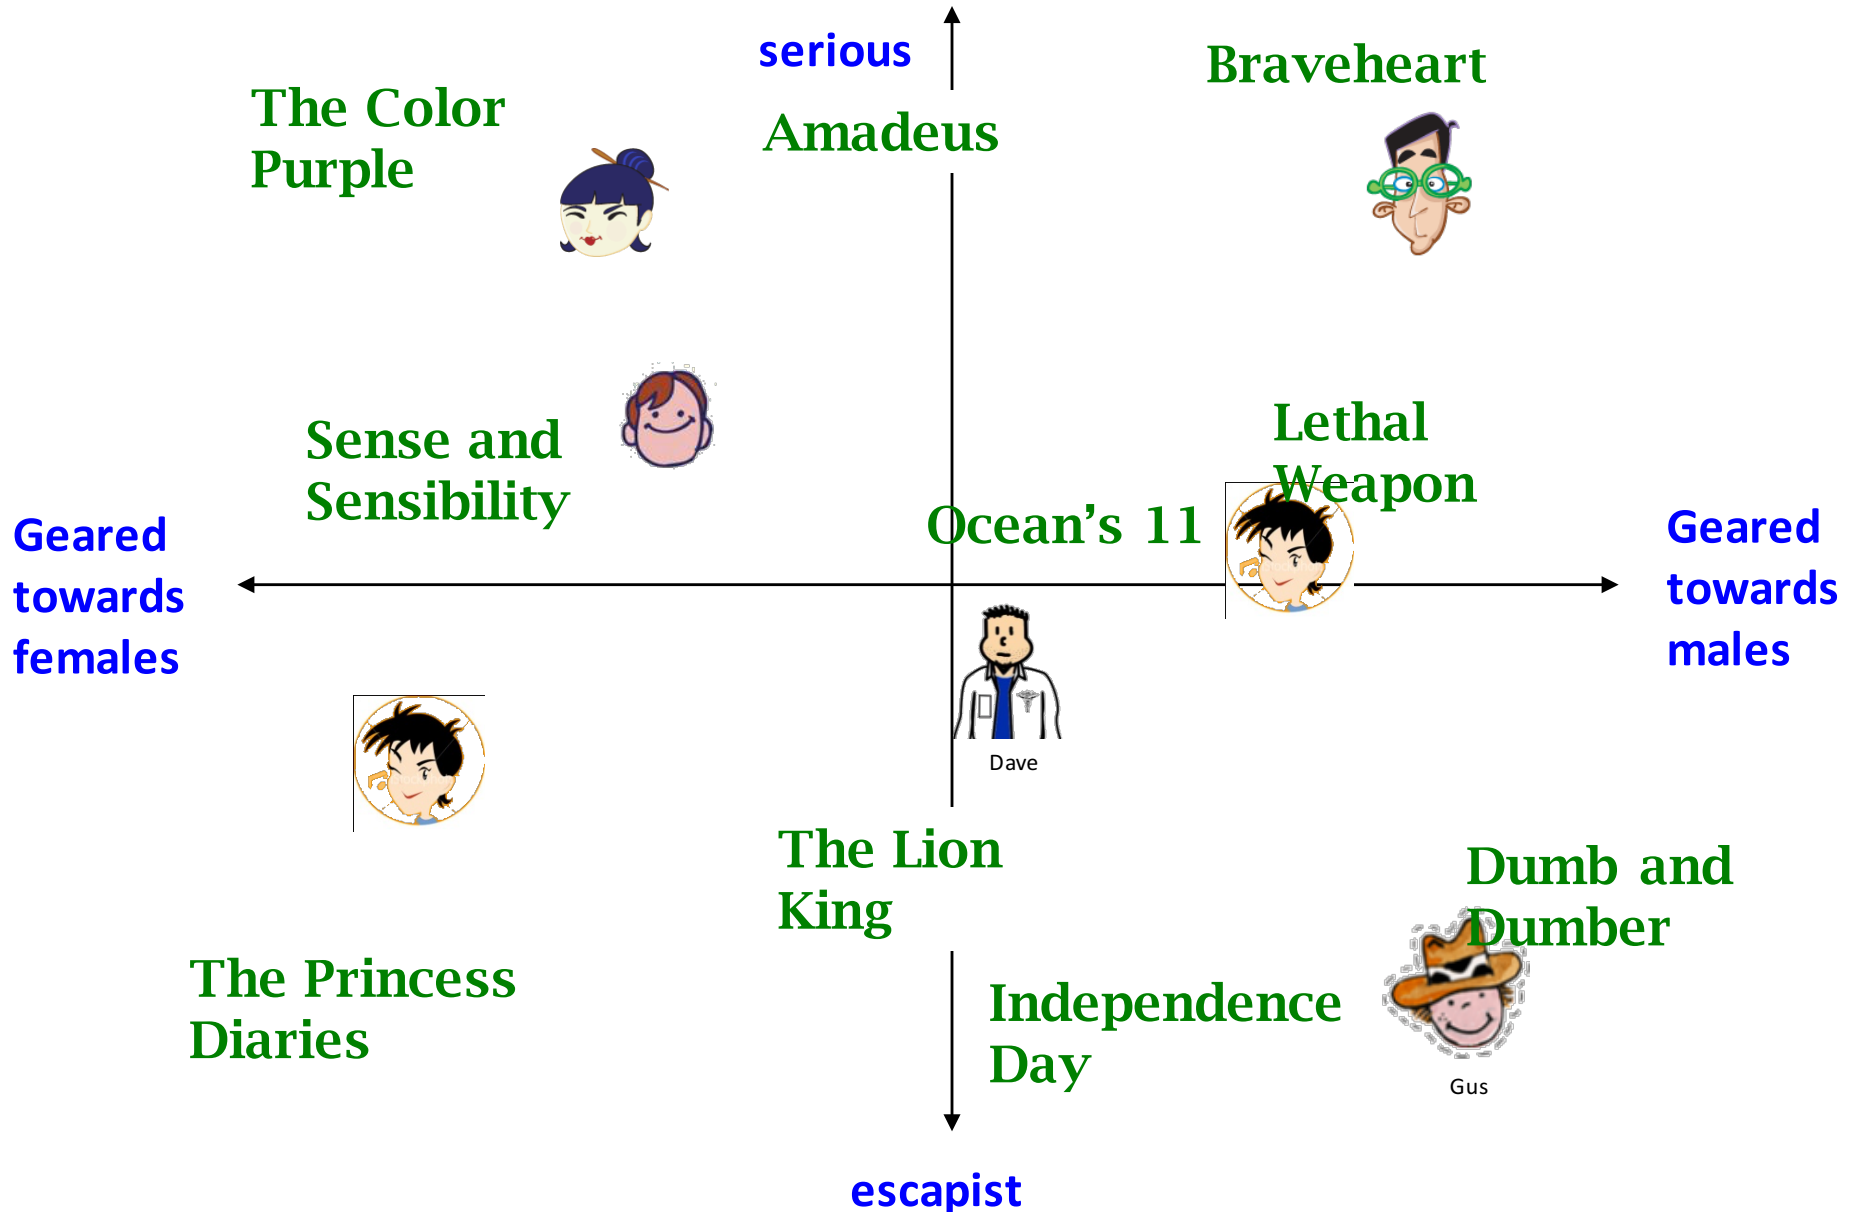
- Simple methods on large data do best

## Add more data

- e.g., add IMDB data on genres

## More data beats better algorithms

# Latent Factor Models (like SVD)



# Famous Historical Example: The Netflix Prize

## Training data

- 100 million ratings, 480,000 users, 17,770 movies
- 6 years of data: 2000-2005

## Test data

- Last few ratings of each user (2.8 million)
- Evaluation criterion: root mean squared error (RMSE)
- Netflix Cinematch RMSE: 0.9514

## Competition

- 2700+ teams
- \$1 million prize for 10% improvement on Cinematch
- BellKor system won in 2009. Combined many factors
  - Overall deviations of users/movies
  - Regional effects
  - Local collaborative filtering patterns
  - Temporal biases

A decorative vertical bar on the left side of the slide, composed of two parallel lines in shades of orange and brown.

# Summary on Recommendation Systems

- The Long Tail
- Content-based Systems
- Collaborative Filtering
- Latent Factors