



Why language is hard

And what Linguistics has to say about it

Natalia Silveira

Participation code: eagles



Language processing is so easy for humans that it is like trying to sell cargo airplanes to eagles.



Language processing is so easy for humans that it is like trying to sell cargo airplanes to eagles. They just don't get what is hard, what is easy and the necessity of infrastructure. "Mr. Eagle, um, well we really need a runway to get the 20 tons of product into the air".



Language processing is so easy for humans that it is like trying to sell cargo airplanes to eagles. They just don't get what is hard, what is easy and the necessity of infrastructure. "Mr. Eagle, um, well we really need a runway to get the 20 tons of product into the air". Mr. Eagle responds with "What are you talking about? I can take off, land and raise a family on a tree branch. Cargo planes are easy because flying is easy for me. So I will give you a fish to do the job."

Breck Baldwin in the LingPipe blog




Overview of today's lecture

- (A few reasons) **Why language is hard**
 - Exploring coarse processing vs. fine processing
 - In between: approaches to sentiment and difficult problems
- Intermission
- (A little bit of) **What Linguistics has to say about it**
 - Part-Of-Speech tagging
 - Dependency parsing



Language is multidimensional





Coarse vs. fine processing

- How do we do text classification?
- How do we do relation extraction?
- Why?



Coarse vs. fine processing

- How do we do text classification?
- How do we do relation extraction?
- Why?



Why language is hard: relation extraction

<PER>Bill Gates</PER> **founded** <ORG>Microsoft</ORG>



Why language is hard: relation extraction

Bill Gates founded Microsoft



Why language is hard: relation extraction

Bill Gates founded Microsoft

Bill Gates, now retired, founded the famous Microsoft



Why language is hard: relation extraction

Bill Gates founded Microsoft

Bill Gates, now retired, founded the famous Microsoft

Microsoft, which Bill Gates founded, is a big software company.



Why language is hard: relation extraction

Bill Gates founded Microsoft

Bill Gates, now retired, founded the famous Microsoft

Microsoft, which Bill Gates founded, is a big software company.

Microsoft was founded by Bill Gates.



Why language is hard: relation extraction

Bill Gates founded Microsoft

Bill Gates, now retired, founded the famous Microsoft

Microsoft, which Bill Gates founded, is a big software company.

Microsoft was founded by Bill Gates.

Microsoft was founded not by Larry Page but by Bill Gates.



Why language is hard: relation extraction

Bill Gates founded Microsoft

Bill Gates, now retired, founded the famous Microsoft

Microsoft, which Bill Gates founded, is a big software company.

Microsoft was founded by Bill Gates.

Microsoft was founded not by Larry Page but by Bill Gates.

Microsoft was founded not only by Bill Gates but also Paul Allen.



Why language is hard: relation extraction

Bill Gates founded Microsoft

Bill Gates, now retired, founded the famous Microsoft

Microsoft, which Bill Gates founded, is a big software company.

Microsoft was founded by Bill Gates.

Microsoft was founded not by Larry Page but by Bill Gates.

Microsoft was founded not only by Bill Gates but also Paul Allen.

And that's not even the hard stuff!



Why language is hard: sentiment analysis

I liked this movie.



Why language is hard: sentiment analysis

I liked this movie.

I didn't like this movie.



Why language is hard: sentiment analysis

I liked this movie.

I didn't like this movie.

I thought I would like this movie.



Why language is hard: sentiment analysis

I liked this movie.

I didn't like this movie.

I thought I would like this movie.

I thought this movie would be great.



Why language is hard: sentiment analysis

I liked this movie.

I didn't like this movie.

I thought I would like this movie.

I thought this movie would be great.

I knew this movie would be great.



Why language is hard: sentiment analysis

I liked this movie.


I didn't like this movie.

I thought I would like this movie.

I thought this movie would be great.

I knew this movie would be great.

I didn't know this movie would be so great.



Part-of-speech tagging

A simple but useful form of linguistic analysis

Slides by Christopher Manning and Ray Mooney

Participation code: eagles

Open class (lexical) words

Nouns

Proper

IBM
Italy

Common

cat / cats
snow

Verbs

Main

see
registered

Adjectives

old older oldest

Adverbs

slowly

Numbers

122,312
one

... more

Closed class (functional)

Determiners

the some

Conjunctions

and or

Pronouns

he its

Modals

can
had

Prepositions

to with

Particles

off up

... more

Interjections

Ow Eh



Open vs. Closed classes

- Open vs. Closed classes
 - Closed:
 - determiners: *a, an, the*
 - pronouns: *she, he, I*
 - prepositions: *on, under, over, near, by, ...*
 - Why “closed”?
 - Open:
 - Nouns, Verbs, Adjectives, Adverbs.



POS Tagging

- Words often have more than one POS: *back*
 - The back door = JJ
 - On my back = NN
 - Win the voters back = RB
 - Promised to back the bill = VB
- The POS tagging problem is to determine the POS tag for a particular instance of a word.



POS Tagging

- Input: Plays well with others
- Ambiguity: NNS/VBZ UH/JJ/NN/RB IN NNS
- Output: Plays/VBZ well/RB with/IN others/NNS
- Uses:
 - Text-to-speech (how do we pronounce “lead”?)
 - Can write regexps like (Det) Adj* N+ over the output for phrases, etc.
 - As input to or to speed up a full parser
 - If you know the tag, you can back off to it in other tasks

Penn
Treebank
POS tags



POS tagging performance

- How many tags are correct? (Tag accuracy)
 - About 97% currently
 - But baseline is already 90%
 - Baseline is performance of stupidest possible method
 - Tag every word with its most frequent tag
 - Tag unknown words as nouns
 - Partly easy because
 - Many words are unambiguous
 - You get points for them (*the*, *a*, etc.) and for punctuation marks!



Deciding on the correct part of speech can be difficult even for people

- Mrs/NNP Shaefer/NNP never/RB got/VBD around/RP to/TO joining/VBG
- All/DT we/PRP gotta/VBN do/VB is/VBZ go/VB around/IN the/DT corner/NN
- Chateau/NNP Petrus/NNP costs/VBZ around/RB 250/CD



How difficult is POS tagging?

- About 11% of the word types in the Brown corpus are ambiguous with regard to part of speech
- But they tend to be very common words. E.g., *that*
 - I know *that* he is honest = IN
 - Yes, *that* play was nice = DT
 - You can't go *that* far = RB
- 40% of the word tokens are ambiguous



Sources of information

- What are the main sources of information for POS tagging?
 - Knowledge of neighboring words
 - Bill saw that man yesterday
 - NNP NN DT NN NN
 - VB VB(D) IN VB NN
 - Knowledge of word probabilities
 - *man* is rarely used as a verb....
- The latter proves the most useful, but the former also helps



More and Better Features → Feature-based tagger

- Can do surprisingly well just looking at a word by itself:
 - Word the: the → DT
 - Lowercased word Importantly: importantly → RB
 - Prefixes unfathomable: un- → JJ
 - Suffixes Importantly: -ly → RB
 - Capitalization Meridian: CAP → NNP
 - Word shapes 35-year: d-x → JJ
- Then build a maxent (or whatever) model to predict tag
 - Maxent $P(t|w)$: 93.7% overall / 82.6% unknown



What else can we do?

- Build better features!

PRP VBD ^{RB} IN RB IN PRP VBD .
 They left as soon as he arrived .

- We could fix this with a feature that looked at the next word

JJ
 NNP NNS VBD VBN .
 Intrinsic flaws remained undetected .

- We could fix this by linking capitalized words to their lowercase versions



Sequence Labeling as Classification

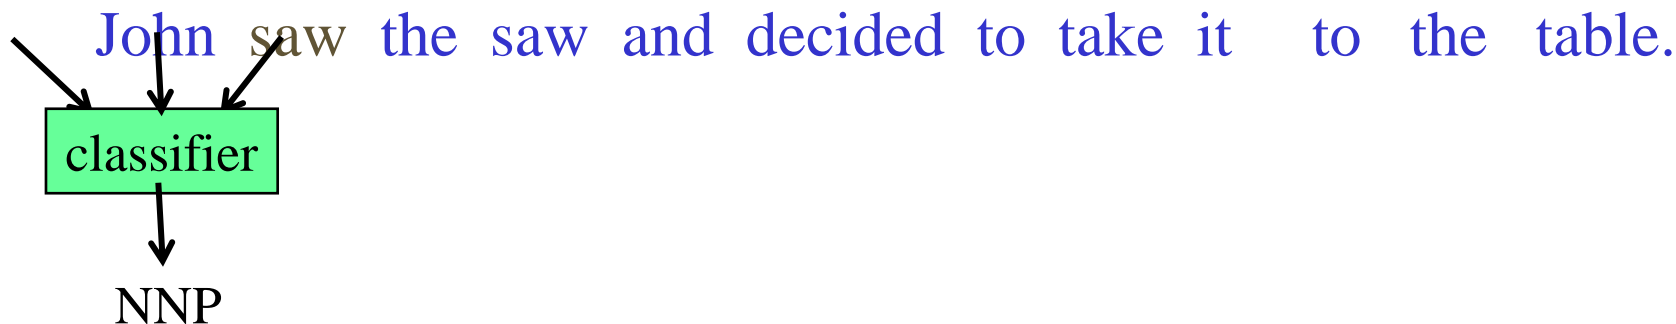
- Classify each token independently but use as input features, information about the surrounding tokens (sliding window).

John saw the saw and decided to take it to the table.



Sequence Labeling as Classification


- Classify each token independently but use as input features, information about the surrounding tokens (sliding window).





Sequence Labeling as Classification

- Classify each token independently but use as input features, information about the surrounding tokens (sliding window).





Sequence Labeling as Classification

- Classify each token independently but use as input features, information about the surrounding tokens (sliding window).

John saw the saw and decided to take it to the table.




NN



Sequence Labeling as Classification

- Classify each token independently but use as input features, information about the surrounding tokens (sliding window).






Sequence Labeling as Classification

- Classify each token independently but use as input features, information about the surrounding tokens (sliding window).

John saw the saw and decided to take it to the table.



CC



Sequence Labeling as Classification

- Classify each token independently but use as input features, information about the surrounding tokens (sliding window).

John saw the saw and decided to take it to the table.

A diagram illustrating the sequence labeling process. A green rectangular box labeled 'classifier' is positioned below the sentence. Three arrows point from the words 'saw', 'and', and 'decided' in the sentence above to the top of the 'classifier' box. An arrow points from the bottom of the 'classifier' box to the label 'VBD' below it.

classifier

VBD



Sequence Labeling as Classification

- Classify each token independently but use as input features, information about the surrounding tokens (sliding window).

John saw the saw and decided to take it to the table.

A diagram illustrating a sliding window for sequence labeling. Three arrows point from the words "saw", "and", and "decided" in the sentence above to a green rectangular box labeled "classifier". An arrow points from the "classifier" box down to the word "TO".

classifier


TO



Sequence Labeling as Classification

- Classify each token independently but use as input features, information about the surrounding tokens (sliding window).

John saw the saw and decided to take it to the table.




VB



Sequence Labeling as Classification

- Classify each token independently but use as input features, information about the surrounding tokens (sliding window).

John saw the saw and decided to take it to the table.




PRP



Sequence Labeling as Classification

- Classify each token independently but use as input features, information about the surrounding tokens (sliding window).

John saw the saw and decided to take it to the table.






Sequence Labeling as Classification

- Classify each token independently but use as input features, information about the surrounding tokens (sliding window).

John saw the saw and decided to take it to the table.




DT



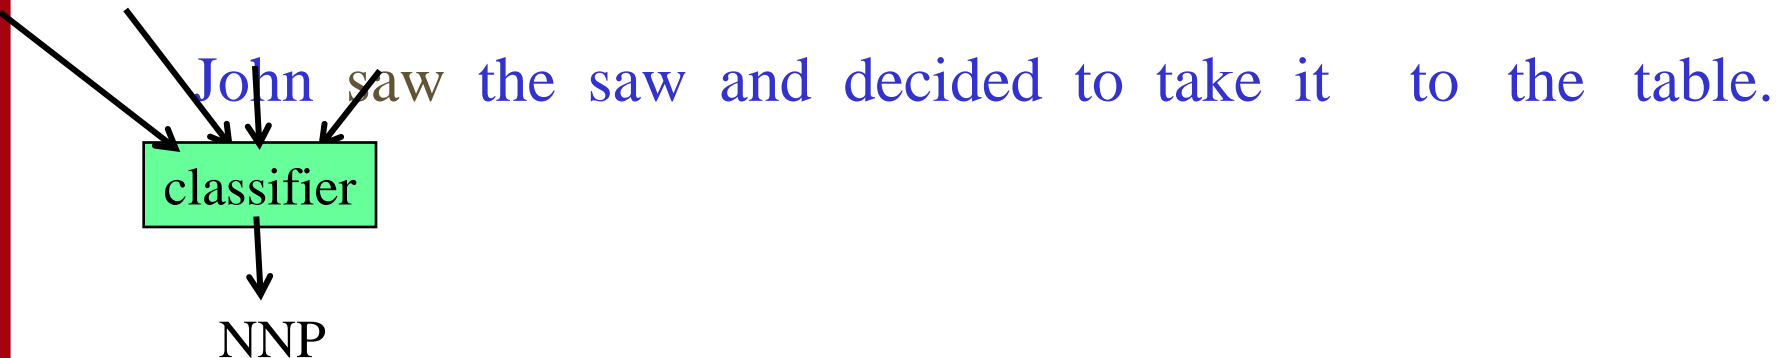
Sequence Labeling as Classification

- Classify each token independently but use as input features, information about the surrounding tokens (sliding window).



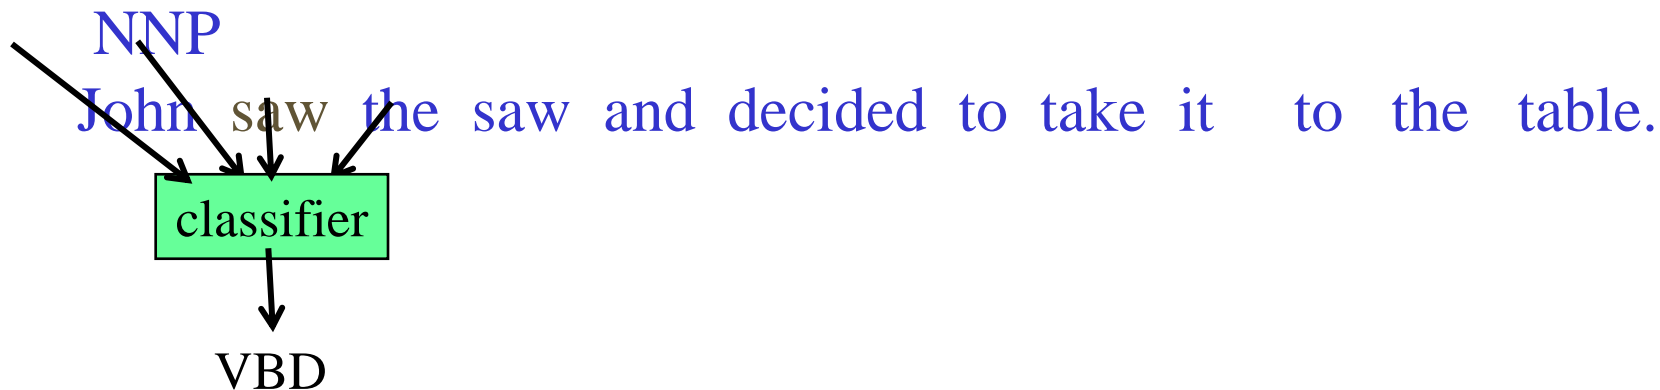


Forward Classification



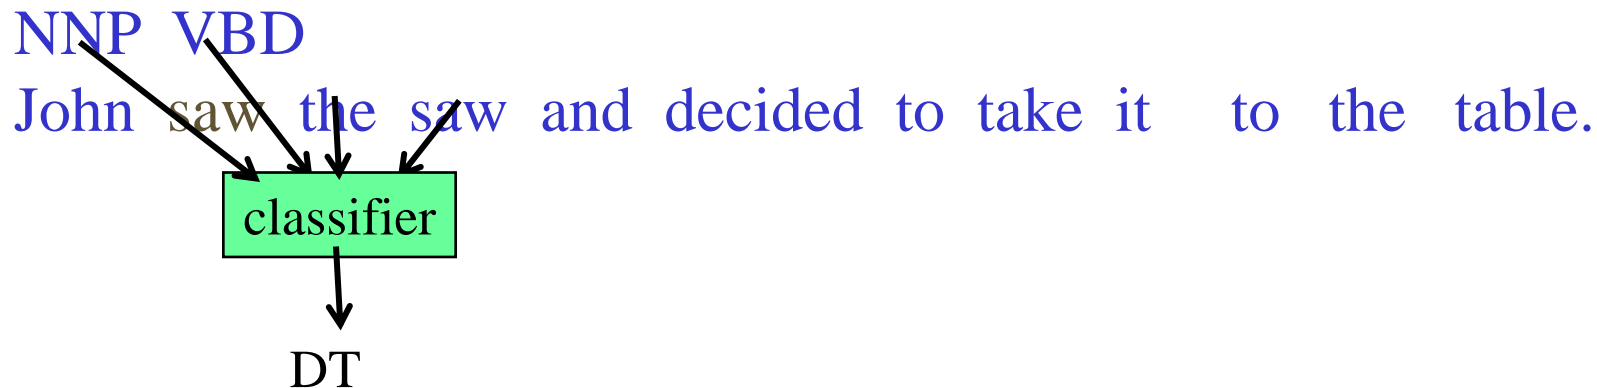


Forward Classification



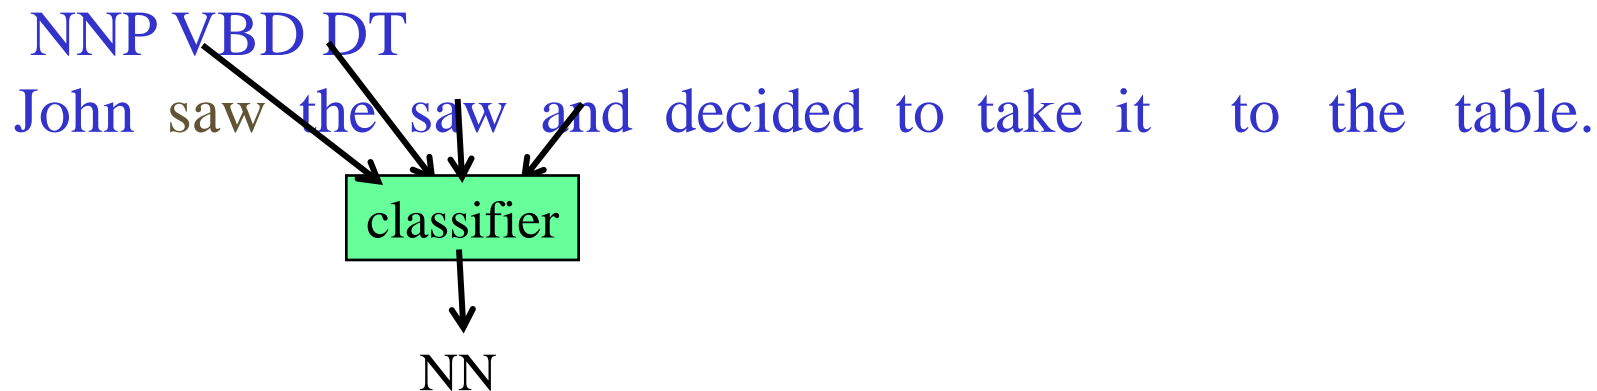


Forward Classification






Forward Classification



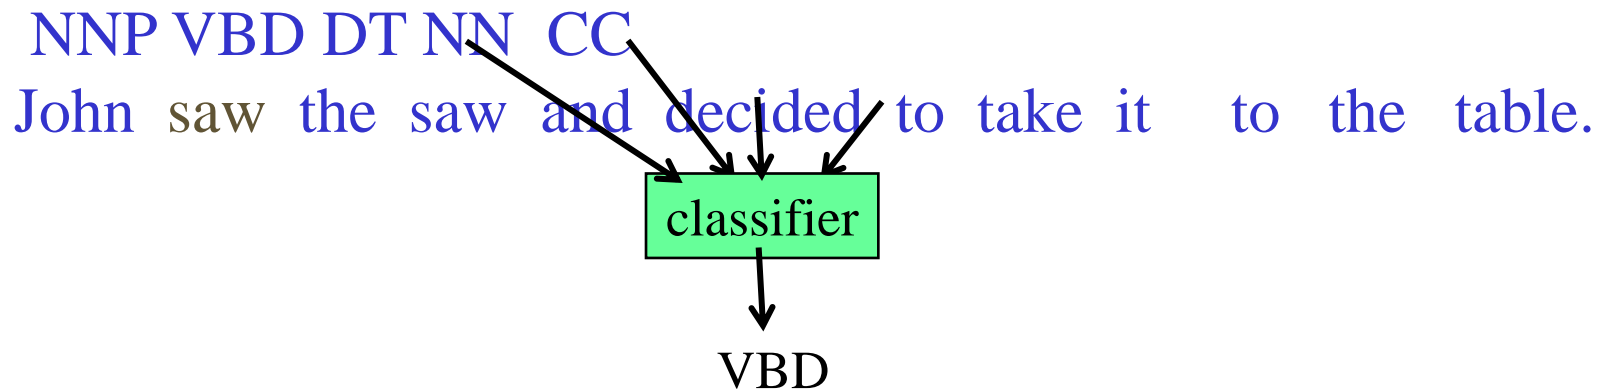


Forward Classification



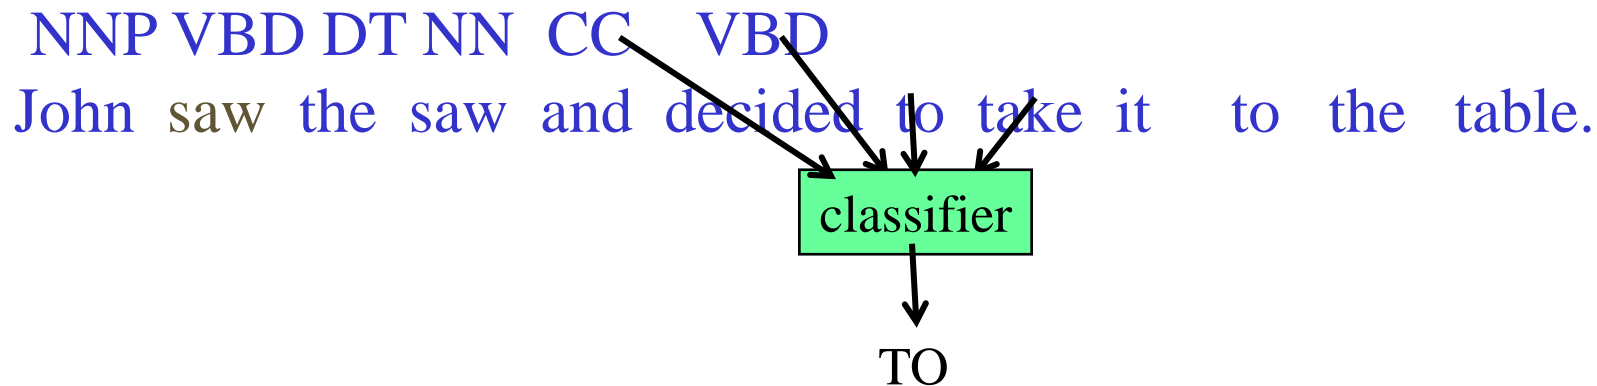


Forward Classification



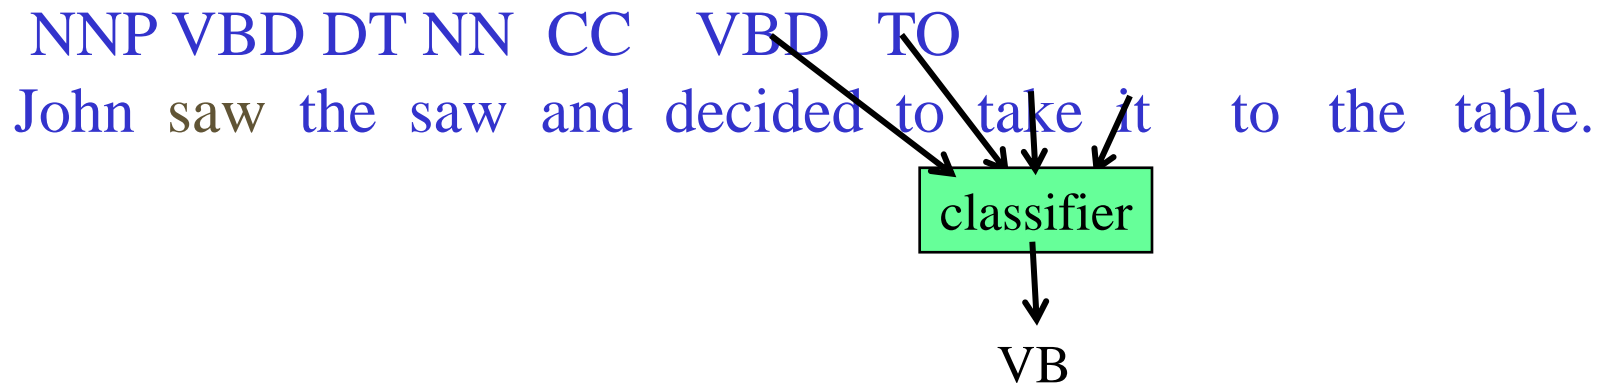


Forward Classification





Forward Classification





Forward Classification

NNP VBD DT NN CC VBD TO VB

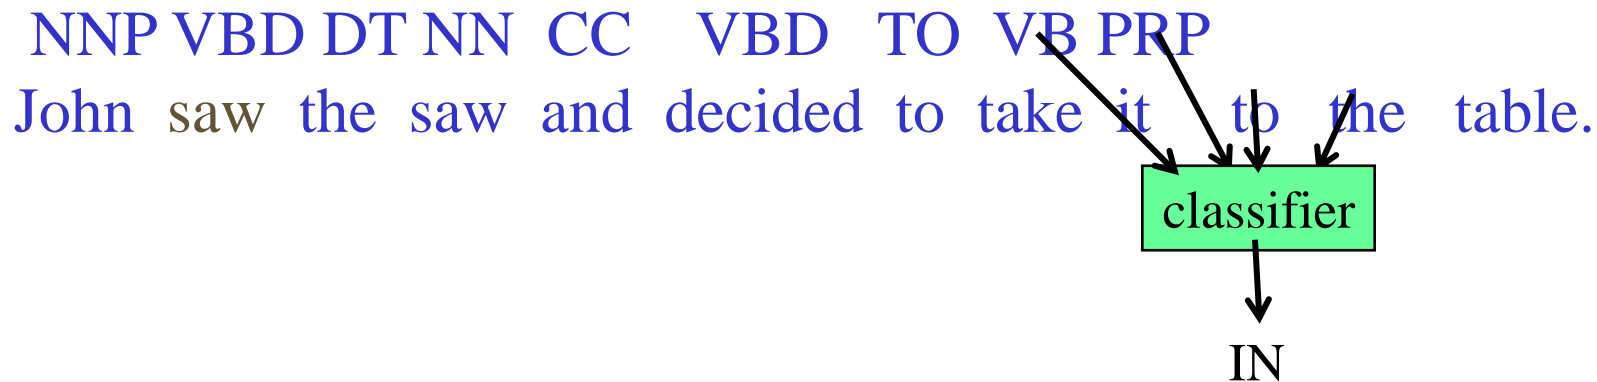
John saw the saw and decided to take it to the table.

classifier

PRP

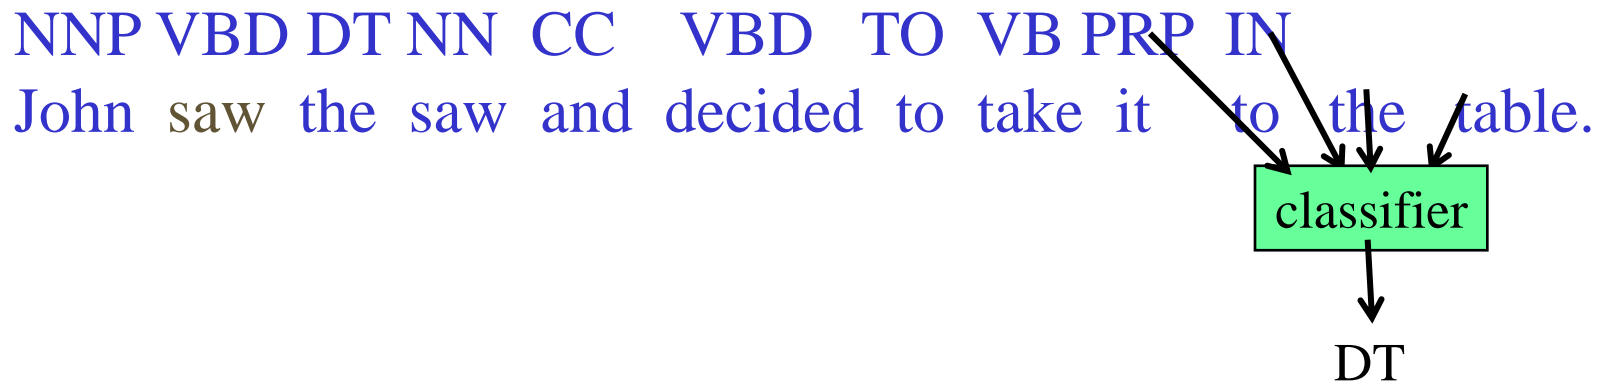


Forward Classification






Forward Classification





Forward Classification





Backward Classification

- Disambiguating “to” in this case would be even easier backward.

John saw the saw and decided to take it to the table.

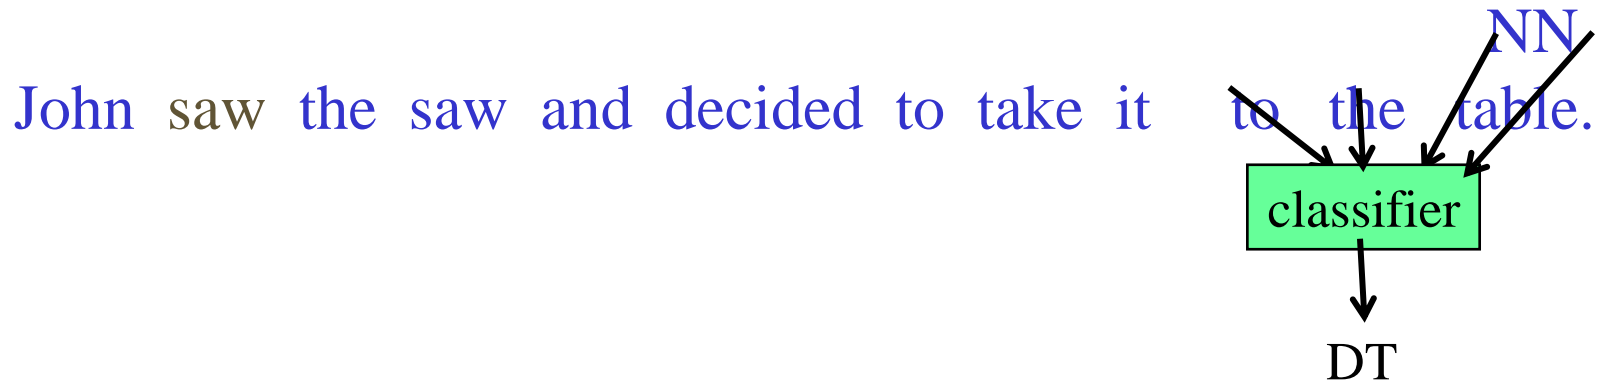
classifier

NN



Backward Classification

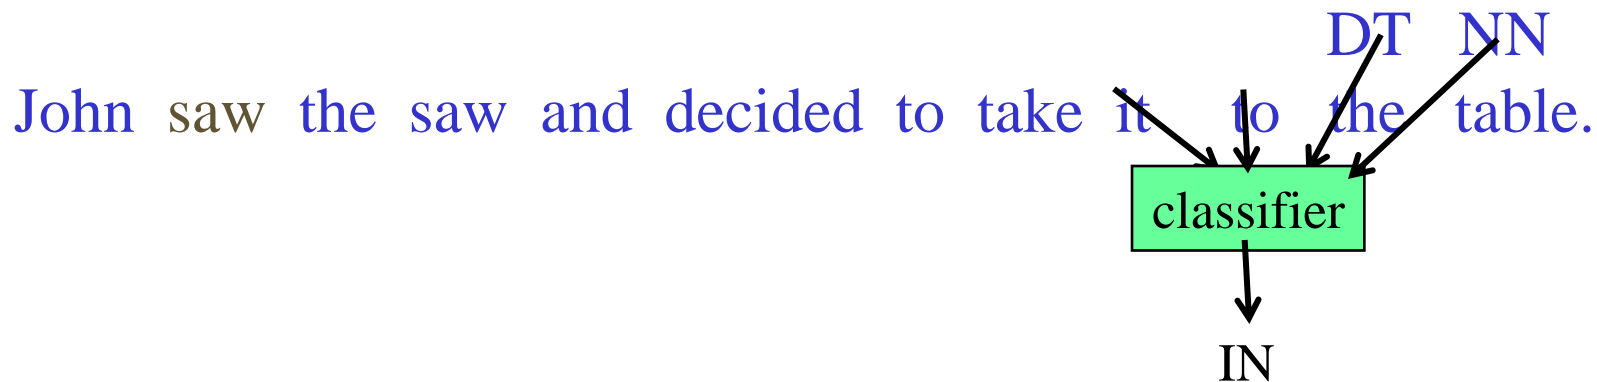
- Disambiguating “to” in this case would be even easier backward.





Backward Classification

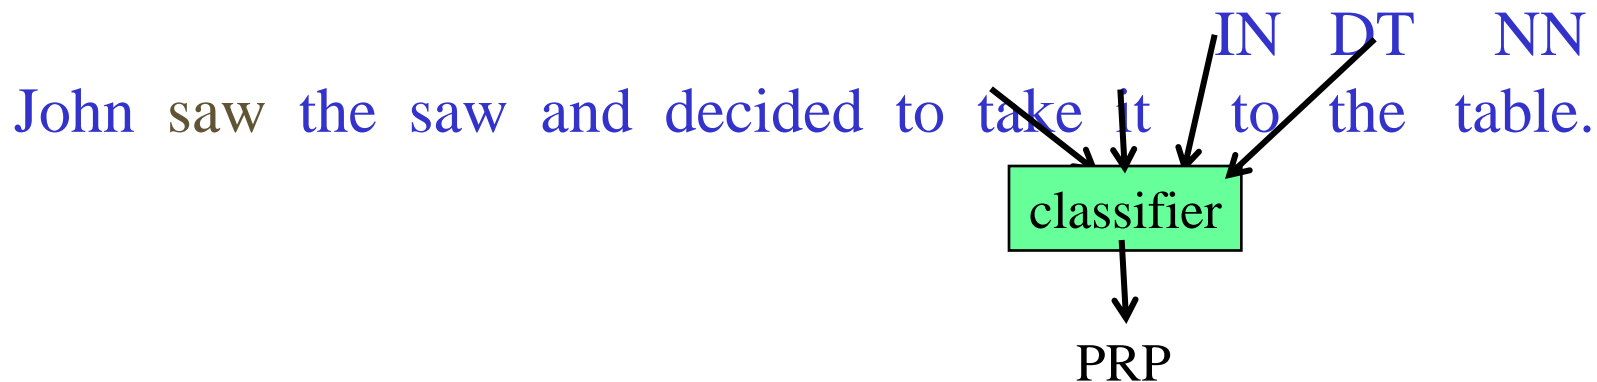
- Disambiguating “to” in this case would be even easier backward.





Backward Classification

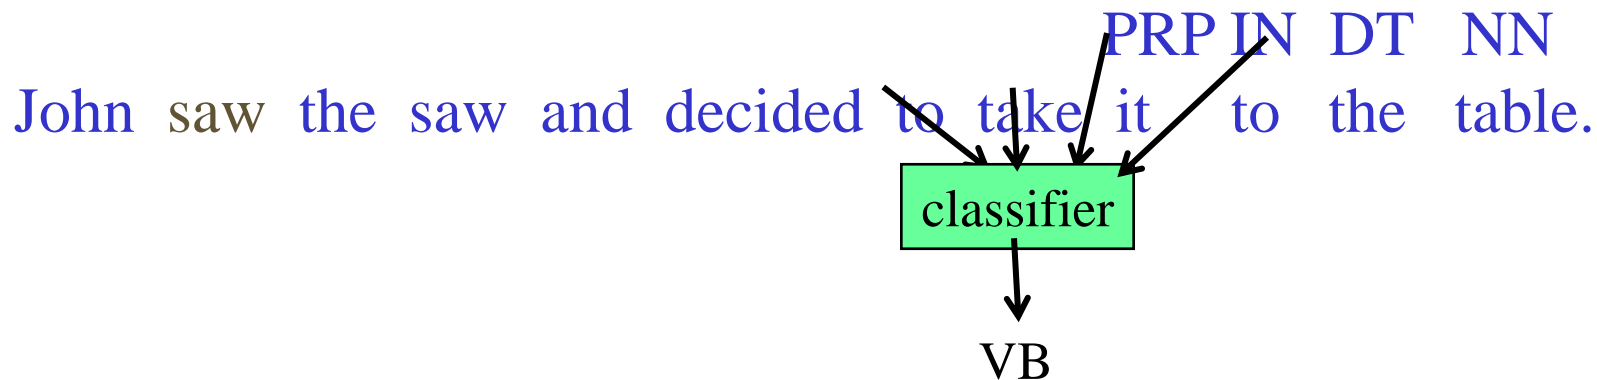
- Disambiguating “to” in this case would be even easier backward.





Backward Classification

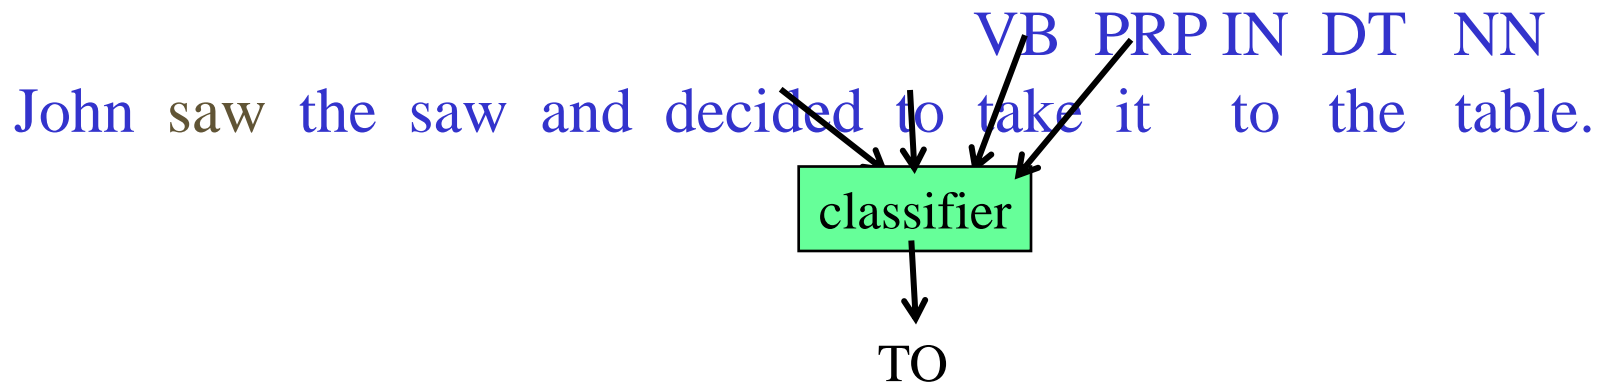
- Disambiguating “to” in this case would be even easier backward.





Backward Classification

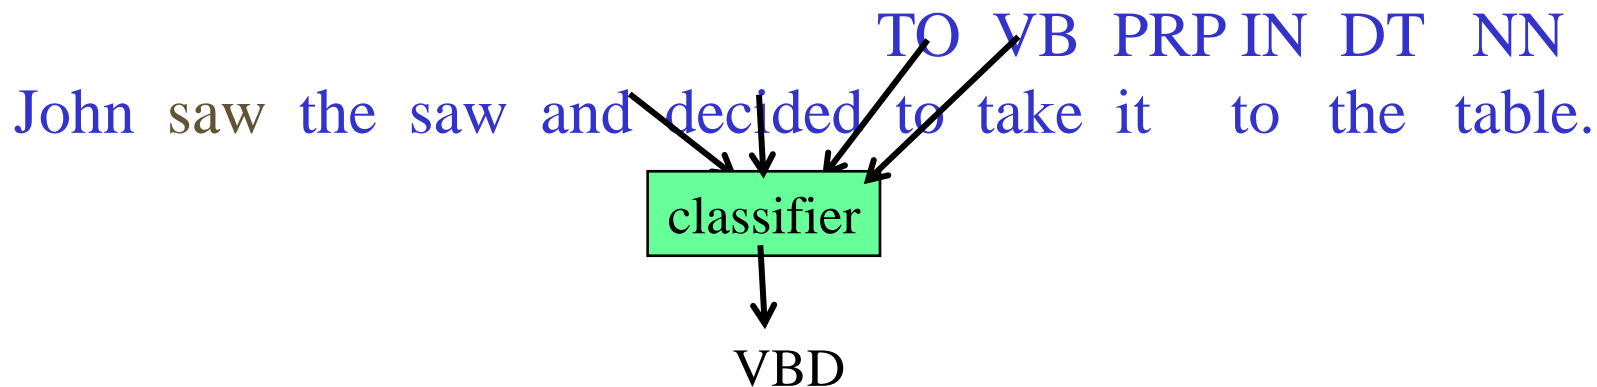
- Disambiguating “to” in this case would be even easier backward.





Backward Classification


- Disambiguating “to” in this case would be even easier backward.





Backward Classification


- Disambiguating “to” in this case would be even easier backward.





Backward Classification

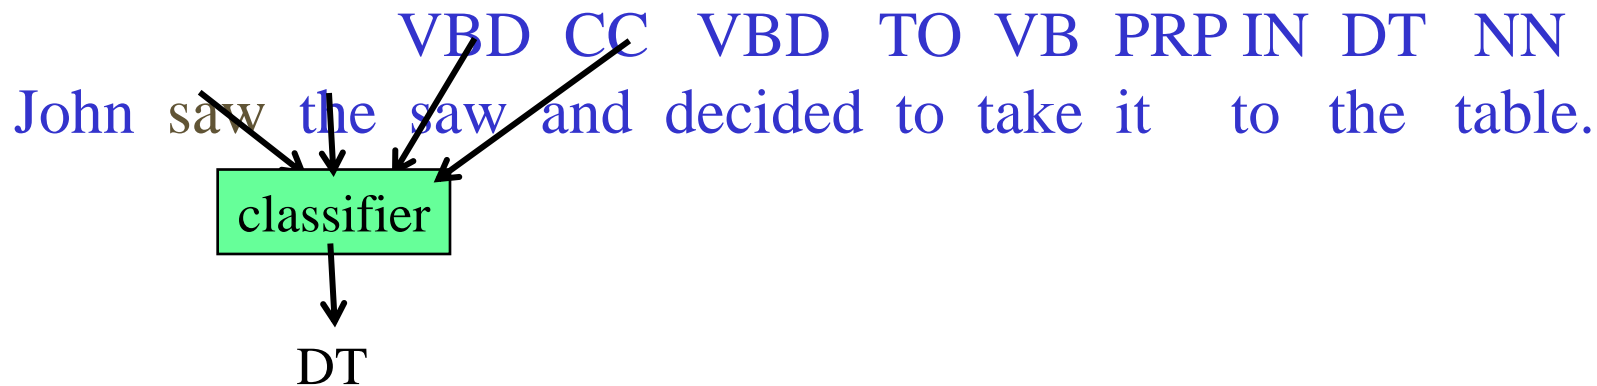
- Disambiguating “to” in this case would be even easier backward.





Backward Classification


- Disambiguating “to” in this case would be even easier backward.





Backward Classification


- Disambiguating “to” in this case would be even easier backward.





Backward Classification

- Disambiguating “to” in this case would be even easier backward.





Overview: POS Tagging Accuracies

- Rough accuracies:

- Most freq tag:

~90% / ~50%

- Trigram HMM:

~95% / ~55%

- Maxent $P(t|w)$:

93.7% / 82.6%

- TnT (HMM++):

96.2% / 86.0%

- MEMM tagger:

96.9% / 86.9%

- Bidirectional dependencies:

97.2% / 90.0%

- Upper bound:


~98% (human agreement)

Most errors
on unknown
words




Tagging Without Sequence Information

Baseline



Three Words



Model	Features	Token	Unknown	Sentence
Baseline	56,805	93.69%	82.61%	26.74%
3Words	239,767	96.57%	86.78%	48.27%

Using words only in a straight classifier works as well as a basic (HMM or discriminative) sequence model!!



Summary of POS Tagging


For tagging, the change from generative to discriminative model **does not by itself** result in great improvement

One profits from models for specifying dependence on **overlapping features of the observation** such as spelling, suffix analysis, etc.

An MEMM allows integration of rich features of the observations, but can suffer strongly from assuming independence from following observations; this effect can be relieved by adding dependence on following words

This additional power (of the MEMM ,CRF, Perceptron models) has been shown to result in improvements in accuracy

The **higher accuracy** of discriminative models comes at the price of **much slower training**



Dependency Parsing

Building structure

Based on slides by Christopher Manning


Participation code: eagles



Dependency Grammar and Dependency Structure

Dependency syntax postulates that syntactic structure consists of lexical items linked by binary asymmetric relations (“arrows”) called dependencies

The arrows are commonly **typed** with the name of grammatical relations (subject, prepositional object, apposition, etc.)






Dependency Grammar and Dependency Structure

Dependency syntax postulates that syntactic structure consists of lexical items linked by binary asymmetric relations (“arrows”) called dependencies

The arrows are commonly **typed** with the name of grammatical relations (subject, prepositional object, apposition, etc.)






Dependency Grammar and Dependency Structure

Dependency syntax postulates that syntactic structure consists of lexical items linked by binary asymmetric relations (“arrows”) called dependencies

The arrows are commonly **typed** with the name of grammatical relations (subject, prepositional object, apposition, etc.)






Dependency Grammar and Dependency Structure

Dependency syntax postulates that syntactic structure consists of lexical items linked by binary asymmetric relations (“arrows”) called dependencies

The arrows are commonly **typed** with the name of grammatical relations (subject, prepositional object, apposition, etc.)






Dependency Grammar and Dependency Structure

Dependency syntax postulates that syntactic structure consists of lexical items linked by binary asymmetric relations (“arrows”) called dependencies

The arrows are commonly **typed** with the name of grammatical relations (subject, prepositional object, apposition, etc.)






Dependency Grammar and Dependency Structure

Dependency syntax postulates that syntactic structure consists of lexical items linked by binary asymmetric relations (“arrows”) called dependencies

The arrows are commonly **typed** with the name of grammatical relations (subject, prepositional object, apposition, etc.)





Methods of Dependency Parsing

1. Dynamic programming (like in the CKY algorithm)

You can do it similarly to lexicalized PCFG parsing: an $O(n^5)$ algorithm

Eisner (1996) gives a clever algorithm that reduces the complexity to $O(n^3)$, by producing parse items with heads at the ends rather than in the middle

2. Graph algorithms

You create a Maximum Spanning Tree for a sentence

McDonald et al.'s (2005) MSTParser scores dependencies independently using a ML classifier (he uses MIRA, for online learning, but it could be MaxEnt)

3. Constraint Satisfaction

Edges are eliminated that don't satisfy hard constraints. Karlsson (1990), etc.

4. "Deterministic parsing"

Greedy choice of attachments guided by machine learning classifiers

MaltParser (Nivre et al. 2008) – discussed in the next segment



Dependency Conditioning Preferences

What are the sources of information for dependency parsing?

1. Bilexical affinities [girl → the] is plausible
2. POS tags [VB → NN] is plausible
3. Dependency distance mostly with nearby words
4. Intervening material

Dependencies rarely span intervening verbs or punctuation

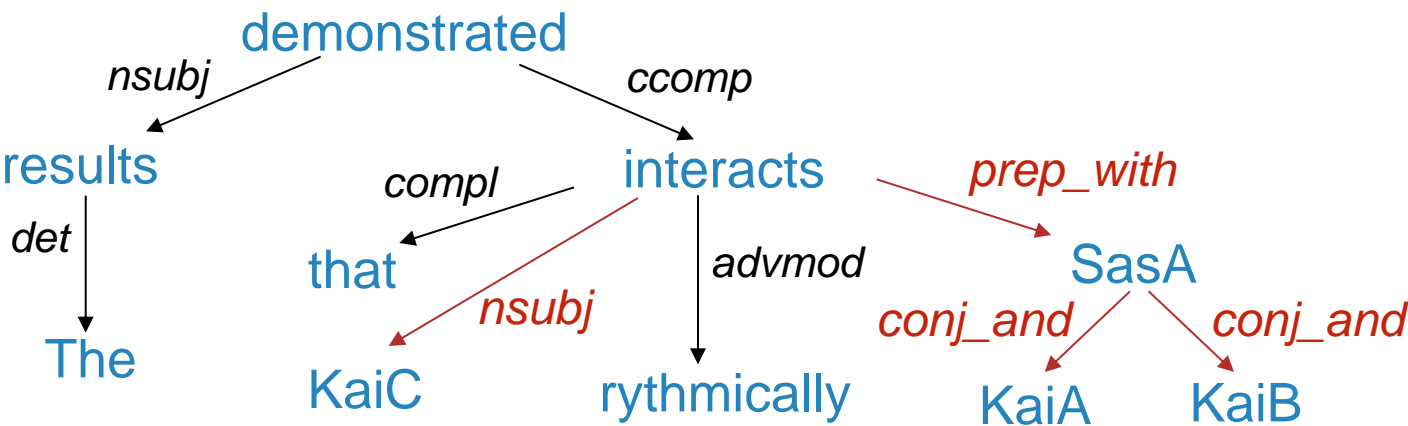
5. Valency of heads

How many dependents on which side are usual for a head?



Dependency paths identify relations like protein interaction

[Erkan et al. EMNLP 07, Fundel et al. 2007]



KaiC \leftarrow nsubj interacts prep_with \rightarrow SasA
KaiC \leftarrow nsubj interacts prep_with \rightarrow SasA conj_and \rightarrow KaiA
KaiC \leftarrow nsubj interacts prep_with \rightarrow SasA conj_and \rightarrow KaiB



BioNLP 2009/2011 relation extraction shared tasks

[Björne et al. 2009]

